

Using the Debugger

Michael Jantz

Dr. Prasad Kulkarni

Debugger

- What is it
 - a powerful tool that supports examination of your program during execution.
- Idea behind debugging programs.
 - Creates additional symbol tables that permit tracking program behavior and relating it back to the source files
- Some common debuggers for UNIX/Linux
 - gdb, sdb ,dbx etc.

GDB

gdb is a tool for debugging C & C++ code

Some capabilities:

- run a program
- stop it on any line or the start of a function
- examine various types of information like values of variables, sequence of function calls
- stop execution when a variable changes
- change values of variables (during execution)
- call a function at any point in execution

Compilation for gdb

- Code must be compiled with the **-g** option
 - `gcc -g -o file1 file1.c file2.c file3.o`
- Which files can you debug?
 - You can debug file1.c, file 2.c, (not file3.o)
- Optimization is not always compatible
 - Due to optimizations which rearrange portions of the code

Building and Testing *bash*

1) Untar and navigate to the bash-4.2 directory:

```
> tar xvzf eecs678-lab-gdb.tar.gz
```

```
> cd gdb/bash-4.2
```

2) Configure *bash* for the build:

```
> ./configure
```

3) Make *bash* using multiple jobs and with CFLAGS=-g

```
> make -j8 CFLAGS=-g
```

4) Test the *bash* executable:

```
> ./bash --version
```

Using GDB with *bash*

Starting GDB:

```
> gdb program
```

The build process created an executable file named *bash*

To start *bash* under GDB do:

```
gdb ./bash-4.2/bash
```

Breakpoints

- break (b)
 - Sets a breakpoint in program execution
 - tbreak (tb) temporary breakpoint. Exists until it is hit for the first time
- Breakpoint syntax
 - **b** *line-number*
 - **b** *function-name*
 - **b** *line-or-function* **if** *condition*
 - **b** *filename: line number*
- info breakpoints – Gives information on all active breakpoints
- delete (d) – Deletes the specified breakpoint number (e.g., d 1)

A Breakpoint in *bash*

- *bash* is a shell program. It provides a convenient command line interface to the OS, interprets commands, sets up pipelines, manages multiple jobs, etc.
- Debugging a running instance of *bash* can help you learn how the shell operates
- As an example, say you want to learn about how *bash* handles and executes commands. Place a breakpoint at the *execute_command* function:

```
gdb> b execute_command
```


Running *bash* Under GDB

run (r) - runs the loaded program under GDB

Can also specify arguments and I/O redirection now, e.g:

```
gdb> r arg1 arg2 < input > output
```

In our case, we can run a script with our *bash* executable:

```
gdb> r ./finder.sh bash-4.2/ execute 20
```

finder.sh

```
find $1 -name '*.ch' | xargs grep -c $2 | sort -t : +1.0 -2.0 --numeric --reverse | head --lines=$3
```

- `find $1 -name '*.ch'` – Find files with .c and .h extensions under the directory given by the first argument.
- `xargs grep -c $2` – Search the set of files on standard input for the string given by the second argument. -c says that instead of printing out each usage in each file, give me the number of times \$2 is used in each file.
- `sort` – Sort standard input and print the sorted order to standard output. -t : +1.0 -2.0 says sort using the second column on each line (delimited by the ':' character) as a key. --numeric says to sort numerically (as opposed to alphabetically). --reverse says sort in reverse order.
- `head` – print only the first *n* lines of standard input. --lines=\$3 lets us set the number of lines with the third argument.

Common GDB Commands

- When you hit the breakpoint you should see:

Breakpoint 1, execute_command (command=0x724088) at execute_cmd.c:376

- Now, gdb has stopped execution of *bash*. Try the following commands:
 - list (l) will list the source code around where execution has stopped. Alternatively, l *n,m* will display the source code in between two given line numbers *n* & *m*.
 - backtrace (bt) prints a backtrace of all stack frames. From this, you can tell how you got to where you are from the main. The output here says you are in `execute_command`, which was called from `reader_loop`, which was called from the main entry point.

Using the Frame Stack

- GDB currently has the `execute_command` frame selected. Use the `info` command to list information about the frame
 - `info args` – print the arguments passed into this frame
 - `info locals` – print the local arguments for this frame
 - `help info` – shows you everything info can tell you
- Additionally, print information about other stack frames using
 - `up [n]` – Select the frame `n` levels up in the call stack (towards `main`). `n=1` if not specified.
 - `down [n]` – Select the frame `n` levels down in the call stack (you must have used `up` in order to come back down)
 - After you select a new frame, use `info` as described above to display information about the frame

Control Flow

- `continue (c)`
 - Continue until the next breakpoint is reached, the program terminates, or an error occurs (Don't use this just yet, we've got a few more commands to try).
- `next (n)`
 - Execute one instruction. Step over function calls.
- `step (s)`
 - Execute one instruction. Step into function calls.
- `kill (k)`
 - Kills the program being debugged (does not exit gdb – preserves everything else from the session, i.e., breakpoints.)

Inspecting and Assigning

- Continue to the end of the `execute_command` function:
 - `finish (fin)` – continue to the end of the function you're currently broken in
- Now, if you read the code in this function, it calls `execute_command_internal` with the current *command*. To look at *command*'s properties (or any object's) use the following:
 - `print (p) t` – Prints the value of some variable
 - `whatis t` – Prints the type of *t*
 - `ptype t` – Prints fields for the type *t*

Inspecting and Assigning (cont.)

- Try these commands to inspect the command object:
 - `gdb> whatis command` – tells us the type of *command*.
 - `gdb> ptype command` – displays all the fields the command type
 - `gdb> p command->value` – prints the value of *command*->*value*
- You can also assign values in gdb:
 - `gdb> set var command=0x0` – sets the *command* pointer to 0x0

Printing Examples

(gdb) p command

\$14 = (COMMAND *) 0x724088

(gdb) ptype command

```
type = struct command {  
    enum command_type type;  
    int flags;  
    int line;  
    REDIRECT *redirects;  
    union {  
        struct for_com *For;  
        ...  
        struct coproc_com *Coproc;  
    } value;  
} *
```


Printing Examples

(gdb) p command->type

\$15= cm_connection

(gdb) p (struct connection *) command->value

\$16 = (struct connection *) 0x724048

(gdb) ptype ((struct connection *) command->value)

type = struct connection {

int ignore;

COMMAND *first;

COMMAND *second;

int connector;

} *

Printing Examples

```
(gdb) p ((struct connection *) command->value)->first  
$17 = (COMMAND *) 0x721108
```

```
(gdb) p ((struct connection *) command->value)->first->type  
$18 = cm_simple
```

```
(gdb) ptype ((struct simple_com *) ((struct connection *) command->value)->first)  
type = struct simple_com {  
    int flags;  
    int line;  
    WORD_LIST *words;  
    REDIRECT *redirects;  
} *
```

Printing Examples

```
(gdb) p ((struct simple_com *) ((struct connection *) command->value)->first)->words  
$19 = (WORD_LIST *) 0xdfdfdfdfdfdfdfdf
```

```
(gdb) ptype ((struct simple_com *) ((struct connection *) command->value)->first)->words  
type = struct word_list {  
    struct word_list *next;  
    WORD_DESC *word;  
} *
```

```
(gdb) ptype ((struct simple_com *) ((struct connection *) command->value)->first)->words->word  
type = struct word_desc {  
    char *word;  
    int flags;  
} *
```

```
(gdb) p ((struct simple_com *) ((struct connection *) command->value)->first)->words->word  
Cannot access memory at address 0xdfdfdfdfdfdfdf7
```

Calling Functions from GDB

- The *call* command allows you to call other functions in your code within GDB.
- Very useful for printing complicated data structures within the debugger

Call Example

(gdb) b execute_simple_command

Breakpoint 2 at 0x4380a4: file execute_cmd.c, line 3650.

(gdb) c

Continuing.

(gdb) p simple_command

\$28 = (SIMPLE_COM *) 0x721148

(gdb) p simple_command->words

\$29 = (WORD_LIST *) 0x721fe8

(gdb) call _print_word_list(simple_command->words, " ", printf)

(gdb) call fflush(stdout)

find \$1 -name '*'.[ch]\$30 = 0

GDB References

- The Unix manual is a good quick reference for common GDB commands:
 - At a terminal, type: `man gdb`
- While running GDB, `help` will give you any information you need for any command:
 - `help (h) command`
- If you need to do some heavy lifting with GDB, the official documentation for users is at this website:
 - http://sourceware.org/gdb/current/onlinedocs/gdb_toc.html

Backup

Multiple Threads

- GDB also has a set of commands for finer control of multi-threaded applications.
- GDB comes with these commands for controlling multiple threads:
 - `info threads` – Print a numbered list of all current threads and their contexts. An asterisk denotes the thread on which GDB is currently focused.
 - `thread <thread #>` - Switch focus to the thread numbered <thread #>.
 - `thread apply (all | <thread # list>) cmd` – Apply cmd to all threads or each thread in the <thread # list>.
- For example, *thread apply all bt* shows the stack trace for each thread.

Automatic Source Navigation

- GDB comes with a tool for automatic source navigation called the Text User Interface (TUI).
- To access the TUI, do C-x, C-a in the shell running GDB.
- It should split the terminal. Now, when you run your program under GDB, the source will be displayed in the screen above your command line.
- To switch between control of the the source code screen and the command line do: C-x, o.
- Alternatively, if you would like a more graphical user interface, you can use the DDD debugger (which is essentially identical to the TUI, but provides more buttons and mouse over actions).

Valgrind

- A heavyweight tool for dynamically catching hard-to-detect errors in programs.
 - Only checks code that it runs
- Valgrind is a virtual machine and quite a bit slower than normal execution.
 - Translates binaries into an intermediate representation performs some modifications and recompiles the transformations at runtime as needed.
 - Allows it to instrument code with safe guards profiling tools.
- Valgrind's main tool is a memory checker that detects memory errors in programs
- Also has [a few other useful tools](#) such as cache and branch miss-prediction profilers, a heap profiler, a multi-thread data race detector, and more.

Valgrind

- Memory checker (memcheck) helps find:
 - Memory leaks
 - Usage of uninitialized variables
 - Often the cause of non-deterministic behavior in single threaded software.
 - Bad frees of heap blocks
 - Double frees
 - Mismatched frees (frees on addresses without associated mallocs)
 - Accesses to unallocated memory
 - Accesses to invalid stack and heap addresses
 - Freed heap addresses
 - Out of bounds heap addresses
 - Out of bounds stack addresses

Valgrind - Memcheck

- To run valgrind use:
 - “valgrind [optional flags] <executable>”
- By default only the memcheck tool is used.

```
1#include <stdlib.h>
2
3void* still_reachable;
4void* possibly_lost;
5
6int main() {
7    int uninitialized_variable; // This variable is never given a value.
8
9    for (; uninitialized_variable < 100; uninitialized_variable++) {
10        void** definitely_lost = (void**) malloc(sizeof(void*)); // allocate a
11                                                                    // pointer on the
12                                                                    // heap.
13
14        *definitely_lost = (void*) malloc(7); // Give the pointer something else to
15                                                // point to on the heap. This will be
16                                                // indirectly lost.
17    }
18
19    // At this point, definitely_lost is out of scope and we can no longer free
20    // it. The pointer pointed to by definitely_lost is indirectly lost since we
21    // were only able to reach the pointer through definitely_lost.
22
23    still_reachable = malloc(42); // This value is never freed but is pointed to
24                                // in the global scope at program completion.
25
26    possibly_lost = malloc(10);
27    possibly_lost += 4; // This is similar to still_reachable except there is a
28                        // pointer pointing to the middle of the allocated block
29                        // but nothing points to the front of the block. This is
30                        // very odd behavior and usually is a memory leak (but not
31                        // always).
32}
```

valgrind_test.c – a bad program that compiles without warnings (gcc -Wall -g valgrind_test.c)

```
[jrobinson@localhost Development]$ valgrind ./test
==2606== Memcheck, a memory error detector
==2606== Copyright (C) 2002-2015, and GNU GPL'd, by Julian Seward et al.
==2606== Using Valgrind-3.11.0 and LibVEX; rerun with -h for copyright info
==2606== Command: ./test
==2606==
==2606== Conditional jump or move depends on uninitialised value(s)
==2606==    at 0x40056A: main (test.c:9)
==2606==
==2606== HEAP SUMMARY:
==2606==    in use at exit: 1,553 bytes in 202 blocks
==2606==    total heap usage: 202 allocs, 0 frees, 1,553 bytes allocated
==2606==
==2606== LEAK SUMMARY:
==2606==    definitely lost: 800 bytes in 100 blocks
==2606==    indirectly lost: 700 bytes in 100 blocks
==2606==    possibly lost: 10 bytes in 1 blocks
==2606==    still reachable: 43 bytes in 1 blocks
==2606==    suppressed: 0 bytes in 0 blocks
==2606== Rerun with --leak-check=full to see details of leaked memory
==2606==
==2606== For counts of detected and suppressed errors, rerun with: -v
==2606== Use --track-origins=yes to see where uninitialised values come from
==2606== ERROR SUMMARY: 101 errors from 1 contexts (suppressed: 0 from 0)
```

Valgrind shows all of the problems with this code 28

Valgrind - Memcheck

- Rerun with “--leak-check=full” and “--track-origins=yes” as suggested by memcheck.

```
[jrobinson@localhost Development]$ valgrind --leak-check=full --track-origins=yes ./test
==16850== Memcheck, a memory error detector
==16850== Copyright (C) 2002-2015, and GNU GPL'd, by Julian Seward et al.
==16850== Using Valgrind-3.11.0 and LibVEX; rerun with -h for copyright info
==16850== Command: ./test
==16850==
==16850== Conditional jump or move depends on uninitialised value(s)
==16850==    at 0x40056A: main (test.c:9)
==16850==    Uninitialised value was created by a stack allocation
==16850==    at 0x400536: main (test.c:6)
==16850==
==16850== HEAP SUMMARY:
==16850==    in use at exit: 1,552 bytes in 202 blocks
==16850==    total heap usage: 202 allocs, 0 frees, 1,552 bytes allocated
==16850==
==16850== 10 bytes in 1 blocks are possibly lost in loss record 1 of 4
==16850==    at 0x4C28C50: malloc (in /usr/lib64/valgrind/vgpreload_memcheck-amd64-linux.so)
==16850==    by 0x400586: main (test.c:26)
==16850==
==16850== 1,500 (800 direct, 700 indirect) bytes in 100 blocks are definitely lost in loss record 4 of 4
==16850==    at 0x4C28C50: malloc (in /usr/lib64/valgrind/vgpreload_memcheck-amd64-linux.so)
==16850==    by 0x400549: main (test.c:10)
==16850==
==16850== LEAK SUMMARY:
==16850==    definitely lost: 800 bytes in 100 blocks
==16850==    indirectly lost: 700 bytes in 100 blocks
==16850==    possibly lost: 10 bytes in 1 blocks
==16850==    still reachable: 42 bytes in 1 blocks
==16850==    suppressed: 0 bytes in 0 blocks
==16850== Reachable blocks (those to which a pointer was found) are not shown.
==16850== To see them, rerun with: --leak-check=full --show-leak-kinds=all
==16850==
==16850== For counts of detected and suppressed errors, rerun with: -v
==16850== ERROR SUMMARY: 103 errors from 3 contexts (suppressed: 0 from 0)
```

Valgrind - Memcheck

- Rerun with “--show-leak-kinds=all” appended to the options to see the still reachable blocks.
- Valgrind is not perfect.
 - It will only catch errors that only occur during a program run.
 - If an error inducing code is not reached or specific input conditions that trigger an error are not met, then Valgrind will not detect it in that run.
 - Importance of thorough unit tests.
 - Custom memory allocators can confuse Memcheck.