

EECS 678: Introduction to Operating Systems


Instructor: Heechul Yun

About Me

- Heechul Yun, Assistant Prof., Dept. of EECS
 - Office: 3040 Eaton, 236 Nichols
 - Email: heechul.yun@ku.edu
- Research Areas
 - Operating systems and architecture support for embedded/real-time systems
 - To improve time predictability, energy efficiency, and throughput
 - Multicore, memory systems
- Previously
 - Systems software engineer, Samsung Electronics, Nvidia
 - mainly worked on Linux kernel
- More Information
 - <http://itcc.ku.edu/~heechul>

About This Class

- Textbook: Operating System Concepts
- Objectives: Learn OS basics and practical system programming skills
 - Understand *how it works!*
- Audience: Senior and Junior undergraduate (grad students)
- Course website:
<http://ittc.ku.edu/~heechul/courses/eecs678/>



About This Class


- Course structure
 - Lecture
 - Office hour: MF 11:00 – 11:50 @ 3040 Eaton
 - Discuss OS concepts and the design of major OS components
 - Lab
 - Hands-on system programming experiences.
 - Each lab includes lab discussion and an assignment
 - TA will help you better understand the concepts learned during the lecture.
 - Programming projects
 - Design and implement some parts of OS (e.g., shell, scheduler)
 - 3 projects are expected, each will be given 2~3 weeks to finish
 - To do in groups of two persons. Solo project needs permission

About This Class

- Grading
 - Exam: 50% (mid:20%, final:30%)
 - Quiz: 5%
 - Lab: 15%
 - Programming projects: 30%

About This Class

- Late submissions
 - 20% off / day
- Cheating
 - You can discuss about code and help find bugs of your peers. However, copying another's code or writing code for someone else is cheating and, if identified, **the involved students will be notified to the department chair**



Operating Systems Are Everywhere

- Computers
- Smart phones
- Cars
- Airplanes
- ...
- Almost everything




What is an Operating System?



QNX SOFTWARE SYSTEMS

What is an Operating System?

- A program that acts as an intermediary **between users** and the computer **hardware**

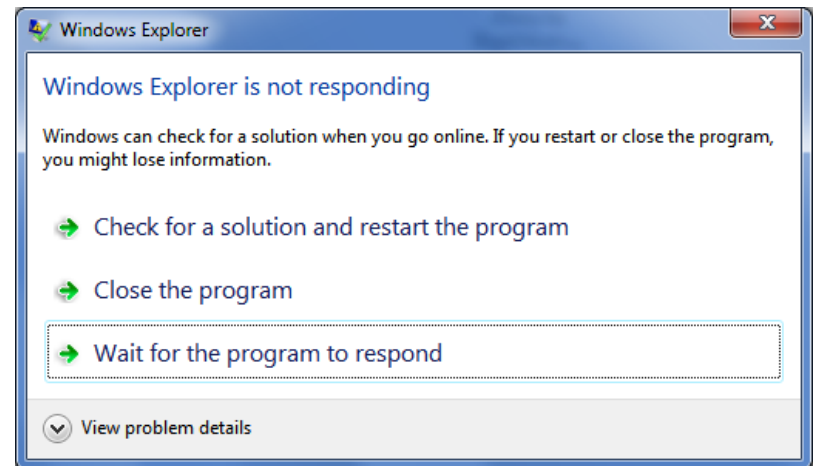


What is an Operating System?

- An easy to use virtual machine – **User's view**
 - Hide complex details for you.
 - What CPU am I using? Intel or AMD?
 - How much memory do I have?
 - Where and how to store my data on the disk?
 - Provide APIs and services
 - read(...), write(..)
 - Virtual memory, filesystems, ...


What is an Operating System?

- A resource manager – **System's view**
 - Make everybody get a fair share of resources
 - Time and space multiplexing hardware resources
 - Monitor/prevent error or improper use



What is an Operating System?

- Is an internet browser part of an OS?
 - Everything that shipped by the OS vendor?
 - What about 'solitaire'?



- The program that always runs
 - Typically in kernel mode (we will learn it later)

Why Needed?

- Programmability
 - You don't need to know hardware details to do stuffs
- Portability
 - You can run the same program on different hardware configurations
- Safety
 - The OS protects your program from faults in other programs
- Efficiency
 - Multiple programs/users can share the same hardware efficiently

What to Study?

- Not “how to use”
 - I’m sure you know better than me about how to use the iOS in your iPhone.
- But “how it works!”
 - We will study the underlying concepts, standard OS components and their designs

OS Design Issues





- Structure
 - How to organize the OS?
- Communication
 - How to exchange data among different programs?
- Performance
 - How to maximize/guarantee performance and fairness?
- Naming
 - How to name/access resources ?
- Protection
 - How to protect with each other?
- Security
 - How to prevent unauthorized access?
- Reliability
 - How to prevent system crash?

Why Study?

- I'm a user
 - Have you ever wondered how it works?
 - You can better tune the OS to improve performance (or save energy)
- I'm a system programmer
 - You can write more efficient programs by knowing how the OS works.
- I'm a hacker
 - You need to know the enemy (the OS) to beat it

Brief History of Computers

- Early computing machines
 - Babbage's analytical engine
 - First programmer: Ada Lovelace
- Vacuum tube machines
 - 1940s ~ 1950s
 - Used to break code in WWII
 - No OS, No PL




Brief History of Computers

- Vacuum tubes → Transistors → IC → VLSI
 - Smaller, faster, and more reliable
 - Enable smaller computers
- 1960s Mainframes
- 1970s Minicomputers
- 1980s Microprocessor, Apple, IBM PC
- 1990s PC, Internet
- 2000s Cloud computing
- 2010s Mobile, Internet-of-things (IoT)

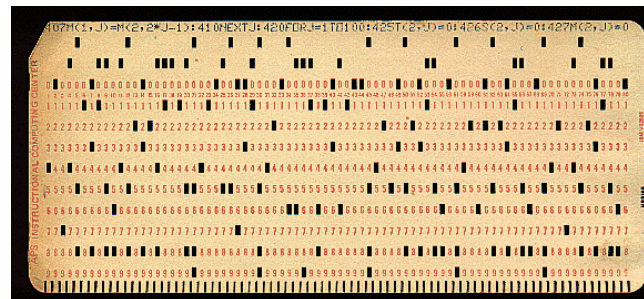


Evolution of Operating Systems

- Batch systems
 - Each user submits her job on punch cards
 - Collect a batch of jobs, read the batch before start processing
 - The 'OS' processes each job at a time
 - Problems
 - No interactivity
 - CPU is underutilize to wait I/O operations



IBM 029 card punch



Evolution of Operating Systems

- Multiprogramming
 - Multiple runnable jobs at a time
 - I/O and compute can overlap
 - OS goal: maximize system throughput
 - IBM OS/360



Evolution of Operating Systems

- Timesharing
 - Multiple interactive users sharing a machine
 - Each user accesses the machine via a terminal
 - Provide each user an illusion of using the entire machine
 - OS goal: optimize response time
 - UNIX

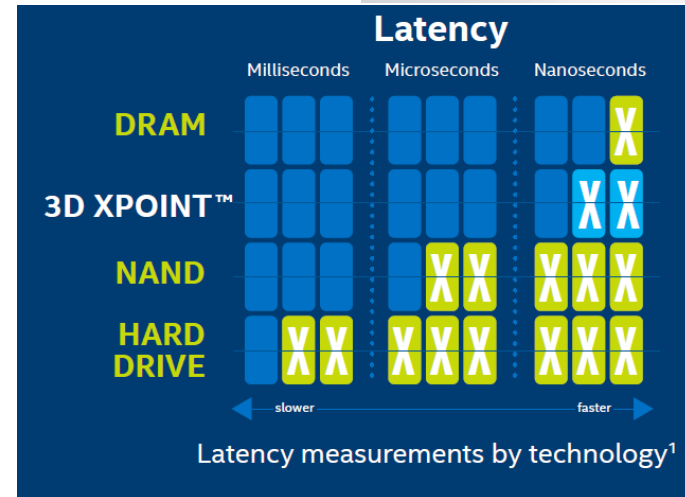



Evolution of Operating Systems

- Parallel computing
 - Use multiple CPUs/cores to speed up performance
 - OS goal: fast synchronization, max utilization
- Distributed computing
 - Physically separate networked computers
- Virtualization
 - Multiple OSes on a single machine

Challenges for Future OS

- New kinds of hardware are keep coming
 - Heterogeneous multicore processors (e.g., ARM big.LITTLE)
 - Storage Class Memory (SCM): non-volatile DRAM-like memories
- New computing paradigms
 - Cloud computing
 - Internet-of-Things (IoT)



Summary

- In this class, you will learn
 - Major OS components
 - Their structure, interface, mechanisms, policies, and algorithms
- This class will (hopefully) help you
 - Understand the foundation of computing systems
 - Understand various engineering trade-offs in designing complex systems you would build in future

Computer Architecture and OS

Recap

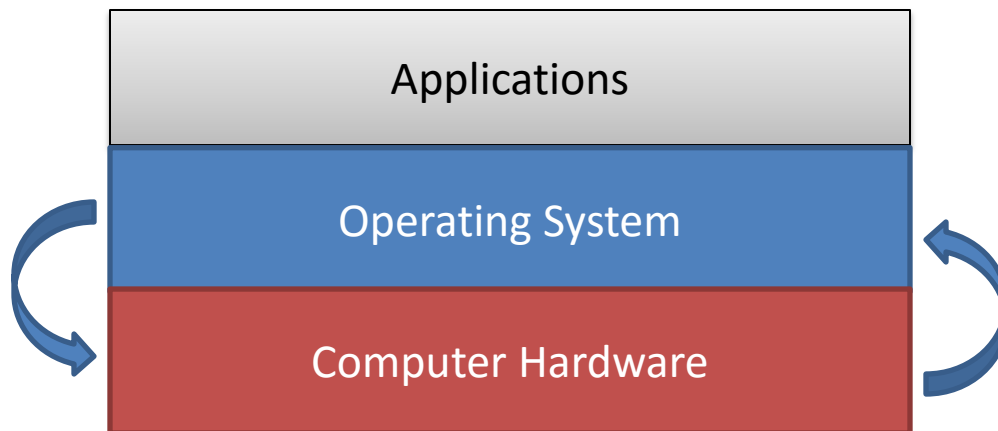
- What is an OS?
 - An intermediary between users and hardware
 - A program that is always running
 - A resource manager
 - Manage resources efficiently and fairly
 - A easy to use virtual machine
 - providing APIs and services

Agenda

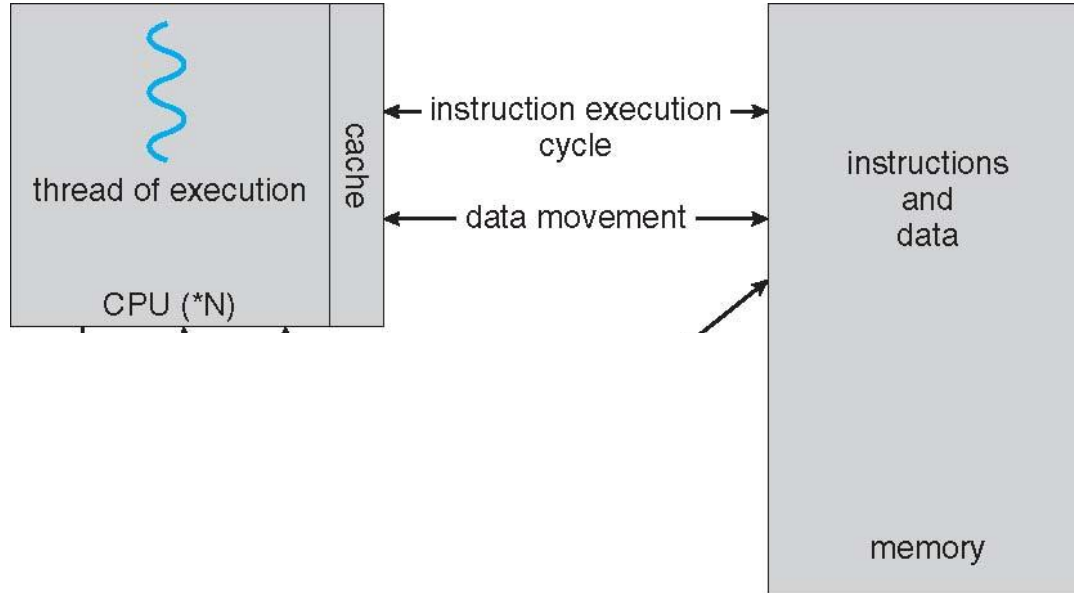
- Computer architecture and OS
 - CPU, memory, disk
 - Architecture trends and their impact to OS
 - Architectural support for OS

Computer Architecture and OS

- OS talks to hardware
 - OS needs to know the hardware features
 - OS drives new hardware features



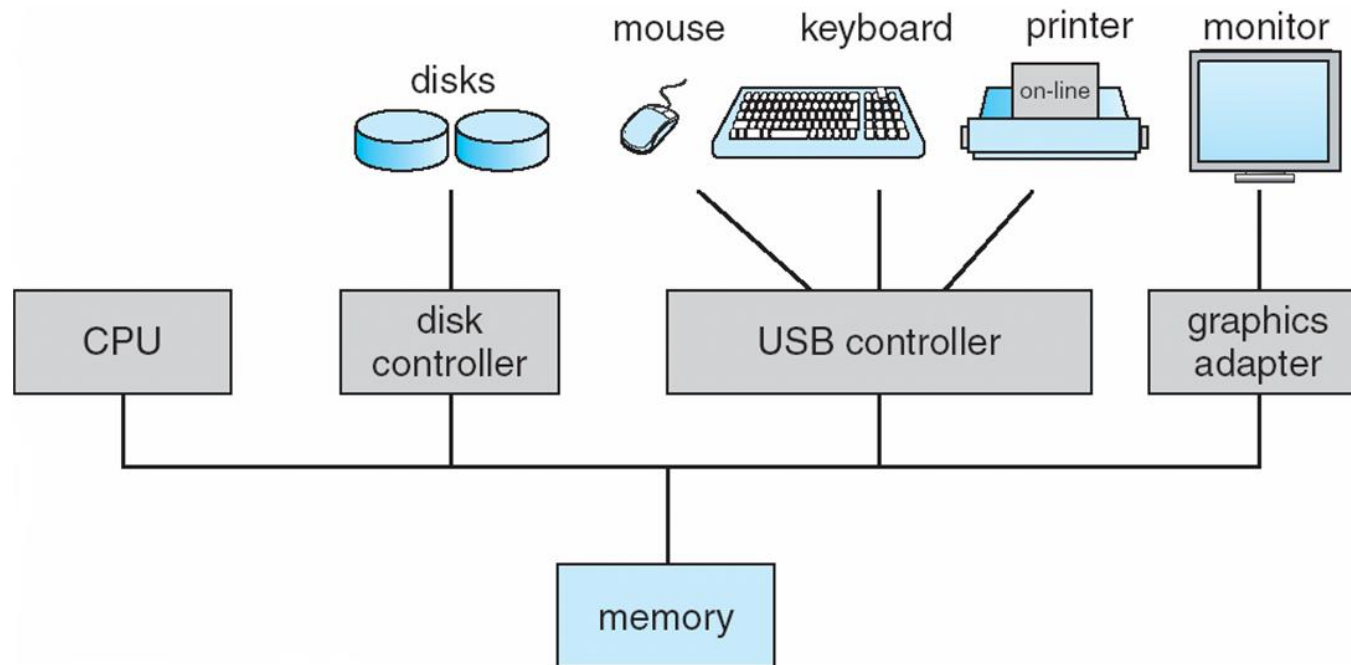
Simplified Computer Architecture



A von Neumann architecture


A Computer System

- **Essentials: CPU, Memory, Disk**
- Others: graphic, USB, keyboard, mouse, ...




Central Processing Unit (CPU)

- The brain of a computer
 - Fetch instruction from memory
 - Decode and execute
 - Store results on memory/registers
- Moore's law
 - Transistors double every 1~2yr
 - 5.56 billion in a 18-core Intel Xeon Haswell-E5




H Sutter, "The Free Lunch Is Over", Dr. Dobb's Journal, 2009

Single-core CPU



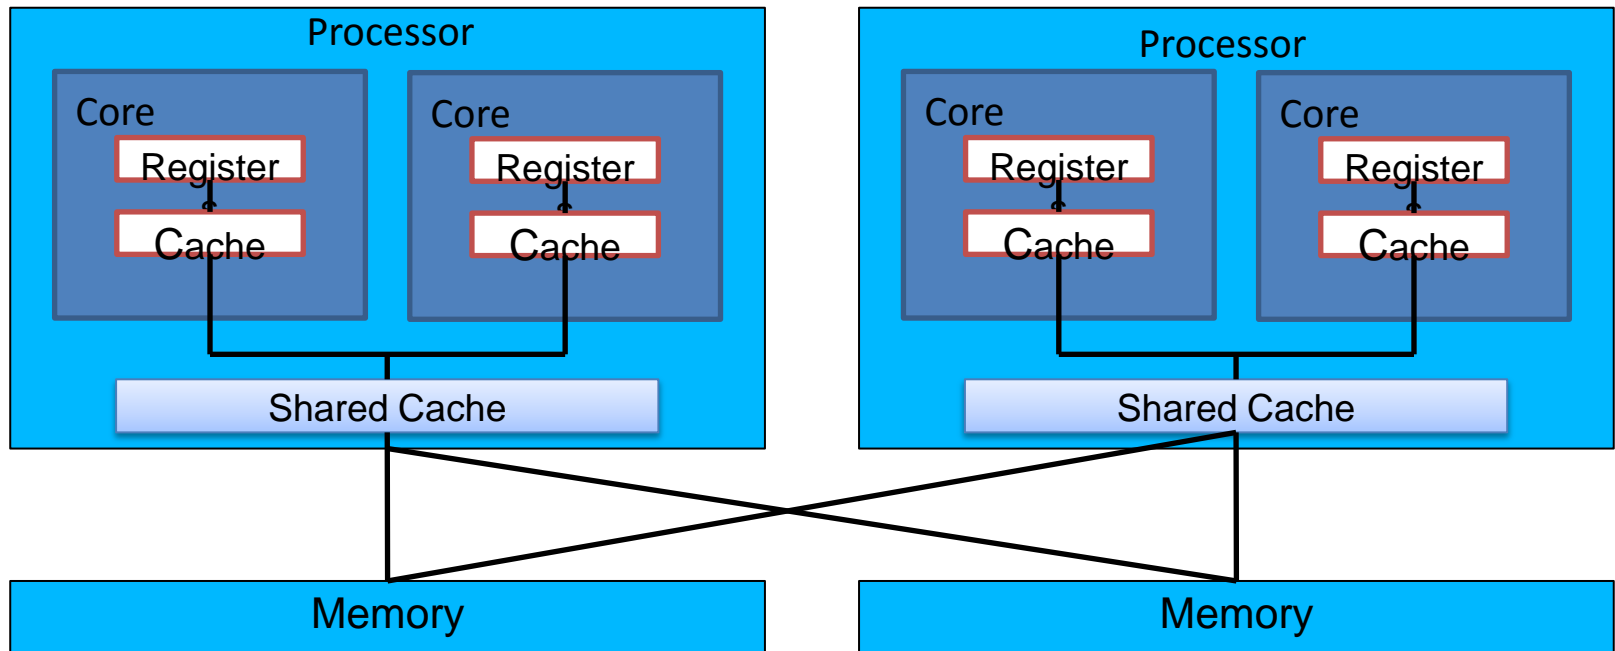
- Time sharing
 - When to schedule which task?

Multicore CPU



- Parallel processing
 - Which tasks to which cores?
 - May have performance implication due to cache contention → contention-aware scheduling


Multiprocessors




- Non-uniform memory access (NUMA) architecture
 - Memory access cost varies significantly: local vs. remote
 - Which tasks to which processors?

Memory Hierarchy

- Main memory
 - DRAM
 - Fast, volatile, expensive
 - CPU has direct access
- Disk
 - Hard disks, solid-state disks
 - Slow, non-volatile, inexpensive
 - CPU doesn't have direct access.



Memory Hierarchy



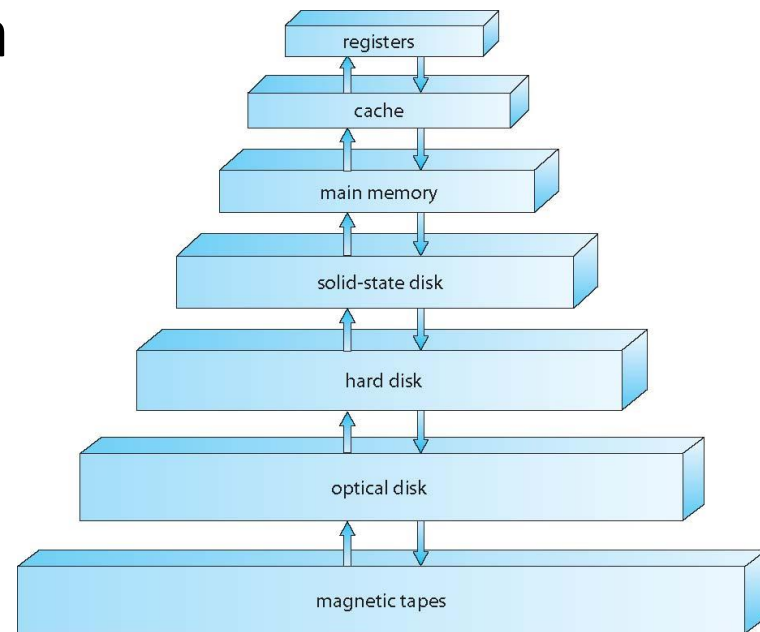
Storage Performance

- Performance of various levels of storage depends on
 - distance from the CPU, size, and process technology used
- Movement between levels of storage hierarchy can be explicit or implicit

Level	1	2	3	4	5
Name	registers	cache	main memory	solid state disk	magnetic disk
Typical size	< 1 KB	< 16MB	< 64GB	< 1 TB	< 10 TB
Implementation technology	custom memory with multiple ports CMOS	on-chip or off-chip CMOS SRAM	CMOS SRAM	flash memory	magnetic disk
Access time (ns)	0.25 - 0.5	0.5 - 25	80 - 250	25,000 - 50,000	5,000,000
Bandwidth (MB/sec)	20,000 - 100,000	5,000 - 10,000	1,000 - 5,000	500	20 - 150
Managed by	compiler	hardware	operating system	operating system	operating system
Backed by	cache	main memory	disk	disk	disk or tape

Caching

- A very important principle applied in all layers of hardware, OS, and software
 - Put frequently accessed data in a small amount of faster memory
 - Fast, most of the time (hit)
 - Copy from slower memory to the cache (miss)



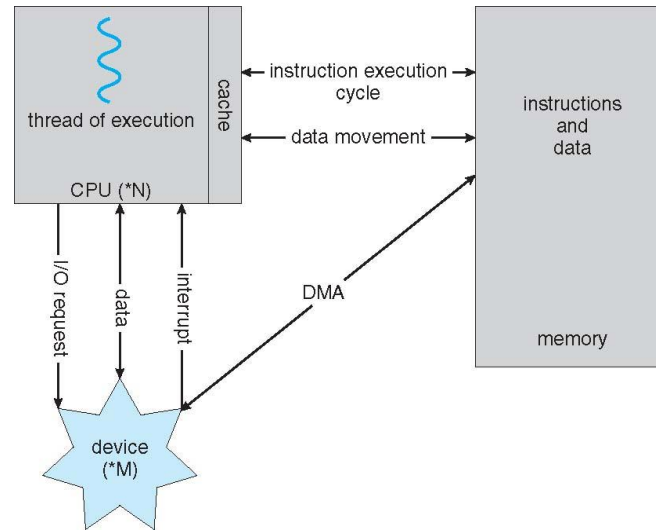
Architectural Support for OS

- Interrupts and exceptions
- Protected modes (kernel/user modes)
- Memory protection and virtual memory
- Synchronization instructions

Interrupt

- What is an interrupt?
 - A signal to the processor telling “do something now!”
- Hardware interrupts
 - Devices (timer, disk, keyboard, ...) to CPU
- Software interrupts (exceptions)
 - Divide by zero, special instructions (e.g., int 0x80)

Interrupt Handling



- save CPU states (registers)
- execute the associated interrupt service routine (ISR)
- restore the CPU states
- return to the interrupted program


Timesharing

- Multiple tasks share the CPU at the same time
 - But there is only one CPU (assume single-core)
 - Want to schedule different task at a regular interval of 10 ms, for example.
- Timer and OS scheduler tick
 - The OS programs a timer to generate an interrupt at every 10 ms.

Dual (User/Kernel) Mode

- Some operations must be restricted to the OS
 - accessing registers in the disk controller
 - updating memory management unit states
 - ...
- User/Kernel mode
 - Hardware support to distinguish app/kernel
 - Privileged instructions are only for kernel mode
 - Applications can enter into kernel mode only via pre-defined **system calls**

User/Kernel Mode Transition



- System calls
 - Programs ask OS services (privileged) via system calls
 - Software interrupt. “int <num>” in Intel x86

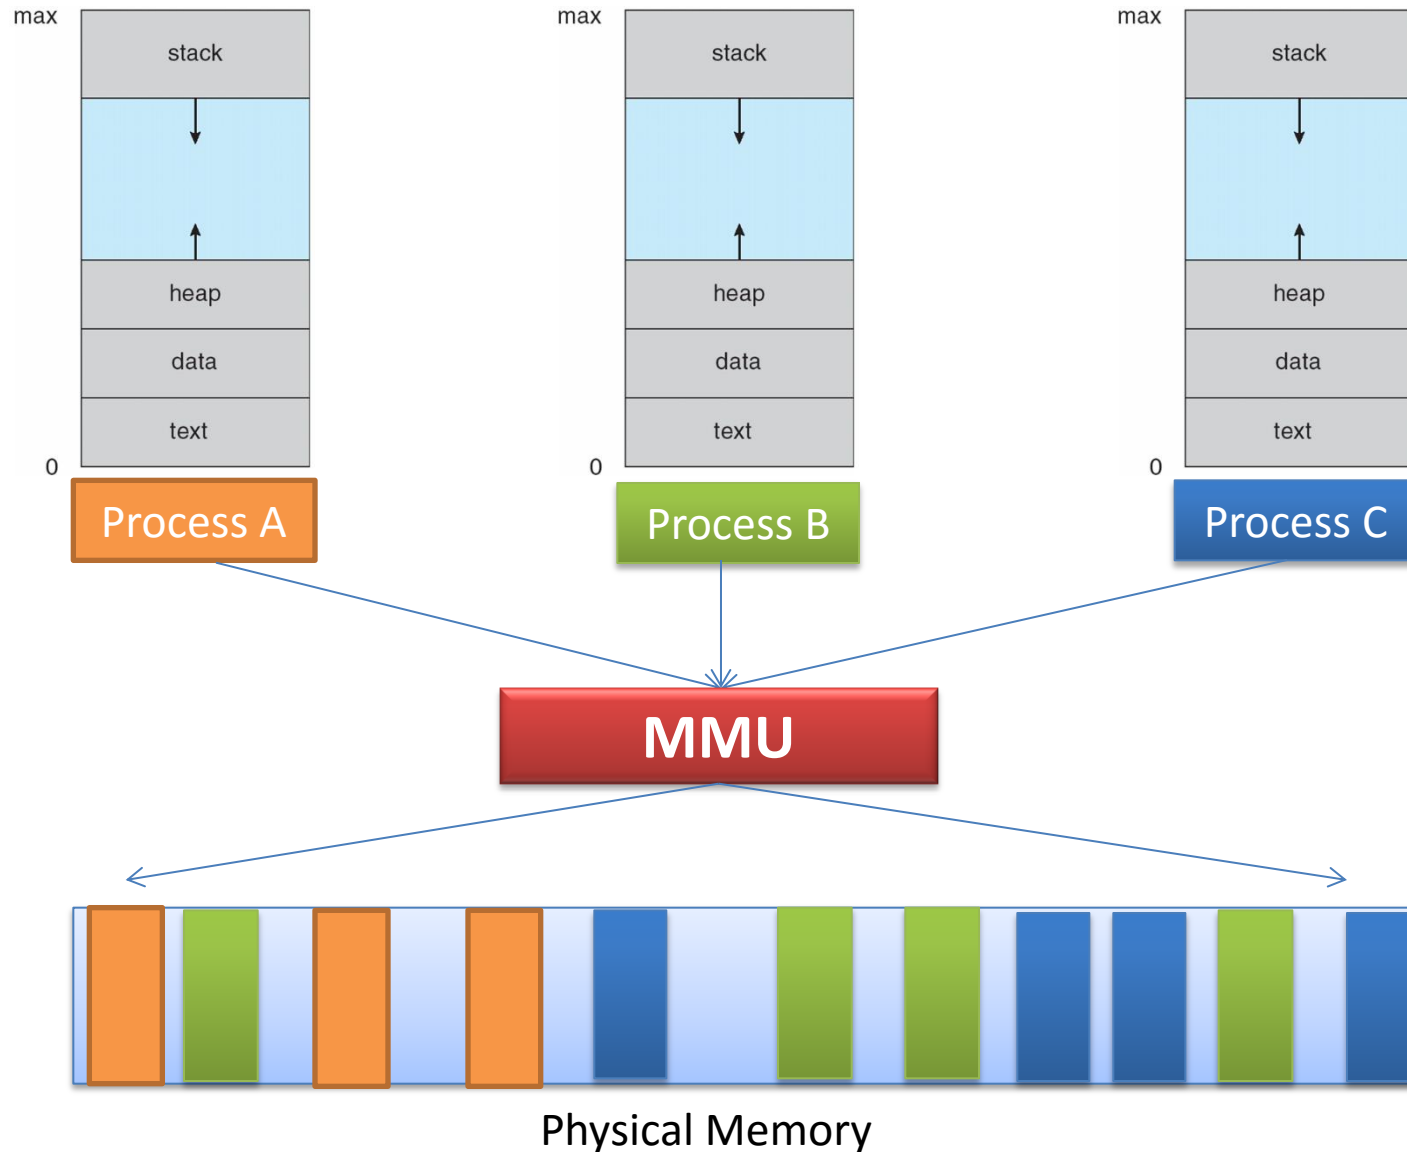
Memory Protection

- How to protect memory among apps/kernel?
 - Applications shouldn't be allowed to access kernel's memory
 - An app shouldn't be able to access another app's memory

Virtual Memory

- How to overcome memory space limitation?
 - Multiple apps must share limited memory space
 - But they want to use memory as if each has dedicated and big memory space
 - E.g.,) 1GB physical memory and 10 programs, each of which wants to have a linear 4GB address space

Virtual Memory



MMU

- Hardware unit that translates *virtual address* to *physical address*
 - Defines the boundaries of kernel/apps
 - Enable efficient use of physical memory



Synchronization

- Synchronization problem with threads

```
Deposit(account, amount) {  
  {  
    account->balance += amount;  
  }  
}
```

Thread 1: *Deposit(acc, 10)*

LOAD R1, account->balance

ADD R1, amount


STORE R1, account->balance

Thread 2: *Deposit(acc, 10)*

LOAD R1, account->balance

ADD R1, amount

STORE R1, account->balance



Synchronization Instructions

- Hardware support for synchronization
 - *TestAndSet*, *CompareAndSwap* instructions
 - Atomic load and store
 - Used to implement **lock** primitives
 - New TSX instruction → hardware transaction
- Another methods to implement locks in single-core systems
 - Disabling interrupts

Summary

- OS needs to understand architecture
 - Hardware (CPU, memory, disk) trends and their implications in OS designs
- Architecture needs to support OS
 - Interrupts and timer
 - User/kernel mode and privileged instructions
 - MMU
 - Synchronization instructions

OS Abstractions

<i>Reality</i>	<i>Abstraction</i>
A single computer	Multiple computers
Limited RAM capacity	Infinite capacity
Mechanical disk	File system
Insecure and unreliable networks	Reliable and secure