

# Process

Heechul Yun

Disclaimer: some slides are adopted from the book authors' slides with permission

# Recap

- OS services
  - Resource (CPU, memory) allocation, filesystem, communication, protection, security, I/O operations
- OS interface
  - System-call interface
- OS structure
  - Monolithic , microkernel
  - Loadable module

# Roadmap


- Beginning of a series of **important** topics:
  - Process
  - Thread
  - Synchronization
- Today
  - Process concept
  - Context switching

# Process

- Process
  - An OS **abstraction** represents a **running** application
- Three main components
  - **Address space**
    - The process's view of memory
    - Includes program code, global variables, dynamic memory, stack
  - **Processor state**
    - Program counter (PC), stack pointer, and other CPU registers
  - **OS resources**
    - Various OS resources that the process uses
    - E.g.) open files, sockets, accounting information

# Process Address Space

- Text
  - Program code
- Data
  - Global variables
- Heap
  - Dynamically allocated memory
    - i.e., Malloc()
- Stack
  - Temporary data
  - Grow at each function call



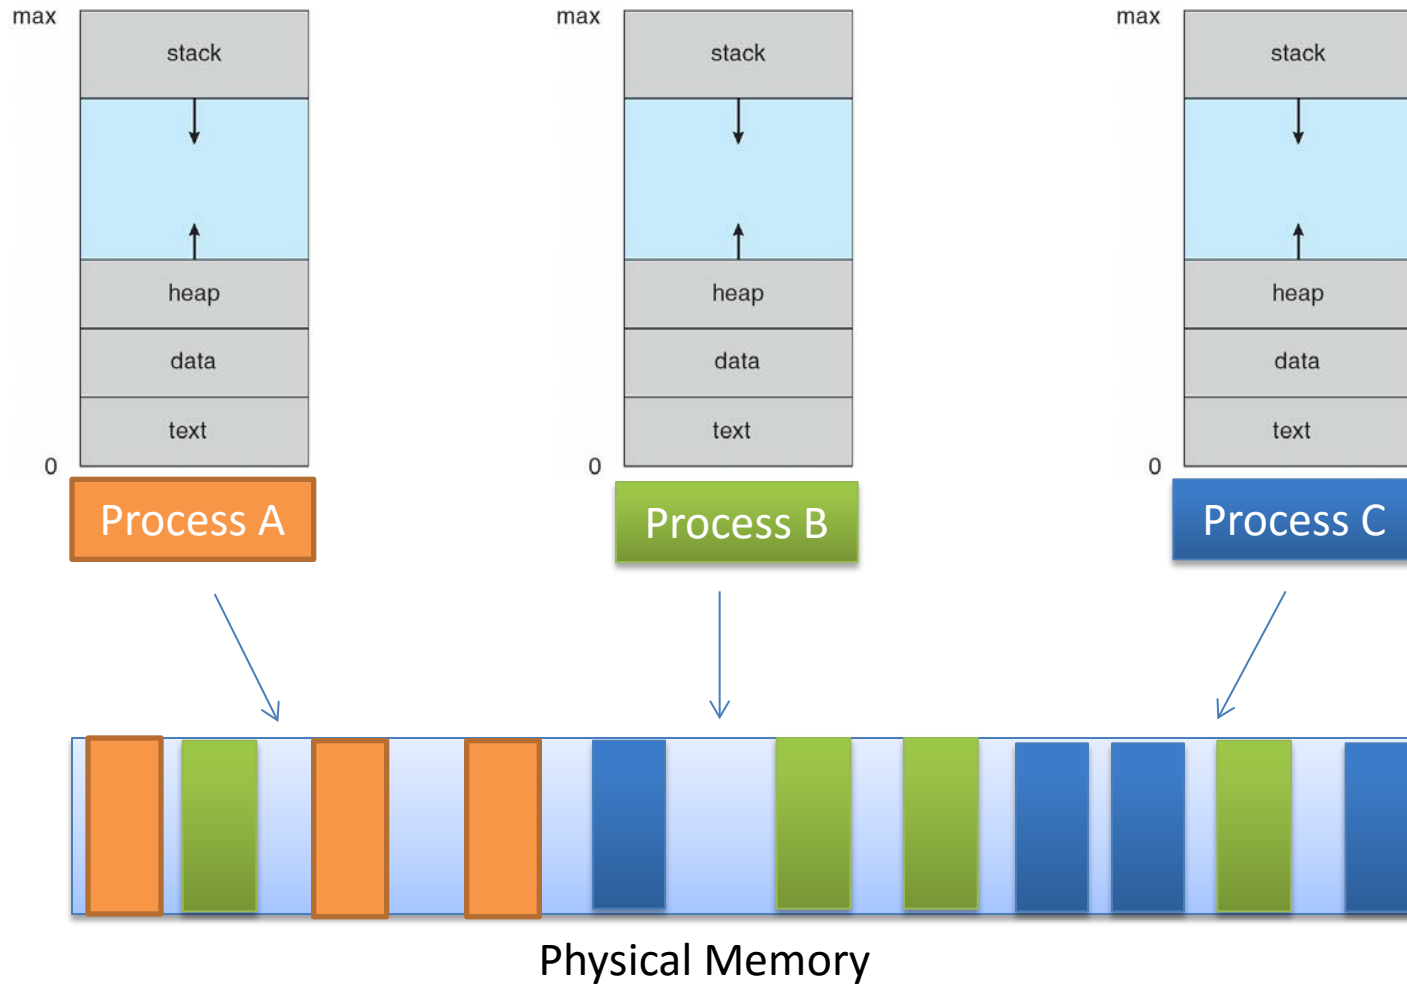


# Process Address Space

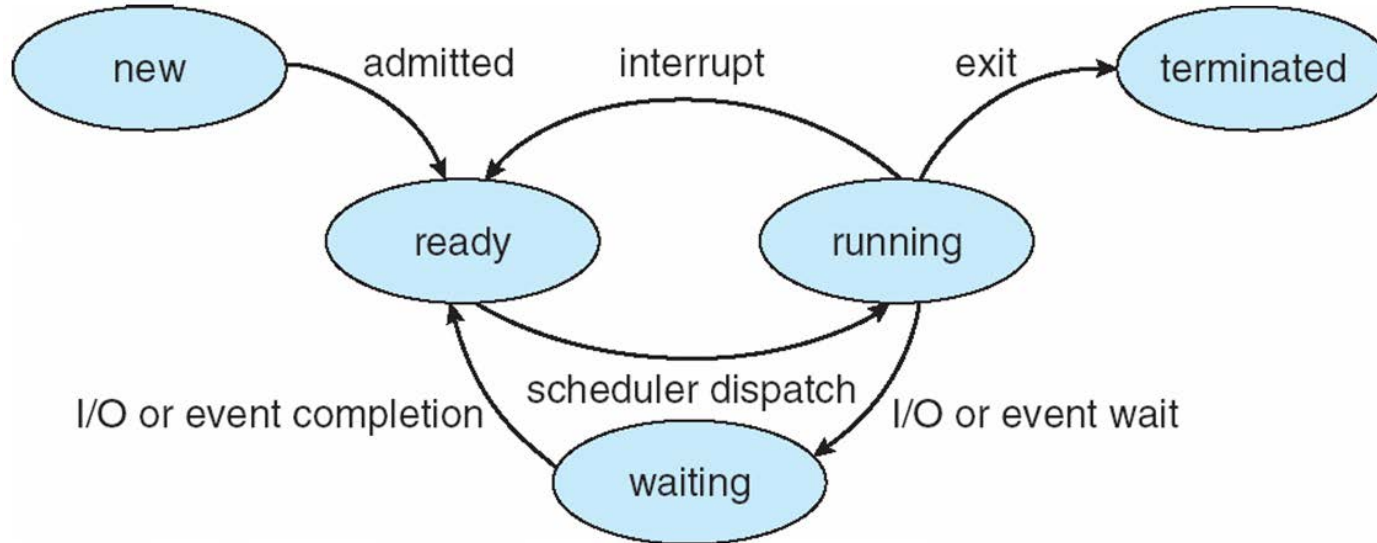
- Each process has its own **private** address space
  - $2^{32}$  (4GB) of **continuous memory** in a 32bit machine
  - Each has same address range (e.g., 0x0 ~ 0xffffffff)
  - How is this possible?
    - What if you have less than 4GB physical DRAM?
    - What if you have 100 processes to run?
- **Virtual memory**
  - An OS mechanism providing this **illusion**
  - We will study it in great detail later in the 2<sup>nd</sup> half of the semester

# Virtual Memory vs. Physical Memory

## Virtual Memory



# Process State



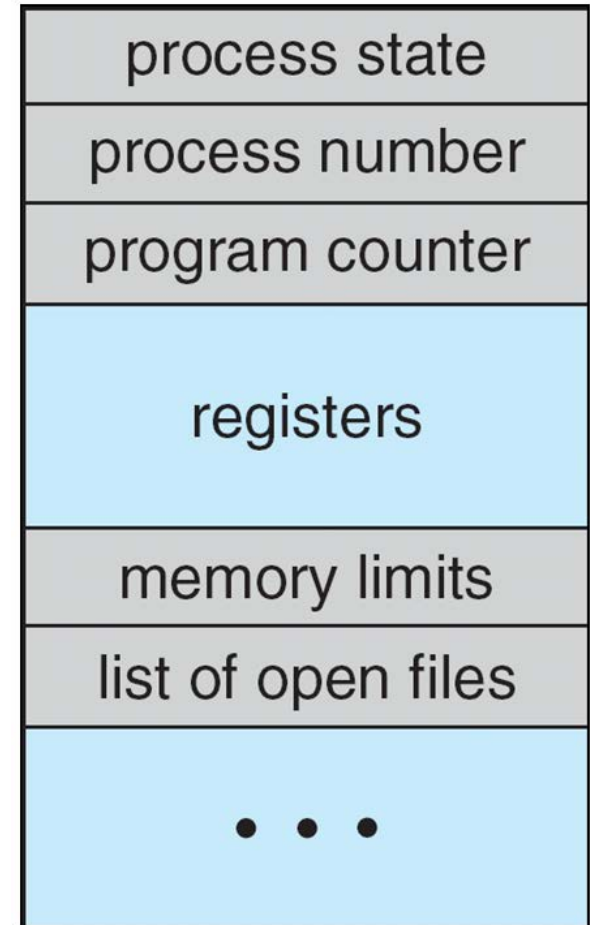
- **running**: Instructions are being executed
- **waiting**: The process is waiting for some event to occur
- **ready**: The process is waiting to be assigned to a processor



# Process Control Block (PCB)

- Information associated with each process

- Process id
- Process state
  - running, waiting, etc.
- Saved CPU registers
  - Register values saved on the last preemption
- CPU scheduling information
  - priorities, scheduling queue pointers
- Memory-management information
  - memory allocated to the process
- Accounting information
  - CPU used, clock time elapsed since start, time limits
- OS resources
  - Open files, sockets, etc.



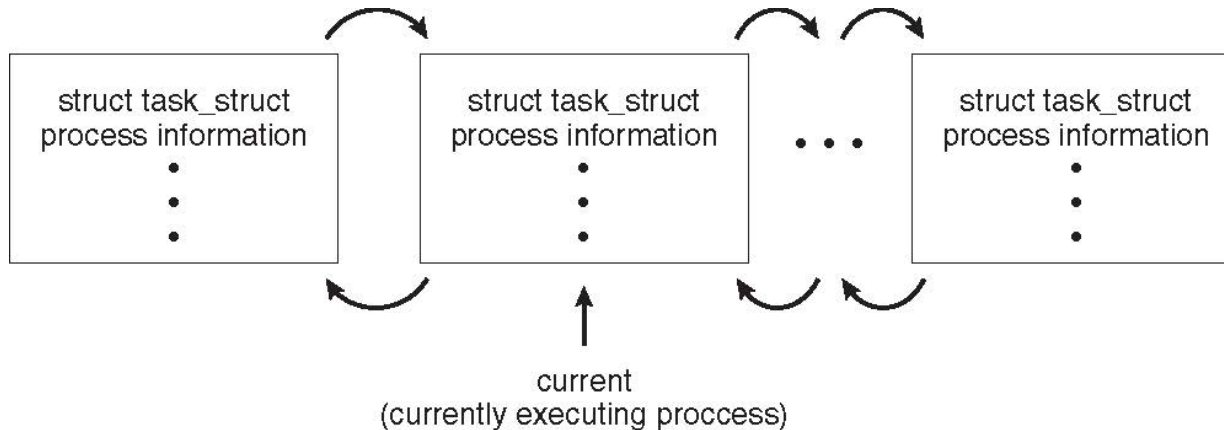
# Process in Linux

Represented by the C structure `task_struct` (`include/linux/sched.h`)

```
pid_t pid;                /* process identifier */
long state;               /* state of the process */
u64 vruntime;             /* CFS scheduling information */
struct files_struct *files; /* list of open files */
struct mm_struct *mm;      /* address space of this process */
cputime_t utime, stime;    /* accounting information */
struct thread_struct thread; /* CPU states */
```

...

(very **big** structure: 5872 bytes in my desktop \*)




(\*) `# cat /sys/kernel/slab/task_struct/object_size`

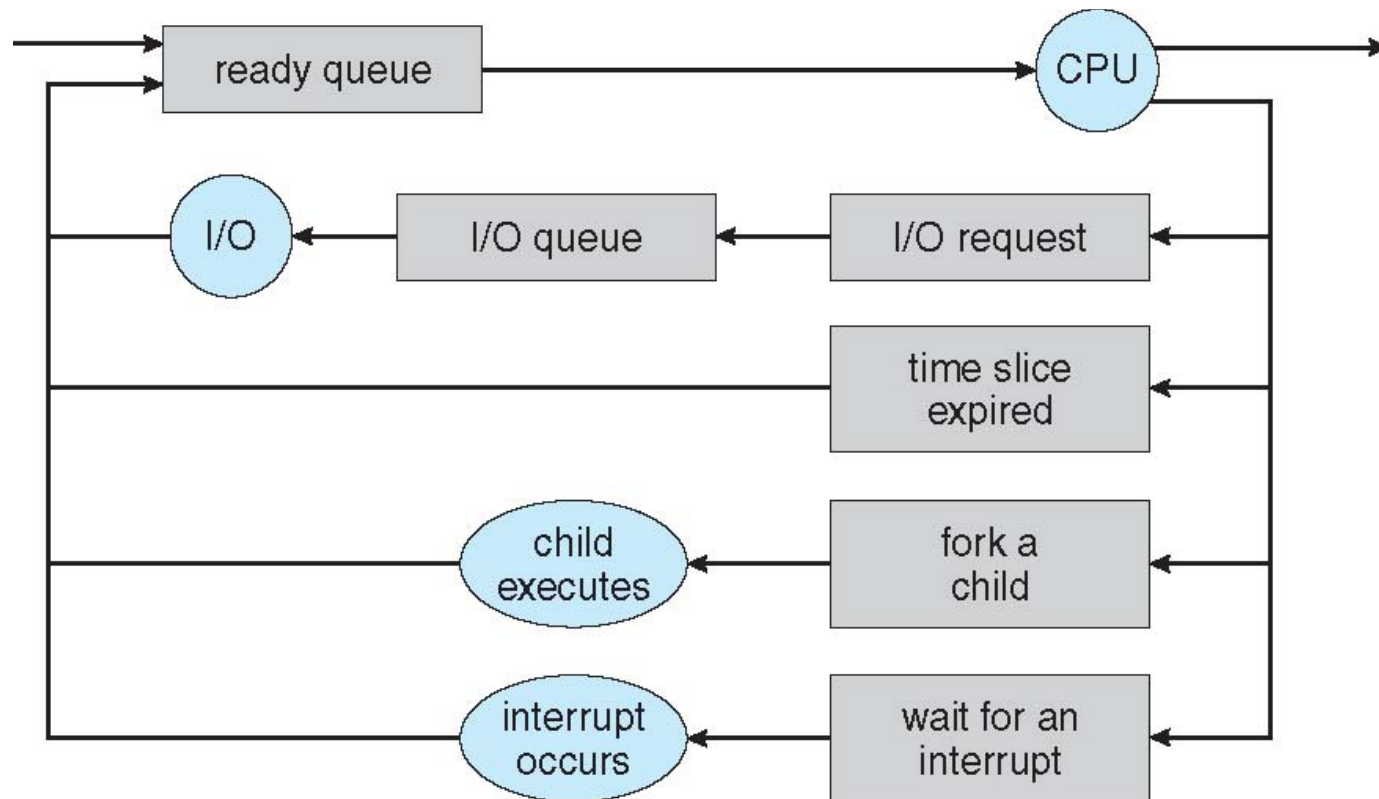
# Process Scheduling

- Decides which process to run **next**
  - Among **ready** processes
- We cover in much more detail later in the class
  - but let's get some **basics**
- OS maintains multiple **scheduling queues**
  - **Ready queue**
    - ready to be executed processes
  - **Device queues**
    - processes waiting for an I/O device
  - Processes **migrate** among the various queues

# Ready Queue and I/O Device Queues

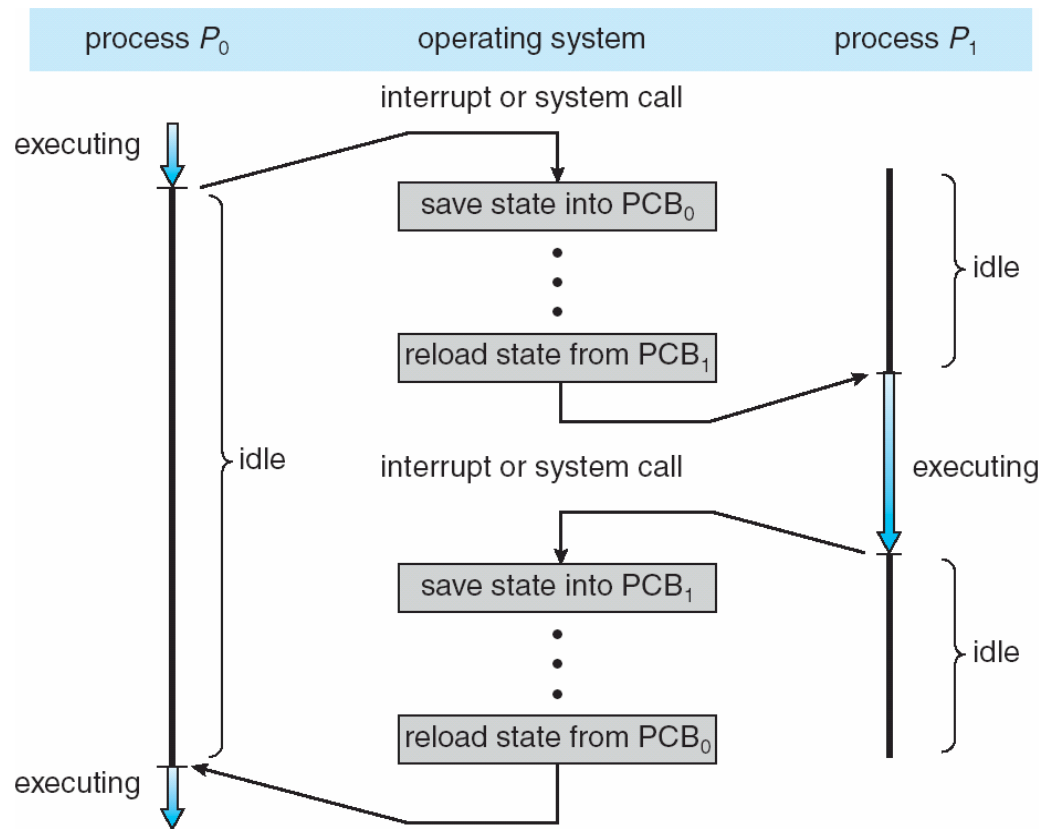


# Process Scheduling: Queuing Representation



# Context Switching

- Suspend the current process and resume a next one from its last suspended state



# Context Switching

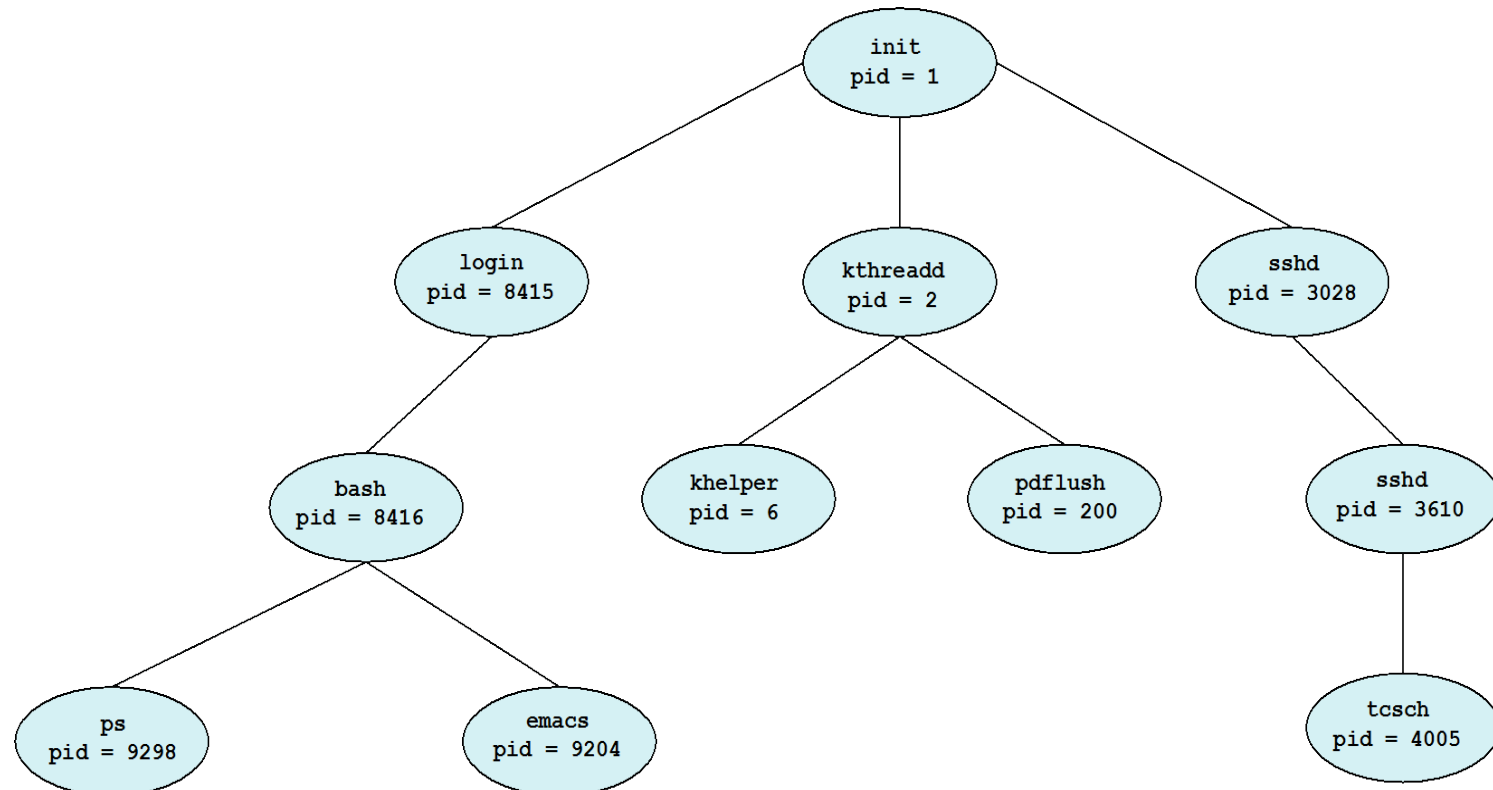
- Overhead
  - Save and restore CPU states
  - Warm up instruction and data cache
    - Cache data of previous process is not useful for new process
- In Linux 3.6.0 on an Intel Xeon 2.8Ghz
  - About 1.8 us
  - ~ 5040 CPU cycles
  - ~ thousands of instructions

# Process Creation

- **Parent** process create **children** processes, which, in turn create other processes, forming a **tree** of processes
- Generally, process identified and managed via a **process identifier (pid)**



# A Process Tree in Linux




```
terminal
root@icecream:/ssd/Dropbox/Teaching/2014-Fall/EECS678# pstree
init--NetworkManager--dhclient
                        dnsmasq
                        2*[{NetworkManager}]
--accounts-daemon---{accounts-daemon}
--acpid
--atd
--automount---3*[{automount}]
--bluetoothd
--colord---2*[{colord}]
--console-kit-dae---64*[{console-kit-dae}]
--cron
--cupsd---dbus
--dbus-daemon
--dropbox---29*[{dropbox}]
--5*[getty]
--irqbalance
--login---bash
--login---bash---nmon
--memcached---5*[{memcached}]
--modem-manager
--nscd---11*[{nscd}]
--nsld---5*[{nsld}]
--php5-fpm---6*[php5-fpm]
--polkitd---{polkitd}
--rpc.idmapd
--rpc.statd
--rpcbind
--rsyslogd---3*[{rsyslogd}]
--rtkit-daemon---2*[{rtkit-daemon}]
--sh---initctl
--sshd---sshd---sshd
                |   |
                |   |--ssh--ssh--bash--emacs--bash
                |   |--ssh--ssh--bash--sudo--bash--pstree
--udev---2*[udev]
--udisks-daemon---udisks-daemon
                  2*[{udisks-daemon}]
--upowerd---2*[{upowerd}]
--upstart-socket-
--upstart-udev-br
--whoopsie---{whoopsie}
--xrdp
--xrdp-sesman
root@icecream:/ssd/Dropbox/Teaching/2014-Fall/EECS678#
```

'pstree' output

# Process Creation

- UNIX examples
  - **fork( )** system call creates new process
  - **exec( )** system call used after a **fork( )** to replace the process' memory space with a new program



# Example: Forking a Process in UNIX

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```

Child

Parent

# Example: Forking a Process in Windows

```
#include <stdio.h>
#include <windows.h>

int main(VOID)
{
    STARTUPINFO si;
    PROCESS_INFORMATION pi;

    /* allocate memory */
    ZeroMemory(&si, sizeof(si));
    si.cb = sizeof(si);
    ZeroMemory(&pi, sizeof(pi));

    /* create child process */
    if (!CreateProcess(NULL, /* use command line */
        "C:\\WINDOWS\\system32\\mspaint.exe", /* command */
        NULL, /* don't inherit process handle */
        NULL, /* don't inherit thread handle */
        FALSE, /* disable handle inheritance */
        0, /* no creation flags */
        NULL, /* use parent's environment block */
        NULL, /* use parent's existing directory */
        &si,
        &pi))
    {
        fprintf(stderr, "Create Process Failed");
        return -1;
    }
    /* parent will wait for the child to complete */
    WaitForSingleObject(pi.hProcess, INFINITE);
    printf("Child Complete");

    /* close handles */
    CloseHandle(pi.hProcess);
    CloseHandle(pi.hThread);
}
```

# Process Termination

- Normal termination via **exit( )** system call.
  - Exit by itself.
  - Returns status data from child to parent (via **wait( )**)
  - Process's resources are deallocated by operating system
- Forced termination via **kill( )** system call
  - Kill someone else (child)
- **Zombie** process
  - If no parent waiting (did not invoke **wait( )**)
- **Orphan** process
  - If parent terminated without invoking **wait**
  - **Q: who will be the parent of a orphan process?**
  - **A: Init process**