

# Recap: Thread

- What is it?
  - Independent flow of control
- What does it need (thread private)?
  - Stack
- What for?
  - Lightweight programming construct for concurrent activities
- How to implement?
  - Kernel thread vs. user thread

# Recap: Process vs. Thread

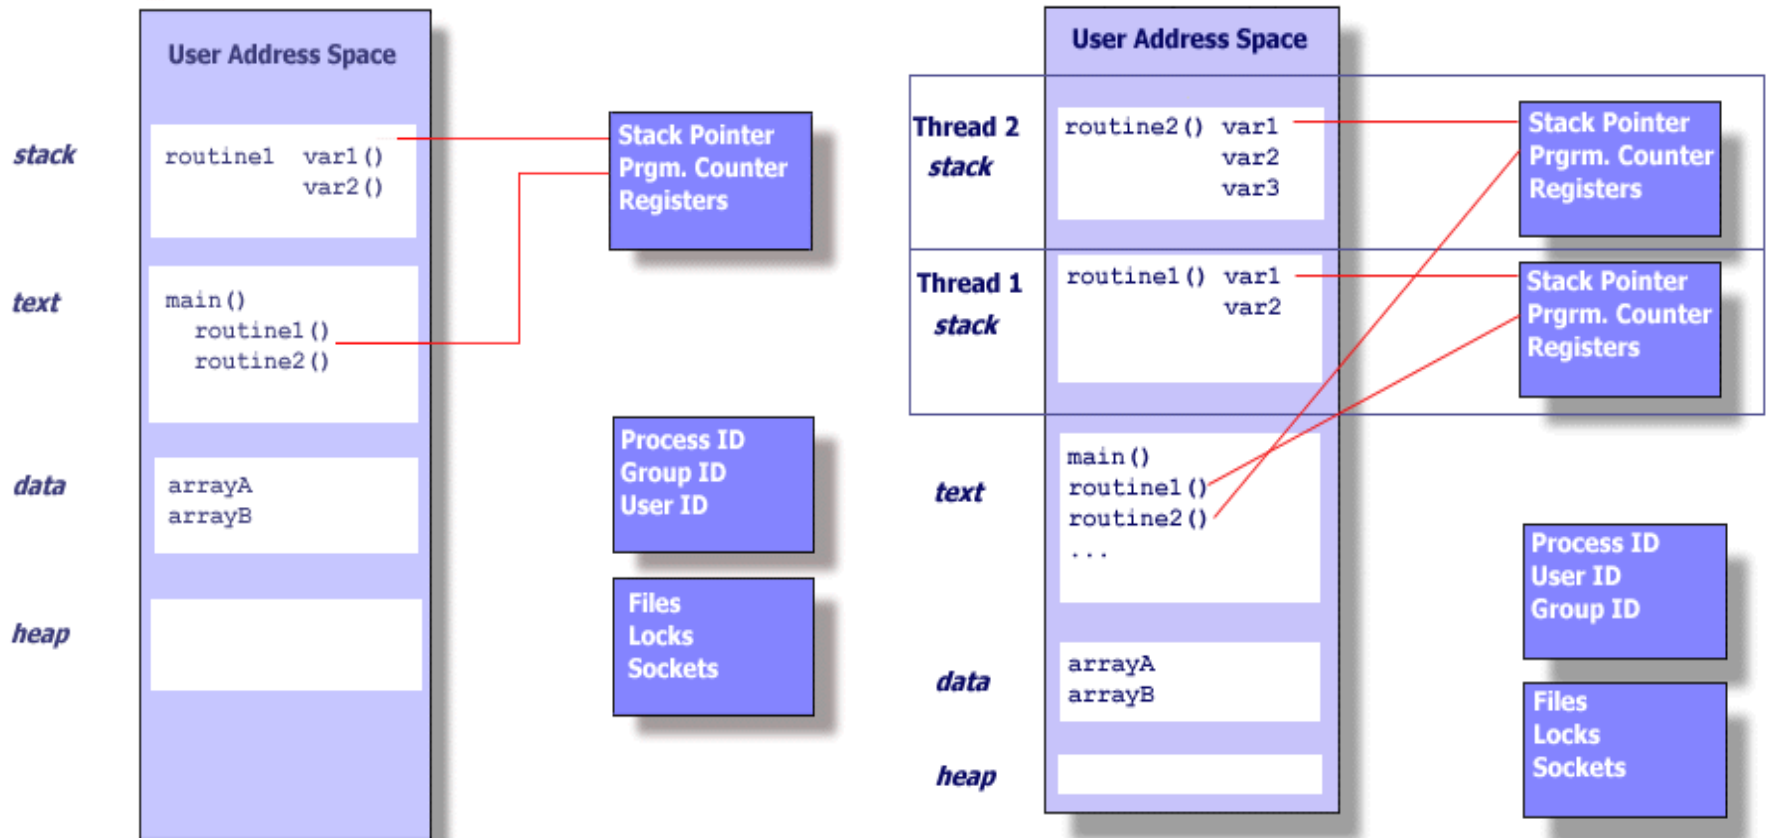


Figure source: <https://computing.llnl.gov/tutorials/pthreads/>

# Multi-threads vs. Multi-processes

- Multi-processes
  - (+) protection
  - (-) performance (?)
- Multi-threads
  - (+) performance
  - (-) protection



Process-per-tab



Single-process multi-threads

# Threads: Advanced Topics

- Semantics of Fork/exec()
- Signal handling
- Thread pool
- Multicore

# Semantics of fork()/exec()

- Remember fork(), exec() system calls?
  - Fork: create a child process (a copy of the parent)
  - Exec: replace the address space with a new pgm.
- Duplicate *all* threads or *the caller* only?
  - Linux: the calling thread only
  - Complicated. [Don't do it!](#)
    - Why? Mutex states, library, ...
    - Exec() immediately after Fork() may be okay.


# Signal Handling

- What is *Singal*?
  - \$ man 7 signal
  - OS to process notification
    - “hey, wake-up, you’ve got a packet on your socket,”
    - “hey, wake-up, your timer is just expired.”
- Which *thread* to deliver a signal?
  - Any thread
    - e.g., kill(pid)
  - Specific thread
    - E.g., pthread\_kill(tid)


# Thread Pool

- Managing threads yourself can be cumbersome and costly
  - Repeat: create/destroy threads as needed.
- Let's create a set of threads ahead of time, and just ask them to execute my functions
  - #of thread  $\sim$  #of cores
  - No need to create/destroy many times
  - Many high-level parallel libraries use this.
    - e.g., Intel TBB (threading building block), ...

# Single Core Vs. Multicore Execution



*Single core execution*



*Multiple core execution*



# Challenges for Multithreaded Programming in Multicore

- How to divide activities?
- How to divide data?
- **How to synchronize accesses to the shared data?**
- How to test and debug?

EECS750

# Summary

- Thread
  - What is it?
    - Independent flow of control.
  - What for?
    - Lightweight programming construct for concurrent activities
  - How to implement?
    - Kernel thread vs. user thread
- Next class
  - How to synchronize?

# Synchronization

Disclaimer: some slides are adopted from the book authors' slides with permission

# Agenda

- **Mutual exclusion**
  - **Peterson's algorithm (Software)**
  - **Synchronization instructions (Hardware)**
- **High-level synchronization mechanisms**
  - **Mutex**
  - **Semaphore**
  - **Monitor**

# Producer/Consumer



# Producer/Consumer

## Producer

```
while (true){  
  
    /* wait if buffer full */  
    while (counter == 10);  
  
    /* produce data */  
    buffer[in] = sdata;  
    in = (in + 1) % 10;  
  
    /* update number of  
       items in buffer */  
    counter++;  
}
```

## Consumer

```
while (true){  
  
    /* wait if buffer empty */  
    while (counter == 0);  
  
    /* consume data */  
    sdata = buffer[out];  
    out = (out + 1) % 10;  
  
    /* update number of  
       items in buffer */  
    counter--;  
}
```

# Producer/Consumer

## Producer

```
while (true){  
  
    /* wait if buffer full */  
    while (counter == 10);  
  
    /* produce data */  
    buffer[in] = sdata;  
    in = (in + 1) % 10;  
  
    /* update number of  
       items in buffer */  
    R1 = load (counter);  
    R1 = R1 + 1;  
    counter = store (R1);  
}
```

## Consumer

```
while (true){  
  
    /* wait if buffer empty */  
    while (counter == 0);  
  
    /* consume data */  
    sdata = buffer[out];  
    out = (out + 1) % 10;  
  
    /* update number of  
       items in buffer */  
    R2 = load (counter);  
    R2 = R2 - 1;  
    counter = store (R2);  
}
```

# Check Yourself

```
int count = 0;
int main()
{
    count = count + 1;
    return count;
}
```

```
$ gcc -O2 -S sync.c
```

```
...
movl    count(%rip), %eax
addl    $1, %eax
movl    %eax, count(%rip)
...
```



# Race Condition

Initial condition: *counter* = 5

Thread 1

```
R1 = load (counter);  
R1 = R1 + 1;  
counter = store (R1);
```


Thread 2

```
R2 = load (counter);  
R2 = R2 - 1;  
counter = store (R2);
```

- What are the possible outcome?

# Race Condition

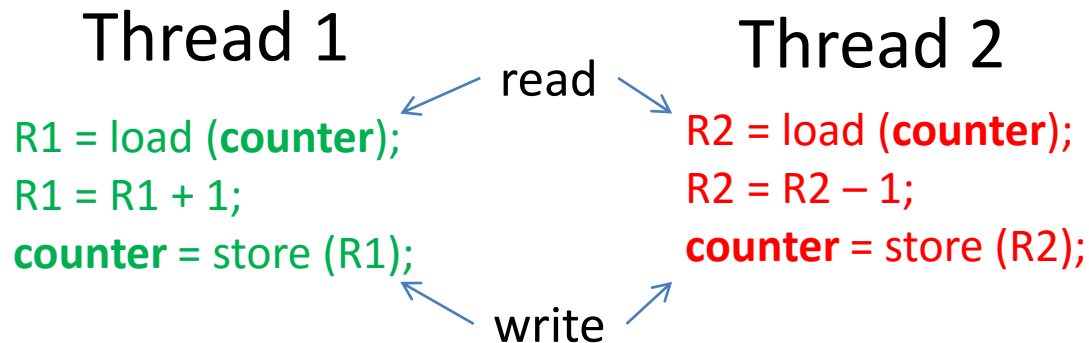
Initial condition: *counter* = 5



- Why this happens?

# Race Condition

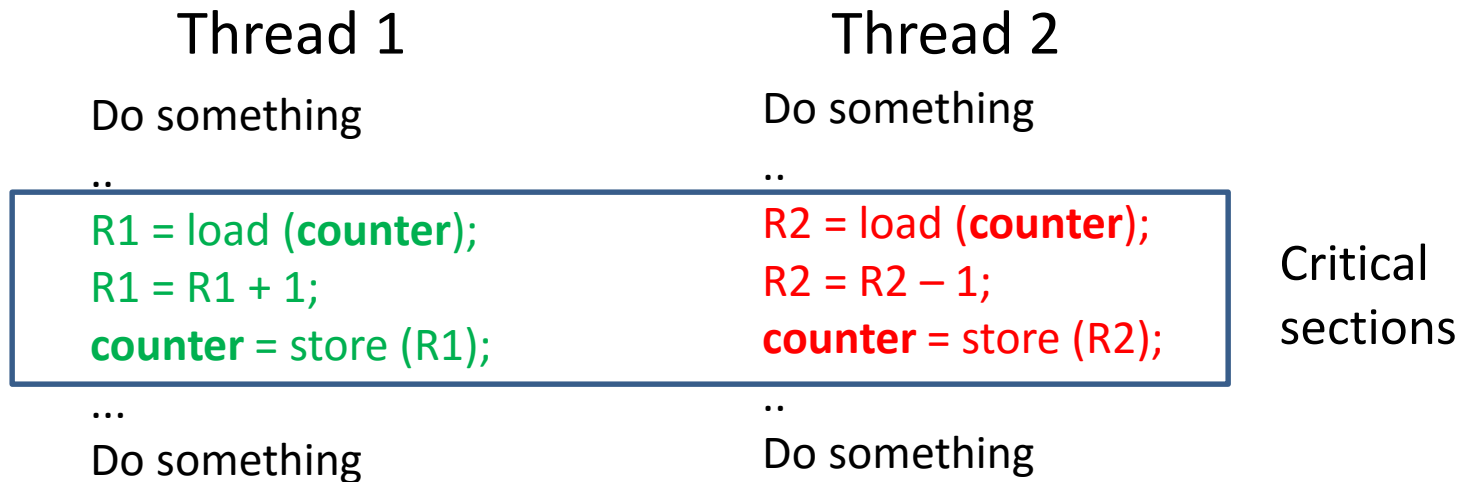
- A situation when two or more threads **read and write** shared data at the same time
- Correctness depends on the execution order



- How to prevent race conditions?

# Critical Section

- Code sections of potential race conditions



# Solution Requirements

- Mutual Exclusion
  - If a thread executes its critical section, *no other threads* can enter their critical sections
- Progress
  - If no one executes a critical section, someone can enter its critical section
- Bounded waiting
  - Waiting (time/number) must be bounded

# Simple Solution (?): Use a Flag

```
// wait
while (in_cs)
    ;
// enter critical section
in_cs = true;

Do something

// exit critical section
in_cs = false;
```

**T1**  
while(in\_cs);  
  
in\_cs = true;  
//enter

**T2**  
  
while(in\_cs);  
in\_cs = true;  
  
//enter

- Mutual exclusion is not guaranteed

# Peterson's Solution

- Software solution (no h/w support)
- Two process solution
  - Multi-process extension exists
- The two processes share two variables:
  - int turn;
    - The variable turn indicates whose turn it is to enter the critical section
  - Boolean flag[2]
    - The flag array is used to indicate if a process is ready to enter the critical section.

# Peterson's Solution

## Thread 1

```
do {  
    flag[0] = TRUE;  
    turn = 1;  
    while (flag[1] && turn==1)  
        ;  
    // critical section  
  
    flag[0] = FALSE;  
  
    // remainder section  
} while (TRUE)
```

## Thread 2

```
do {  
    flag[1] = TRUE;  
    turn = 0;  
    while (flag[0] && turn==0)  
        ;  
    // critical section  
  
    flag[1] = FALSE;  
  
    // remainder section  
} while (TRUE)
```

- Solution meets all three requirements
  - Mutual exclusion: P0 and P1 cannot be in the critical section at the same time
  - Progress: if P0 does not want to enter critical region, P1 does no waiting
  - Bounded waiting: process waits for at most one turn



# Peterson's Solution

- Only supports two processes
  - generalizing for more than two processes has been achieved, but not very efficient
- Assumes that the LOAD and STORE instructions are atomic
- Assumes that memory accesses are not reordered
  - your compiler re-orders instructions (gcc -O2, -O3, ...)
  - your processor re-orders instructions (memory consistency models )

# Reordering by the CPU

*Initially  $X = Y = 0$*

Thread 0   Thread 1

$X = 1$     $Y = 1$   
 $R1 = Y$     $R2 = X$

Thread 0   Thread 1

$R1 = Y$   
 $X = 1$   
 $R2 = X$   
 $Y = 1$

- Possible values of R1 and R2?
  - 0,1
  - 1,0
  - 1,1
  - **0,0 ← possible on PC**

# Lock-Based Solutions

- General solution to the critical section problem
  - critical sections are protected by locks
  - process must acquire lock before entry
  - process releases lock on exit

```
do {  
    acquire lock;  
  
    critical section  
  
    release lock;  
  
    remainder section  
} while(TRUE);
```

# Hardware Support for Lock-Based Solutions

## – Uniprocessors

- For uniprocessor systems
  - concurrent processes cannot be overlapped, only *interleaved*
  - process runs until it is *interrupted*
- Disable interrupts !
  - active process will run without preemption

do {

    disable interrupts;  
        *critical section*  
    enable interrupts;

        remainder section  
} while(TRUE);

# Hardware Support for Lock-Based Solutions

## – Multiprocessors

- In multiprocessors
  - several processes share memory
  - processors behave independently in a peer manner
- Disabling interrupt based solution will not work
  - too inefficient
  - OS using this not broadly scalable
- Provide hardware support in the form of **atomic** instructions
  - atomic *test-and-set* instruction
  - atomic *compare-and-swap* instruction
- Atomic execution of a set of instructions means that the instructions are treated as a single step that cannot be interrupted.
  - **All or nothing** property

# *TestAndSet* Instruction

- Pseudo code definition of *TestAndSet*

```
boolean TestAndSet (boolean *target)
{
    boolean rv = *target;
    *target = TRUE;
    return rv;
}
```

# Mutual Exclusion using *TestAndSet*

```
int mutex;  
init_lock (&mutex);  
  
do {  
    lock (&mutex);  
    critical section  
    unlock (&mutex);  
    remainder section  
} while(TRUE);
```

```
void init_lock (int *mutex)  
{  
    *mutex = 0;  
}  
  
void lock (int *mutex)  
{  
    while(TestAndSet(mutex))  
        ;  
}  
  
void unlock (int *mutex)  
{  
    *mutex = 0;  
}
```

# *CAS (Compare & Swap)* Instruction

## ■ Psuedo code definition of CAS instruction

```
int CAS(int *value, int oldval, int newval)
{
    int temp = *value;
    if (*value == oldval)
        *value = newval;
    return temp;
}
```



# Mutual Exclusion using CAS

```
int mutex;  
init_lock (&mutex);  
  
do {  
  
    lock (&mutex);  
        critical section  
    unlock (&mutex);  
  
        remainder section  
} while(TRUE);
```

```
void init_lock (int *mutex) {  
    *mutex = 0;  
}  
  
void lock (int *mutex) {  
    while(CAS(&mutex, 0, 1) != 0);  
}  
  
void unlock (int *mutex) {  
    *mutex = 0;  
}
```

*Fairness not guaranteed by any implementation !*

# Preemptive Vs. Non-preemptive Kernels

- Kernel is full of important shared data
  - structures for maintaining file systems, memory allocation, interrupt handling, etc.
- How to ensure the OS is free from race conditions ?
- Non-preemptive kernels
  - process executing in kernel mode cannot be preempted
  - disable interrupts when process is in kernel mode
  - what about multiprocessor systems ?
- Preemptive kernels
  - process executing in kernel mode can be preempted
  - suitable for real-time programming
  - more responsive

# Roadmap

- Solutions for mutual exclusion
  - Peterson's algorithm (Software)
  - Synchronization instructions (Hardware)
- **High-level synchronization mechanisms**
  - **Mutex**
  - **Semaphore**
  - **Monitor**

# Spinlock using TestAndSet

```
void init_lock (int *mutex)
{
    *mutex = 0;
}

void lock (int *mutex)
{
    while(TestAndSet(mutex))
        ;
}

void unlock (int *mutex)
{
    *mutex = 0;
}
```

# What's Wrong With **Spinlocks**?

- Very wasteful
  - Waiting thread continues to use CPU cycles
  - While doing absolutely nothing but wait
  - 100% CPU utilization, but no useful work done
  - Power consumption, fan noise, ...
- Useful when
  - You hold the lock only *briefly*
- Otherwise
  - A better solution is needed

# Mutex – Blocking Lock

- Instead of spinning
  - Let the thread sleep
    - There can be multiple waiting threads
  - In the meantime, let other threads use the CPU
  - When the lock is released, wake-up one thread
    - Pick one if there multiple threads were waiting

```
void mutex_init (mutex_t *lock)
```

```
{
```

```
    lock->value = 0;
```

```
    list_init(&lock->wait_list);
```

```
    spin_lock_init(&lock->wait_lock);
```

```
}
```

← Thread waiting list

← To protect waiting list

```
void mutex_lock (mutex_t *lock)
```

```
{
```

```
...
```

```
    while(TestAndSet(&lock->value)) {
```

```
...
```

```
...
```

```
...
```

```
...
```

```
...
```

```
    }
```

```
...
```

```
}
```

```
void mutex_unlock (mutex_t *lock)
```

```
{
```

```
...
```

```
    lock->value = 0;
```

```
...
```

```
...
```

```
...
```

```
}
```

```
void mutex_init (mutex_t *lock)
```

```
{
```

```
    lock->value = 0;
```

```
    list_init(&lock->wait_list);
```

```
    spin_lock_init(&lock->wait_lock);
```

```
}
```

← Thread waiting list

← To protect waiting list

```
void mutex_lock (mutex_t *lock)
```

```
{
```

```
...
```

```
    while(TestAndSet(&lock->value)) {
```

```
        current->state = WAITING;
```

```
        list_add(&lock->wait_list, current);
```

```
...
```

```
        schedule();
```

```
...
```

```
    }
```

```
...
```

```
}
```

← Thread state change

← Add the current thread to the waiting list

← Sleep or schedule another thread

```
void mutex_unlock (mutex_t *lock)
```

```
{
```

```
...
```

```
    lock->value = 0;
```

```
...
```

```
...
```

```
...
```

```
}
```



```
void mutex_init (mutex_t *lock)
```

```
{
```

```
    lock->value = 0;
```

```
    list_init(&lock->wait_list);
```

```
    spin_lock_init(&lock->wait_lock);
```

```
}
```

← Thread waiting list

← To protect waiting list

```
void mutex_lock (mutex_t *lock)
```

```
{
```

```
    spin_lock(&lock->wait_lock);
```

```
    while(TestAndSet(&lock->value)) {
```

```
        current->state = WAITING;
```

```
        list_add(&lock->wait_list, current);
```

```
    ...
```

```
        schedule();
```

```
    ...
```

```
    }
```

```
    spin_unlock(&lock->wait_lock);
```

```
}
```

← Thread state change

← Add the current thread to the waiting list

← Sleep or schedule another thread

```
void mutex_unlock (mutex_t *lock)
```

```
{
```

```
...
```

```
    lock->value = 0;
```

```
...
```

```
...
```

```
...
```

```
}
```

```
void mutex_init (mutex_t *lock)
```

```
{
```

```
    lock->value = 0;
```

```
    list_init(&lock->wait_list);
```

```
    spin_lock_init(&lock->wait_lock);
```

```
}
```

← Thread waiting list

← To protect waiting list

```
void mutex_lock (mutex_t *lock)
```

```
{
```

```
    spin_lock(&lock->wait_lock);
```

```
    while(TestAndSet(&lock->value)) {
```

```
        current->state = WAITING;
```

```
        list_add(&lock->wait_list, current);
```

```
        spin_unlock(&lock->wait_lock);
```

```
        schedule();
```

```
        spin_lock(&lock->wait_lock);
```

```
    }
```

```
    spin_unlock(&lock->wait_lock);
```

```
}
```

← Thread state change

← Add the current thread to the waiting list

← Sleep or schedule another thread

```
void mutex_unlock (mutex_t *lock)
```

```
{
```

```
...
```

```
    lock->value = 0;
```

```
...
```

```
...
```

```
...
```

```
}
```

```
void mutex_init (mutex_t *lock)
```

```
{
```

```
    lock->value = 0;
```

```
    list_init(&lock->wait_list);
```

```
    spin_lock_init(&lock->wait_lock);
```

```
}
```

← Thread waiting list

← To protect waiting list

```
void mutex_lock (mutex_t *lock)
```

```
{
```

```
    spin_lock(&lock->wait_lock);
```

```
    while(TestAndSet(&lock->value)) {
```

```
        current->state = WAITING;
```

```
        list_add(&lock->wait_list, current);
```

```
        spin_unlock(&lock->wait_lock);
```

```
        schedule();
```

```
        spin_lock(&lock->wait_lock);
```

```
    }
```

```
    spin_unlock(&lock->wait_lock);
```

```
}
```

← Thread state change

← Add the current thread to the waiting list

← Sleep or schedule another thread

```
void mutex_unlock (mutex_t *lock)
```

```
{
```

```
...
```

```
    lock->value = 0;
```

```
    if (!list_empty(&lock->wait_list))
```

```
        wake_up_process(&lock->wait_list)
```

← Someone is waiting for the lock

← Wake-up a waiting thread

```
...
```

```
}
```

```
void mutex_init (mutex_t *lock)
```

```
{
```

```
    lock->value = 0;
```

```
    list_init(&lock->wait_list);
```

```
    spin_lock_init(&lock->wait_lock);
```

```
}
```

← Thread waiting list

← To protect waiting list

```
void mutex_lock (mutex_t *lock)
```

```
{
```

```
    spin_lock(&lock->wait_lock);
```

```
    while(TestAndSet(&lock->value)) {
```

```
        current->state = WAITING;
```

```
        list_add(&lock->wait_list, current);
```

```
        spin_unlock(&lock->wait_lock);
```

```
        schedule();
```

```
        spin_lock(&lock->wait_lock);
```

```
    }
```

```
    spin_unlock(&lock->wait_lock);
```

```
}
```

← Thread state change

← Add the current thread to the waiting list

← Sleep or schedule another thread

```
void mutex_unlock (mutex_t *lock)
```

```
{
```

```
    spin_lock(&lock->wait_lock);
```

```
    lock->value = 0;
```

```
    if (!list_empty(&lock->wait_list))
```

```
        wake_up_process(&lock->wait_list)
```

```
    spin_unlock(&lock->wait_lock);
```

```
}
```

← Someone is waiting for the lock

← Wake-up a waiting thread

# High-level Synchronization Primitives

- Lock (mutex) is great, but...
  - Too low-level primitive
  - Sometimes we need more **powerful primitives**
- Semaphore
  - Binary/integer semaphore
- Monitor
  - Condition variable

# Semaphore

- High-level synchronization primitive
  - Designed by Dijkstra in 1960'
- Definition
  - Semaphore is an integer variable
  - Only two operations are possible:
    - P() or wait() or down()
      - **Wait** until the semaphore value to become  $> 0$ , then **decrements** it by 1.
    - V() or signal() or up()
      - **Increments** the semaphore value by 1