

Recap

- Mutual exclusion
- Peterson's algorithm
 - s/w solution
 - Assume program order execution
- Synchronization instructions
 - Test&set

TestAndSet Instruction

- Pseudo code

```
boolean TestAndSet (boolean *target)
{
    boolean rv = *target;
    *target = TRUE;
    return rv;
}
```

Mutual Exclusion using *TestAndSet*

```
int mutex;  
init_lock (&mutex);  
  
do {  
    lock (&mutex);  
    critical section  
    unlock (&mutex);  
    remainder section  
} while(TRUE);
```

```
void init_lock (int *mutex)  
{  
    *mutex = 0;  
}  
  
void lock (int *mutex)  
{  
    while(TestAndSet(mutex))  
        ;  
}  
  
void unlock (int *mutex)  
{  
    *mutex = 0;  
}
```

CAS (Compare & Swap) Instruction

- Pseudo code

```
int CAS(int *value, int oldval, int newval)
{
    int temp = *value;
    if (*value == oldval)
        *value = newval;
    return temp;
}
```

Mutual Exclusion using CAS

```
int mutex;  
init_lock (&mutex);  
  
do {  
    lock (&mutex);  
    critical section  
    unlock (&mutex);  
    remainder section  
} while(TRUE);
```

```
void init_lock (int *mutex) {  
    *mutex = 0;  
}  
  
void lock (int *mutex) {  
    while(CAS(&mutex, 0, 1) != 0);  
}  
  
void unlock (int *mutex) {  
    *mutex = 0;  
}
```

Roadmap

- Solutions for mutual exclusion
 - Peterson's algorithm (Software)
 - Synchronization instructions (Hardware)
- **High-level synchronization mechanisms**
 - **Mutex**
 - **Semaphore**
 - **Monitor**

Spinlock using TestAndSet

```
void init_lock (int *mutex)
{
    *mutex = 0;
}

void lock (int *mutex)
{
    while(TestAndSet(mutex))
        ;
}

void unlock (int *mutex)
{
    *mutex = 0;
}
```

What's Wrong With **Spinlocks**?

- Very wasteful
 - Waiting thread continues to use CPU cycles
 - While doing absolutely nothing but wait
 - 100% CPU utilization, but no useful work done
 - Power consumption, fan noise, ...
- Useful when
 - You hold the lock only *briefly*
- Otherwise
 - A better solution is needed

Mutex – Blocking Lock

- Instead of spinning
 - Let the thread sleep
 - There can be multiple waiting threads
 - In the meantime, let other threads use the CPU
 - When the lock is released, wake-up one thread
 - Pick one if there multiple threads were waiting

```
void mutex_init (mutex_t *lock)
```

```
{
```

```
    lock->value = 0;
```

```
    list_init(&lock->wait_list);
```

```
    spin_lock_init(&lock->wait_lock);
```

```
}
```

← Thread waiting list

← To protect waiting list

```
void mutex_lock (mutex_t *lock)
```

```
{
```

```
...
```

```
    while(TestAndSet(&lock->value)) {
```

```
...
```

```
...
```

```
...
```

```
...
```

```
...
```

```
    }
```

```
...
```

```
}
```

```
void mutex_unlock (mutex_t *lock)
```

```
{
```

```
...
```

```
    lock->value = 0;
```

```
...
```

```
...
```

```
...
```

```
}
```

```
void mutex_init (mutex_t *lock)
```

```
{
```

```
    lock->value = 0;
```

```
    list_init(&lock->wait_list);
```

```
    spin_lock_init(&lock->wait_lock);
```

```
}
```

← Thread waiting list

← To protect waiting list

```
void mutex_lock (mutex_t *lock)
```

```
{
```

```
...
```

```
    while(TestAndSet(&lock->value)) {
```

```
        current->state = WAITING;
```

```
        list_add(&lock->wait_list, current);
```

```
...
```

```
        schedule();
```

```
...
```

```
    }
```

```
...
```

```
}
```

← Thread state change

← Add the current thread to the waiting list

← Sleep or schedule another thread

```
void mutex_unlock (mutex_t *lock)
```

```
{
```

```
...
```

```
    lock->value = 0;
```

```
...
```

```
...
```

```
...
```

```
}
```

```
void mutex_init (mutex_t *lock)
```

```
{
```

```
    lock->value = 0;
```

```
    list_init(&lock->wait_list);
```

```
    spin_lock_init(&lock->wait_lock);
```

```
}
```

← Thread waiting list

← To protect waiting list

```
void mutex_lock (mutex_t *lock)
```

```
{
```

```
    spin_lock(&lock->wait_lock);
```

```
    while(TestAndSet(&lock->value)) {
```

```
        current->state = WAITING;
```

```
        list_add(&lock->wait_list, current);
```

```
    ...
```

```
        schedule();
```

```
    ...
```

```
    }
```

```
    spin_unlock(&lock->wait_lock);
```

```
}
```

← Thread state change

← Add the current thread to the waiting list

← Sleep or schedule another thread

```
void mutex_unlock (mutex_t *lock)
```

```
{
```

```
...
```

```
    lock->value = 0;
```

```
...
```

```
...
```

```
...
```

```
}
```

```
void mutex_init (mutex_t *lock)
```

```
{
```

```
    lock->value = 0;
```

```
    list_init(&lock->wait_list);
```

```
    spin_lock_init(&lock->wait_lock);
```

```
}
```

← Thread waiting list

← To protect waiting list

```
void mutex_lock (mutex_t *lock)
```

```
{
```

```
    spin_lock(&lock->wait_lock);
```

```
    while(TestAndSet(&lock->value)) {
```

```
        current->state = WAITING;
```

```
        list_add(&lock->wait_list, current);
```

```
        spin_unlock(&lock->wait_lock);
```

```
        schedule();
```

```
        spin_lock(&lock->wait_lock);
```

```
    }
```

```
    spin_unlock(&lock->wait_lock);
```

```
}
```

← Thread state change

← Add the current thread to the waiting list

← Sleep or schedule another thread

```
void mutex_unlock (mutex_t *lock)
```

```
{
```

```
...
```

```
    lock->value = 0;
```

```
...
```

```
...
```

```
...
```

```
}
```

```
void mutex_init (mutex_t *lock)
```

```
{
```

```
    lock->value = 0;
```

```
    list_init(&lock->wait_list);
```

```
    spin_lock_init(&lock->wait_lock);
```

```
}
```

← Thread waiting list

← To protect waiting list

```
void mutex_lock (mutex_t *lock)
```

```
{
```

```
    spin_lock(&lock->wait_lock);
```

```
    while(TestAndSet(&lock->value)) {
```

```
        current->state = WAITING;
```

```
        list_add(&lock->wait_list, current);
```

```
        spin_unlock(&lock->wait_lock);
```

```
        schedule();
```

```
        spin_lock(&lock->wait_lock);
```

```
    }
```

```
    spin_unlock(&lock->wait_lock);
```

```
}
```

← Thread state change

← Add the current thread to the waiting list

← Sleep or schedule another thread

```
void mutex_unlock (mutex_t *lock)
```

```
{
```

```
...
```

```
    lock->value = 0;
```

```
    if (!list_empty(&lock->wait_list))
```

```
        wake_up_process(&lock->wait_list)
```

← Someone is waiting for the lock

← Wake-up a waiting thread

```
...
```

```
}
```

```
void mutex_init (mutex_t *lock)
```

```
{
```

```
    lock->value = 0;
```

```
    list_init(&lock->wait_list);
```

```
    spin_lock_init(&lock->wait_lock);
```

```
}
```

← Thread waiting list

← To protect waiting list

```
void mutex_lock (mutex_t *lock)
```

```
{
```

```
    spin_lock(&lock->wait_lock);
```

```
    while(TestAndSet(&lock->value)) {
```

```
        current->state = WAITING;
```

```
        list_add(&lock->wait_list, current);
```

```
        spin_unlock(&lock->wait_lock);
```

```
        schedule();
```

```
        spin_lock(&lock->wait_lock);
```

```
    }
```

```
    spin_unlock(&lock->wait_lock);
```

```
}
```

← Thread state change

← Add the current thread to the waiting list

← Sleep or schedule another thread

```
void mutex_unlock (mutex_t *lock)
```

```
{
```

```
    spin_lock(&lock->wait_lock);
```

```
    lock->value = 0;
```

```
    if (!list_empty(&lock->wait_list))
```

```
        wake_up_process(&lock->wait_list)
```

```
    spin_unlock(&lock->wait_lock);
```

```
}
```

← Someone is waiting for the lock

← Wake-up a waiting thread

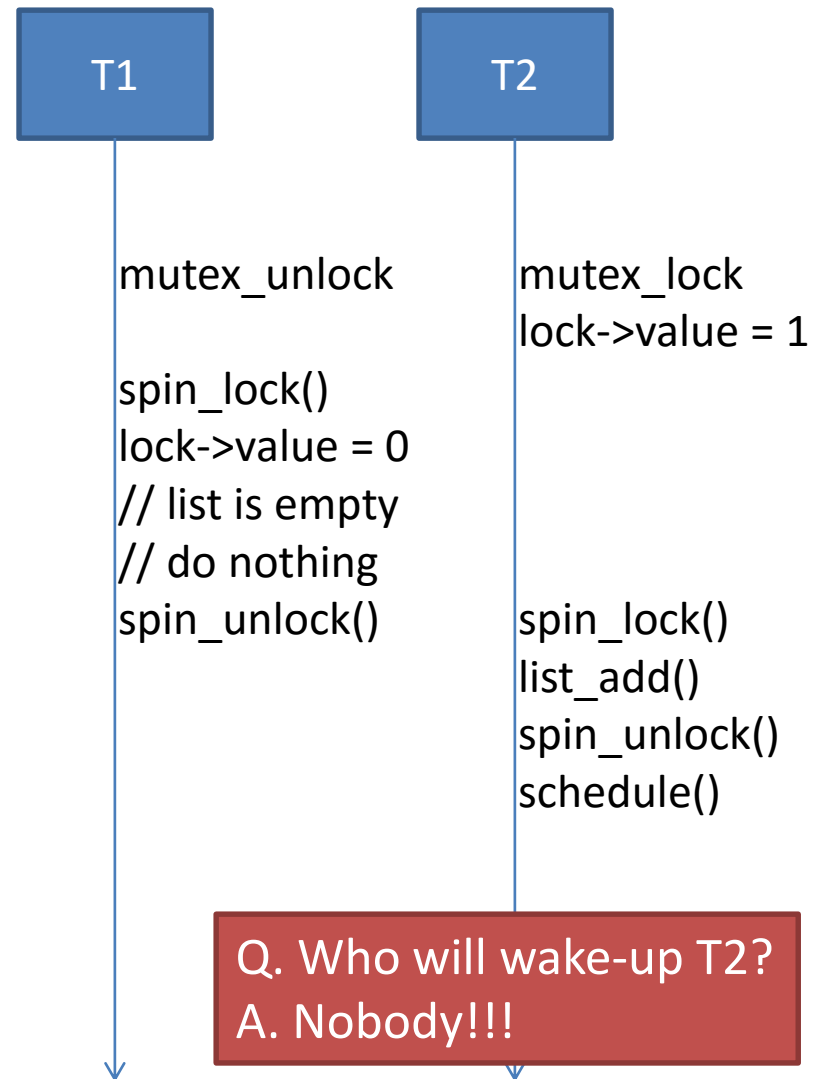
```
void mutex_init (mutex_t *lock)
{
    lock->value = 0;
    list_init(&lock->wait_list);
    spin_lock_init(&lock->wait_lock);
}
```

```
void mutex_lock (mutex_t *lock)
{
    while(TestAndSet(&lock->value)) {
        current->state = WAITING;
        spin_lock(&lock->wait_lock);
        list_add(&lock->wait_list, current);
        spin_unlock(&lock->wait_lock);
        schedule();
    }
}
```

Correct?

```
void mutex_unlock (mutex_t *lock)
{
    spin_lock(&lock->wait_lock);
    lock->value = 0;
    if (!list_empty(&lock->wait_list))
        wake_up_process(&lock->wait_list)
    spin_unlock(&lock->wait_lock);
}
```

More reading: [mutex.c in Linux](#)



Agenda

- High-level synchronization mechanisms
 - Mutex
 - **Semaphore**
 - **Monitor**

High-level Synchronization Primitives

- Lock (mutex) is great, but...
 - Too low-level primitive
 - Sometimes we need more **powerful primitives**
- Semaphore
 - Binary/integer semaphore
- Monitor
 - Condition variable

Semaphore

- High-level synchronization primitive
 - Designed by Dijkstra in 1960'
- Definition
 - Semaphore is an integer variable
 - Only two operations are possible:
 - P() or wait() or down()
 - V() or signal() or up()

Simple Semaphore Implementation

- P() operation

```
P(semaphore *S) {  
    while(S->value == 0) {  
        S->list->addQ(P);  
        schedule( );  
    }  
    S->value— ;  
}
```

*schedule () – schedule
another thread*

- V() operation

```
V(semaphore *S) {  
    if(!IsEmpty(&S->list)) {  
        P = delQ(&S->list);  
        wakeup(P);  
    }  
    S->value++ ;  
}
```

wakeup() – wake up a thread

What's wrong with the code?

Simple Semaphore Implementation

- P() operation

```
P(semaphore *S) {  
    S->lock->Acquire();  
    while(S->value == 0) {  
        addQ(&S->list, P);  
        S->lock->Release();  
        schedule( );  
        S->lock->Acquire();  
    }  
    S->value-- ;  
    S->lock->Release();  
}
```

- V() operation

```
V(semaphore *S) {  
    S->lock->Acquire();  
    if(!isEmpty(&S->list)) {  
        P = delQ(&S->list);  
        wakeup(P);  
    }  
    S->value++ ;  
    S->lock->Release();  
}
```