

Synchronization

Heechul Yun

Disclaimer: some slides are adopted from the book authors and Dr. Kulkani

Recap

- Semaphore
 - Blocking
 - Binary semaphore = mutex
 - Integer semaphore
 - Solved synchronization problems
 - Bounded-buffer
 - 2 Integer semaphores, 1 binary semaphore
 - Train control
 - 1 integer semaphore
 - Reader/write
 - 2 binary semaphores

Recap: Train Control Problem

- No more than two trains can enter the rail


```
Enter() {  
  Leave()  
}
```



Recap: Train Control Problem

- No more than two trains can enter the rail

```
Enter() { P(); }  
Leave() { V(); }
```



Quiz

- Semaphore mutex = 1;
- Semaphore **full** = 0;
- Semaphore **empty** = N;

Producer

```
do {  
    Produce new resource  
    _____;  
    mutex.P();  
    Add resource to next buffer  
    mutex.V();  
    _____;  
} while (TRUE);
```

Consumer

```
do {  
    _____;  
    mutex.P();  
    Remove resource from buffer  
    mutex.V();  
    _____;  
    Consume resource  
} while (TRUE);
```

Quiz

- Semaphore mutex = 1;
- Semaphore **full** = 0;
- Semaphore **empty** = N;

Producer

```
do {  
    Produce new resource  
    empty.P();  
    mutex.P();  
    Add resource to next buffer  
    mutex.V();  
    full.V();  
} while (TRUE);
```

Consumer


```
do {  
    full.P();  
    mutex.P();  
    Remove resource from buffer  
    mutex.V();  
    empty.V();  
    Consume resource  
} while (TRUE);
```

Monitor

- Monitor
 - A lock (mutual exclusion) + condition variables (scheduling)
 - Some languages like Java natively support this, but you can use monitors in other languages like C/C++
- Lock: mutual exclusion
 - Protects the shared data structures inside the monitor
 - Always acquire it to enter the monitor
 - Always release it to leave the monitor
- Condition Variable: scheduling
 - Allow thread to wait on certain events inside the monitor
 - Key idea: to wait (sleep) inside the monitor, it first releases the lock and go to sleep atomically

Monitor

- Lock: mutual exclusion
 - Only one thread can execute any monitor procedure at a time.
 - Other threads invoking a monitor procedure when one is already executing some monitor procedure must wait.
 - When the active thread exits the monitor procedure, one other waiting thread can enter.



A Simple Monitor Example

C++

```
Mutex lock;  
Queue queue;  
  
produce (item)  
{  
    lock.acquire();  
    queue.enqueue(item);  
    lock.release();  
}  
  
consume()  
{  
    lock.acquire();  
    item = queue.dequeue(item);  
    lock.release();  
    return item;  
}
```

A Simple Monitor Example

C++

```
Mutex lock;
Queue queue;

produce (item)
{
    lock.acquire();
    queue.enqueue(item);
    lock.release();
}

consume()
{
    lock.acquire();
    item = queue.dequeue(item);
    lock.release();
    return item;
}
```

Java

```
Queue queue;

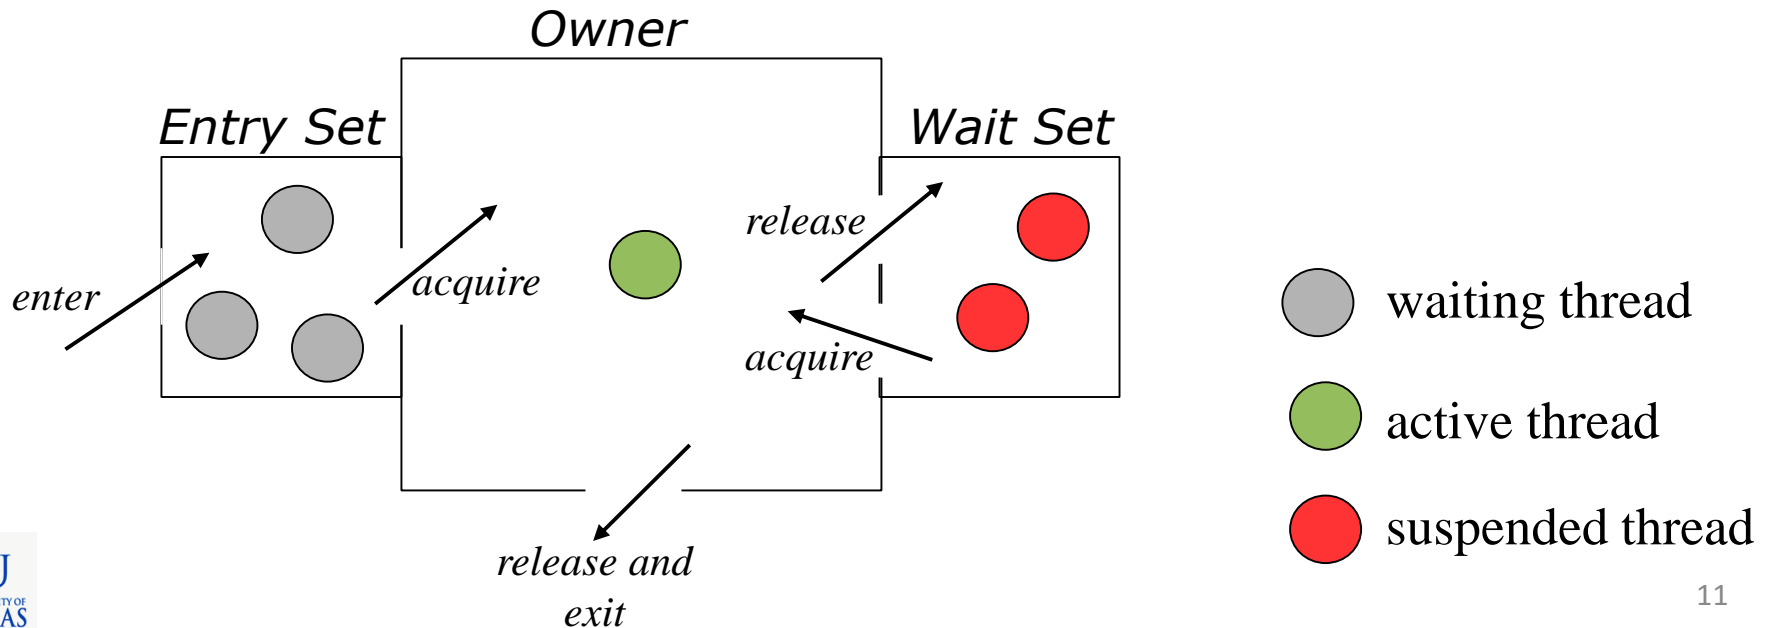
Synchronized produce (item)
{
    queue.enqueue(item);
}

Synchronized consume()
{
    item = queue.dequeue(item);

    return item;
}
```

Monitor

- What if a thread needs to wait inside a monitor?
- Condition variable: Scheduling
 - Wait(&lock): atomically release the lock and sleep; Re-acquire the lock on returning.
 - Signal(): wake-up one waiter, if exists
 - Broadcast(): wake-up all waiters



Monitor with Condition Variable

```
Mutex lock;  
Condition full;  
Queue queue;  
  
produce (item)  
{  
    lock.acquire();  
    queue.enqueue(item);  
    full.signal();  
    lock.release();  
}  
  
consume()  
{  
    lock.acquire();  
    while (queue.isEmpty())  
        full.wait(&lock);  
    item = queue.dequeue(item);  
    lock.release();  
    return item;  
}
```

Why not 'if'?



Semantics

- Hoare monitors (original)
 - Signaler immediately switches to the waiting thread
 - Waiter's condition is guaranteed to hold when it resumes
- Mesa monitors
 - Waiter is simply placed on ready queue, signaler continues
 - Waiter's condition may no longer be true when it resumes
- Almost always Mesa style in practice

Bounded Buffer Problem Revisit

```
Mutex lock;  
Condition full, empty;  
  
produce (item)  
{  
    lock.acquire();  
    ...  
    ...  
    queue.enqueue(item);  
    full.signal();  
    lock.release();  
}  
  
consume()  
{  
    lock.acquire();  
    while (queue.isEmpty())  
        full.wait(&lock);  
    item = queue.dequeue(item);  
    ...  
    lock.release();  
    return item;  
}
```

Bounded Buffer Problem Revisit

```
Mutex lock;  
Condition full, empty;  
  
produce (item)  
{  
    lock.acquire();  
    while (queue.isFull())  
        empty.wait(&lock);  
    queue.enqueue(item);  
    full.signal();  
    lock.release();  
}  
  
consume()  
{  
    lock.acquire();  
    while (queue.isEmpty())  
        full.wait(&lock);  
    item = queue.dequeue(item);  
    empty.signal();  
    lock.release();  
    return item;  
}
```

Bounded Buffer Problem Revisit

Monitor version

```
Mutex lock;  
Condition full, empty;  
  
produce (item)  
{  
    lock.acquire();  
    while (queue.isFull())  
        empty.wait(&lock);  
    queue.enqueue(item);  
    full.signal();  
    lock.release();  
}  
  
consume()  
{  
    lock.acquire();  
    while (queue.isEmpty())  
        full.wait(&lock);  
    item = queue.dequeue(item);  
    empty.signal();  
    lock.release();  
    return item;  
}
```

Semaphore version

```
Semaphore mutex = 1, full = 0,  
empty = N;  
  
produce (item)  
{  
    P(&empty);  
    P(&mutex);  
    queue.enqueue(item);  
    V(&mutex);  
    V(&full);  
}  
  
consume()  
{  
    P(&full);  
    P(&mutex);  
    item = queue.dequeue();  
    V(&mutex);  
    V(&empty);  
    return item;  
}
```


More Synchronization Primitives

- RCU (Read-Copy-Update)
 - Optimized for frequent read, infrequent write
 - Read is zero cost
 - Write can be costly
 - Heavily used in Linux kernel
- Transactional memory (Intel TSX instruction)
 - Opportunistic synchronization
 - Declare a set of instructions as a transaction
 - If no other CPUs update the shared data while executing the transaction, it is committed
 - If not (i.e., someone tries to modify in the middle of the transaction), the transaction will be aborted
 - If successful, there's no synchronization overhead

Summary

- Synchronization
 - Spinlock
 - Implement using h/w instructions (e.g., test-and-set)
 - Mutex
 - Sleep instead of spinning.
 - Semaphore
 - powerful tool, but often difficult to use
 - **Monitor**
 - Powerful and (relatively) easy to use

Team Project

- At company, you will likely to work in a team.
So, better to practice now
- Ideal team size = ??
 - Two is the minimum and a very good one
 - Communication, communication, communication
- A couple of tips
 - Work division
 - Coding standard
 - Version control software

Work Division

- Function based
 - Person 1: built-in commands (cd,...)
 - Person 2: pipe and I/O redirection,..
- Job based
 - Person 1: coder
 - Person 2: tester

Coding Standard

- Function naming

```
PrintData()  
print_data()
```

- indentation, commenting, ...

```
/**  
 * commenting style 1  
 */  
If (condition) {  
    // 4 spaces or  
    // 8 spaces  
}
```

Version Control Software

- Git
 - Most popular these days.
 - Relatively high learning curve
 - <https://git.eecs.ku.edu/> or github
- SVN
 - Still very popular
 - Relatively simple to learn
 - <http://wiki.eecs.ku.edu/subversion>

Agenda

- Famous Synchronization Bugs
 - THERAC-25
 - Mars Pathfinder

Therac 25



Image source: http://idg.bg/test/cwd/2008/7/14/21367-radiation_therapy.JPG

- Computer controlled medical X-ray treatments
- Six people died/injured due to massive overdoses (1985-1987)

Accident History

Date	What happened
June 1985	First overdose
July-Dec 1985	2nd and 3rd overdoses. Lawsuit against the manufacturer and hospital
Jan-Feb 1986	Manufacturer denied the possibility of overdoses
Mar-Apr 1986	Two more overdoses
May-Dec 1986	FDA orders corrective action plans to the manufacturer
Jan 1987	Sixth overdose
Nov 1988	Final safety analysis report

The Problem




Image source: <http://radonc.wdfiles.com/local--files/radiation-accident-therac25/Therac25.png>


- X-ray must be dosed with the filter in place
- But sometimes, X-ray was dosed w/o the filter

The Bug

```
unsigned char in_progress = 1;
```

```
Thread 1 : // tray movement thread (periodic)
    if (system_is_ready())
        in_progress = 0;
    else
        in_progress++;
```

```
Thread 2 : // X-ray control thread.
    if (in_progress == 0)
        start_radiation();
```



tray including the target, a flattening filter, the collimator jaws and an ion chamber was moved OUT for "electron" mode, and IN for "photon" mode.


- Can you spot the bug?

Fixed Code

```
unsigned char in_progress;

Thread 1 : // tray movement thread (periodic)
    if (system_is_ready())
        in_progress = 0;
    else
        in_progress = 1;

Thread 2 : // X-ray control thread.
    if (in_progress == 0)
        start_radiation();
```



- Can you do better using a monitor?

Monitor Version

```
Mutex lock;  
Condition ready;  
unsigned char in_progress;
```

```
Thread 1 : // on finishing tray movement  
    lock.acquire();  
    in_progress = 0;  
    ready.signal();  
    lock.release();
```

```
Thread 2 : // X-ray control thread.  
    lock.acquire();  
    while (in_progress)  
        ready.wait(&lock);  
    start_radiation();  
    lock.release();
```

- No periodic check is needed.

Mars Pathfinder




- Landed on Mars, July 4, 1997
- After operating for a while, it rebooted itself

The Bug

- Three threads with priorities
 - Weather data thread (low priority)
 - Communication thread (medium priority)
 - Information bus thread (high priority)
- Each thread obtains a lock to write data on the shared memory
- High priority thread can't acquire the lock for a very long time → something must be wrong. Let's reboot!

Priority Inversion



- High priority thread is delayed by the medium priority thread (potentially) indefinitely!!!


Solution

- Priority inheritance protocol [Sha'90]
 - If a high priority thread is waiting on a lock, boost the priority of the lock owner thread (low priority) to that of the high priority thread.
- Remotely patched the code
 - To use the priority inheritance protocol in the lock
 - First-ever(?) interplanetary remote debugging

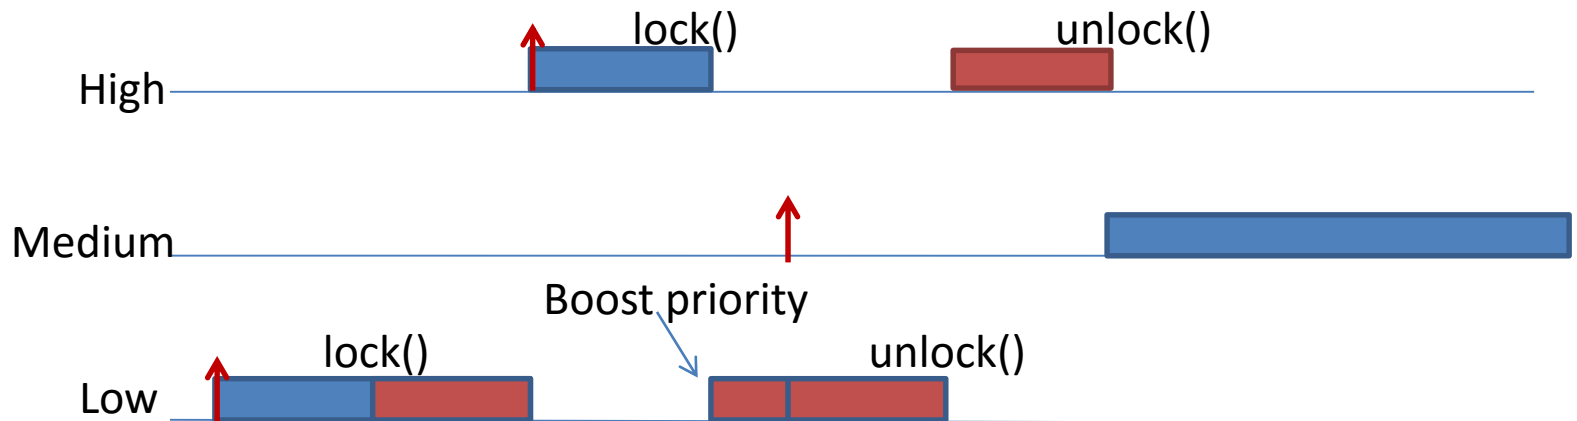
L. Sha, R. Rajkumar, and J. P. Lehoczky. Priority Inheritance Protocols: An Approach to Real-Time Synchronization. In IEEE Transactions on Computers, vol. 39, pp. 1175-1185, Sep. 1990.

Priority Inheritance

Old



New



Summary

- Race condition
 - A situation when two or more threads **read and write** shared data at the same time
- Critical section
 - Code sections of potential race conditions
- Mutual exclusion
 - If a thread executes its critical section, *no other threads* can enter their critical sections
- Peterson's solution
 - Software only solution providing mutual exclusion

Summary

- Spinlock
 - Spin on waiting
 - Use synchronization instructions (test&set)
- Mutex
 - Sleep on waiting
- Semaphore
 - Powerful tool, but often difficult to use
- Monitor
 - Powerful and (relatively) easy to use

Quiz

```
Mutex lock;  
Condition full, empty;  
  
produce (item)  
{  
    _____  
    while (queue.isFull())  
        empty.wait(&lock);  
    queue.enqueue(item);  
    full.signal();  
    _____  
}  
  
consume()  
{  
    _____  
    while (queue.isEmpty())  
        _____  
    item = queue.dequeue(item);  
    _____  
    return item;  
}
```

Quiz

```
Mutex lock;  
Condition full, empty;  
  
produce (item)  
{  
    lock.acquire();  
    while (queue.isFull())  
        empty.wait(&lock);  
    queue.enqueue(item);  
    full.signal();  
    lock.release();  
}  
  
consume()  
{  
    lock.acquire();  
    while (queue.isEmpty())  
        full.wait(&lock);  
    item = queue.dequeue(item);  
    empty.signal();  
    lock.release();  
    return item;  
}
```