

# Deadlock

Disclaimer: some slides are adopted from Dr. Kulkarni's and book authors' slides with permission

# Recap: Synchronization

- Race condition
  - A situation when two or more threads **read and write** shared data at the same time
- Critical section
  - Code sections of potential race conditions
- Mutual exclusion
  - If a thread executes its critical section, *no other threads* can enter their critical sections
- Peterson's solution
  - Software only solution providing mutual exclusion


# Recap: Synchronization

- Spinlock
  - Spin on waiting
  - Use synchronization instructions (test&set)
- Mutex
  - Sleep on waiting
- Semaphore
  - Powerful tool, but often difficult to use
- Monitor
  - Powerful and (relatively) easy to use

# Agenda

- Deadlock
  - Starvation vs. deadlock
  - Deadlock conditions
  - General solutions: detection and prevention
  - Detection algorithm
  - Banker's algorithm

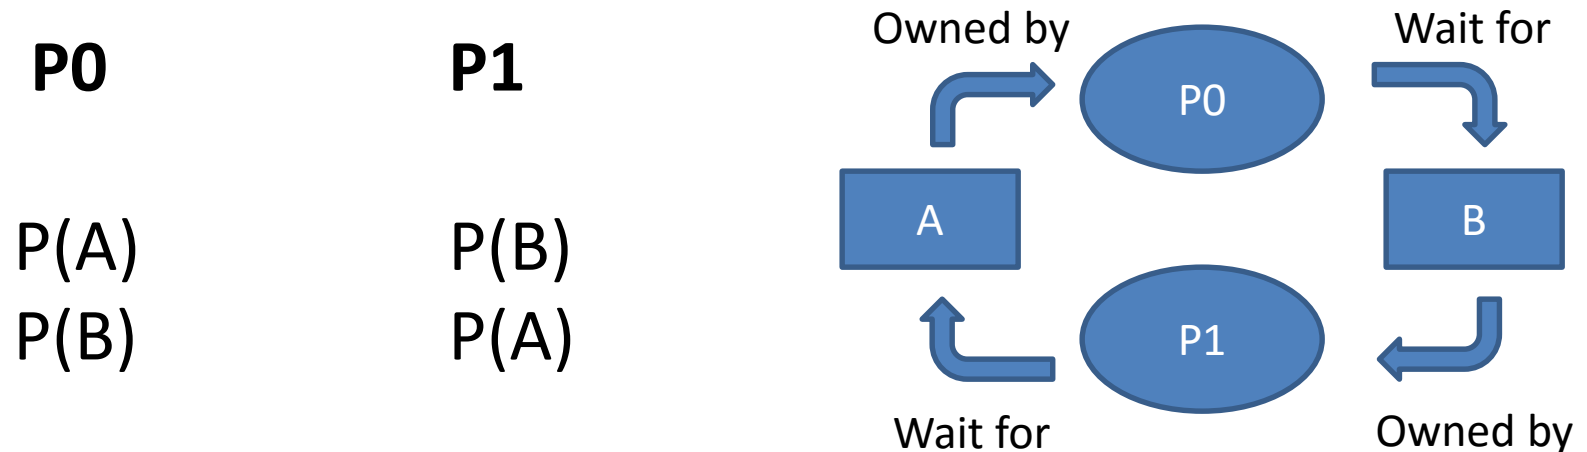
# Starvation



- Starvation
  - Wait potentially **indefinitely**, but it **can end**


# Starvation vs. Deadlock

- Deadlock: circular waiting for resources
  - Example: semaphore A = B = 1



- Deadlock → Starvation
  - But reverse is not true
  - Deadlock can't end but starvation can


# Bridge Crossing



- Traffic only in one direction
- Each section of a bridge can be viewed as a resource
- If a deadlock occurs, how to fix it?
  - Make one car backs up
  - Several cars may have to be backed up if a deadlock occurs

# Dining Philosophers

- Problem synopsis
  - Need two chopsticks to eat
  - Grab one chopsticks at a time
- What happens if all grab left chopstick at the same time??
  - Deadlock!!!
- How to fix it?
- How to avoid it?





# Conditions for Deadlocks

- Mutual exclusion
  - only one process at a time can use a resource
- No preemption
  - resources cannot be preempted, release must be voluntary
- Hold and wait
  - a process must be holding at least one resource, and waiting to acquire additional resources held by other processes
- Circular wait
  - There must be a circular dependency. For example, A waits B, B waits C, and C waits A.
- **All four conditions must simultaneously hold**

# Resource-Allocation Graph

- To illustrate deadlock conditions.
- Graph consists of a set of vertices  $V$  and a set of edges  $E$
- $V$  is partitioned into two types:
  - $P = \{P_1, P_2, \dots, P_n\}$ , the set consisting of all the processes in the system
  - $R = \{R_1, R_2, \dots, R_m\}$ , the set consisting of all resource types in the system
- request edge – directed edge  $P_i \rightarrow R_j$
- assignment edge – directed edge  $R_j \rightarrow P_i$

# Resource-Allocation Graph


- Process




- Resource Type with 4 instances






- $P_i$  requests instance of  $R_j$



- $P_i$  is holding an instance of  $R_j$



# Resource Allocation Graph



- request edge – directed edge  $P_i \rightarrow R_j$
- assignment edge – directed edge  $R_j \rightarrow P_i$

# Methods for Handling Deadlocks

- Detection and recovery
  - Allow a system to enter a deadlock and then recover
    - Need a *detection algorithm*
    - Somehow “preempt” resources
- Prevention and avoidance
  - Ensure a system never enter a deadlock
  - Possible solutions
    - have “Infinite resources”
    - prevent “hold and wait”
    - prevent “circular wait”

Recall four deadlock conditions:

(1) Mutual exclusion, (2) no preemption, (3) hold and wait, (4) circular wait


# Deadlock Detection

- Deadlock detection algorithms
  - Single instance for each resource type
  - Multiple instances for each resource type


# Single Instance Per Resource

- Each resource is unique
  - E.g., one printer, one audio card, ...
- Wait-for-graph
  - Variant of the simplified resource allocation graph
  - Remove resource nodes, collapse corresponding edges
- Detection algorithm
  - Searches for a cycle in the wait-for graph
  - Presence of a cycle points to the existence of a deadlock

# Wait-for Graph



Resource-Allocation Graph




Corresponding wait-for graph



# Multiple Instances Per Resource

- **n** processes, **m** resources
- **FreeResources**: resource vector (of size m)
  - indicates the number of available resources of each type
  - $[R_1, R_2] = [0, 0]$
- **Alloc[i]**: process i's allocated resource vector
  - defines the number of resources of each type currently allocated to each process
  - $\text{Alloc}[1] = [0, 1]$ ,
  - $\text{Alloc}[2] = [1, 0], \dots$
- **Request[i]**: process i's requesting resource vector
  - indicates the resources each process requests
  - $\text{Request}[1] = [1, 0]$ ,
  - $\text{Request}[2] = [0, 0], \dots$



# Detection Algorithm

1. Initialize **Avail** and **Finish** vectors

Avail = FreeResources;

For  $i = 1, 2, \dots, n$ , Finish[i] = false

2. Find an index  $i$  such that

Finish[i] == false AND Request[i]  $\leq$  Avail

If no such  $i$  exists, go to step 4

3. Avail = Avail + Alloc[i], Finish[i] = true

Go to step 2


4. If Finish[i] == false, for some  $i$ ,  $1 \leq i \leq n$ ,

(a) then the system is in deadlock state

■ **FreeResources:** resource vector  
[R1, R2] = [0,0]

■ **Alloc[i]:** process  $i$ 's allocated resource vector:  
Alloc[1] = [0,1], Alloc[2] = [1, 0]

■ **Request[i]:** process  $i$ 's requesting vector:  
Request[1] = [1,0]  
Request[2] = [0,0]



# Recovery from Deadlock

- Terminate
  - Preempt the resources
  - Bridge example: throw the car to the river
  - Kill the deadlocked threads and return the resources
- Rollback
  - Return to a known safe state
  - Bridge example: move one car backward
  - Dining philosopher: make one philosopher give up a chopstick
- Not always possible!

# Deadlock Prevention

- Break any of the four deadlock conditions
  - Mutual exclusion
  - No preemption
  - Hold and wait
  - Circular wait

# Deadlock Prevention

- Break any of the four deadlock conditions
  - **Mutual exclusion → allow sharing**
    - Well, not all resources are sharable
  - No preemption
  - Hold and wait
  - Circular wait

# Deadlock Prevention

- Break any of the four deadlock conditions
  - Mutual exclusion → allow sharing
    - Well, not all resources are sharable
  - **No preemption → allow preemption**
    - This is also quite hard (kill the threads)
  - Hold and wait
  - Circular wait

# Deadlock Prevention

- Break any of the four deadlock conditions
  - Mutual exclusion → allow sharing
    - Well, not all resources are sharable
  - No preemption → allow preemption
    - This is also quite hard (kill the threads)
  - **Hold and wait → get all resources at once**
    - Dining philosopher: get *both* chopsticks or *none*
  - Circular wait

# Deadlock Prevention

- Break any of the four deadlock conditions
  - Mutual exclusion → allow sharing
    - Well, not all resources are sharable
  - No preemption → allow preemption
    - This is also quite hard (kill the threads)
  - Hold and wait → get all resources at once
    - Dining philosopher: get *both* chopsticks or *none*
  - **Circular wait → prevent cycle**
    - Dining philosopher: change the chopstick picking order;  
**if grabbing a chopstick will form a cycle, prevent it.**



# Banker's Algorithm

- General idea
  - Assume that each process's maximum resource demand is known in advance
    - $\text{Max}[i]$  : process  $i$ 's maximum resource demand vector
  - **Pretend** each request is granted, then run the deadlock detection algorithm
  - If a deadlock is detected, the do not grant the request to keep the system in a **safe** state

# Banker's Algorithm

## 1. Initialize **Avail** and **Finish** vectors

Avail = FreeResources;

For  $i = 1, 2, \dots, n$ , Finish[i] = false

## 2. Find an index $i$ such that

Finish[i] == false AND

$$\mathbf{Max[i] - Alloc[i] \leq Avail}$$

If no such  $i$  exists, go to step 4

## 3. Avail = Avail + Alloc[i], Finish[i] = true

Go to step 2

## 4. If Finish[i] == false, for some $i$ , $1 \leq i \leq n$ ,

(a) then the system is in deadlock state

(b) if Finish[i] == false, then  $P_i$  is deadlocked

■ FreeResources: resource vector  
[R1, R2] = [0,0]


■ Alloc[i]: process  $i$ 's allocated resource vector:  
Alloc[1] = [0,1], Alloc[2] = [1, 0]

■ Request[i]: process  $i$ 's requesting vector:  
Request[1] = [1,0]  
Request[2] = [0,0]

■ **Max[i]**: process  $i$ 's maximum resource demand vector

# Example

Free = [1,1,1,1,1]



# Example

Free = [0,0,0,0,1]

Avail = [0,0,0,0,1]


Max[2] = [1,1,0,0,0]  
Alloc[2] = [0,1,0,0,0]

Max[3] = [0,1,1,0,0]  
Alloc[3] = [0,0,1,0,0]

Max[1] = [1,0,0,0,1]  
Alloc[1] = [1,0,0,0,0]

Max[4] = [0,0,1,1,0]  
Alloc[4] = [0,0,0,1,0]

Max[5] = [0,0,0,1,1]  
Alloc[5] = [0,0,0,0,0]



- Philosopher 5 requested R5.
- **Safe or Unsafe?**

# Example

Free = [0,0,0,0,0]

Avail = [0,0,0,0,0]


Max[2] = [1,1,0,0,0]  
Alloc[2] = [0,1,0,0,0]

Max[3] = [0,1,1,0,0]  
Alloc[3] = [0,0,1,0,0]

Max[1] = [1,0,0,0,1]  
Alloc[1] = [1,0,0,0,0]

Max[4] = [0,0,1,1,0]  
Alloc[4] = [0,0,0,1,0]

Max[5] = [0,0,0,1,1]  
Alloc[5] = [0,0,0,0,1]



2. Find an index  $i$  such that  
Finish[ $i$ ] == false AND  
**Max[ $i$ ] - Alloc[ $i$ ] ≤ Avail**  
If no such  $i$  exists, go to step 4

# Quiz

- Using Banker's algorithm, determine whether this state is safe or unsafe.

Total resources: 10

**Avail** resources: 1

Process	Max	Alloc
$P_0$	10	4
$P_1$	3	1
$P_2$	6	4

# Quiz

- Using Banker's algorithm, determine whether this state is safe or unsafe.

Total resources: 10

**Avail** resources: 1

Process	Max	Alloc
P <sub>0</sub>	10	4
P <sub>1</sub>	3	1
P <sub>2</sub>	6	4

$$10 - 4 \not\leq 1$$

$$3 - 1 \not\leq 1$$

$$6 - 4 \not\leq 1$$

Unsafe

# Summary

- Deadlock
  - Four deadlock conditions:
    - Mutual exclusion
    - No preemption
    - Hold and wait
    - Circular wait
  - Detection
  - Avoidance