

# CPU Scheduling

# Administrative

- Midterm
  - Mar. 11, 2016
  - Closed book, in-class
  - Review class: Mar. 9, 2016


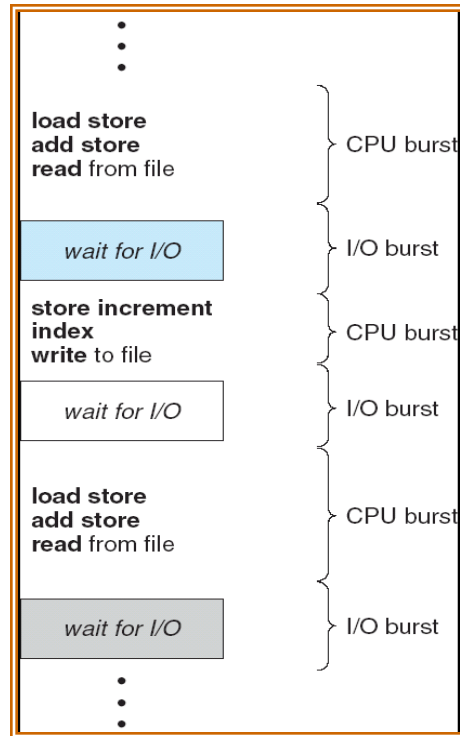
# Agenda

- Introduction to CPU scheduling
- Classical CPU scheduling algorithms

# CPU Scheduling

- CPU scheduling is a **policy** to decide
  - **Which** thread to run next?
  - **When** to schedule the next thread?
  - **How long?**
- Context switching is a **mechanism**
  - To change the running thread


# Assumption: CPU Bursts



- Execution model
  - Program uses the CPU for a while and then does some I/O, back to use CPU, ..., keep alternating


# CPU Scheduler

- An OS component that determines which thread to run, at what time, and how long
  - Among threads in the **ready queue**



# CPU Scheduler

- When the scheduler runs?



- The running thread finishes
- The running thread voluntarily gives up the CPU
  - yield, block on I/O, ...
- The OS **preempts** the current running thread
  - quantum expire (timer interrupt)


# Performance Metrics for CPU Scheduling

- CPU utilization
  - % of time the CPU is busy doing something
- Throughput
  - #of jobs done / unit time
- **Turnaround time**
  - Time to complete a task (ready -> complete)
- **Waiting time**
  - Time spent on waiting in the ready queue
- **Response time**
  - Time to schedule a task (ready -> first scheduled)



# Example


- Assumption: A, B, C are released at time 0



- The times of Process **A**
  - Turnaround time: 9
  - Wait time: 5
  - Response time: 0

# Example


- Assumption: A, B, C are released at time 0



- The times of Process **B**
  - Turnaround time: 5
  - Wait time: 3
  - Response time: 1

# Example

- Assumption: A, B, C are released at time 0



- The times of Process **C**
  - Turnaround time: 10
  - Wait time: 6
  - Response time: 2


# Workload Model and Gantt Chart

- Workload model

Process	Arrival Time	Burst Time
P1	0	8
P2	1	4
P3	1	10
P4	6	2

- Gantt chart

- bar chart to illustrate a particular schedule



# Scheduling Policy Goals

- Maximize throughput
  - High throughput (#of jobs done / time) is always good
- Minimize response/completion time
  - Important to interactive applications (games, editor, ...)
- Fairness
  - Make all threads progress equally
- Goals often conflicts
  - Frequent context switching may be good for reducing response time, but not so much for maximizing throughput

# First-Come, First-Served (FCFS)

- FCFS
  - Assigns the CPU based on the order of the requests.
  - Implemented using a FIFO queue.




© bnpdesignstudio \* [www.ClipartOf.com/226258](http://www.ClipartOf.com/226258)

# FCFS

- Example

Process	Arrival Time	Burst Time
P1	0	24
P2	0	3
P3	0	3

– Suppose that the processes arrive in the order:  $P_1, P_2, P_3$



– Waiting time?

- $P1 = 0$ ;  $P2 = 24$ ;  $P3 = 27$

– Average waiting time


- $(0 + 24 + 27)/3 = 17$

# FCFS

- Example 2

Process	Arrival Time	Burst Time
P1	0	24
P2	0	3
P3	0	3

– Suppose that the processes arrive in the order:  $P_2, P_3, P_1$



– Waiting time?

- $P_1 = 6$ ;  $P_2 = 0$ ;  $P_3 = 3$

– Average waiting time

- $(6 + 0 + 3)/3 = 3$

– Much better than previous case → performance varies greatly depending on the scheduling order



# Shortest Job First (SJF)


- Can we always do the best FIFO?
  - Yes: **if** you know the tasks' CPU burst times
- Shortest Job First (SJF)
  - Order jobs based on their burst lengths
  - Executes the job with the shortest CPU burst first
  - SJF is optimal
    - Achieves minimum average waiting time

# Shortest Job First (SJF)

- Example

Process	Arrival Time	Burst Time
P1	0	6
P2	0	8
P3	0	7
P4	0	3

- Gantt chart



- Average waiting time?


- $(3 + 16 + 9 + 0) / 4 = 7$

- How to know the CPU burst time **in advance**?

# Determining CPU Burst Length

- Can only estimate the length
  - Next CPU burst similar to previous CPU bursts ?
  - Predict based on the past history
- Exponential weighted moving average (EWMA)
  - of past CPU bursts

1.  $t_n$  = actual length of  $n^{th}$  CPU burst
2.  $\tau_{n+1}$  = predicted value for the next CPU burst
3.  $\alpha, 0 \leq \alpha \leq 1$
4. Define:  $\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n$ .



CPU burst ( $t_i$ )	6	4	6	4	13	13	13	...	
"guess" ( $\tau_i$ )	10	8	6	6	5	9	11	12	...

# Recap

- CPU Scheduling
  - Decides which thread, when, and how long?
- Metrics
  - Turnaround time
    - Time to complete a task (ready -> complete)
  - Waiting time
    - Time spent on waiting in the ready queue
  - Response time
    - Time to schedule a task (ready -> first scheduled)
- FIFO
- SJF


# Administrivia

- Project 1 grading policy is updated.
  - Deadline: Tonight
- Quiz1 is posted
  - Chapter 1-4
  - Due: Friday

# Shortest Job First (SJF)

- What if jobs don't arrive at the same time?

Process	Arrival Time	Burst Time
P1	0	8
P2	1	4
P3	2	9
P4	3	5




– Average waiting time

- $(0+7+15+8)/4 = 7.5$

# Shortest Remaining Time First (SRTF)

- Preemptive version of SJF
- New shorter job preempt longer running job

Process	Arrival Time	Burst Time
P1	0	8
P2	1	4
P3	2	9
P4	3	5

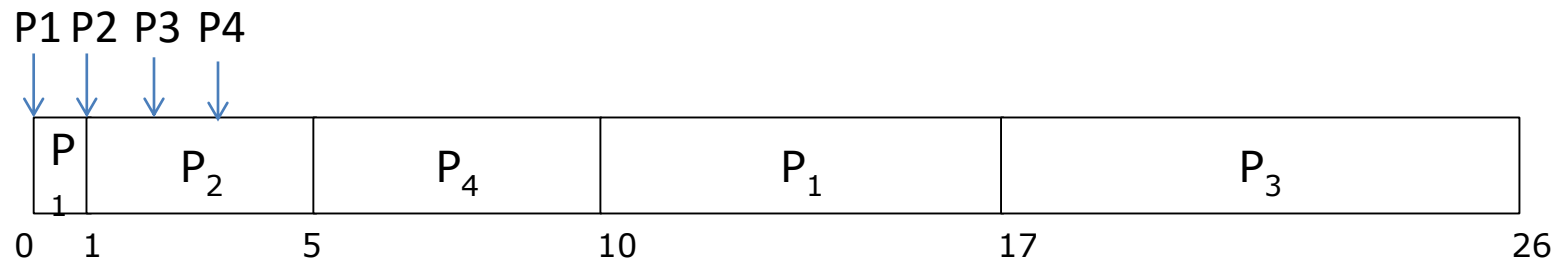


- Average waiting time  
–  $(9 + 0 + 15 + 2) / 4 = 6.5$

# Quiz: SRTF

- Average waiting time?

Process	Arrival Time	Burst Time
P1	0	8
P2	1	4
P3	2	9
P4	3	5



- $(9 + 0 + 15 + 2) / 4 = 6.5$



# Summary

- FIFO
  - In the order of arrival
  - Non-preemptive
- SJF
  - Shortest job first.
  - Non preemptive
- SRTF
  - Preemptive version of SJF

# Issues

- FIFO
  - Bad average turn-around time
- SJF/SRTF
  - Good average turn-around time
  - IF you know or can predict the future
- Time-sharing systems
  - Multiple users share a machine
  - Need high interactivity → low **response time**

# Round-Robin (RR)

- FIFO with preemption
- Simple, fair, and easy to implement
- Algorithm
  - Each job executes for a fixed time slice: **quantum**
  - When quantum expires, the scheduler preempts the task
  - Schedule the next job and continue...


# Round-Robin (RR)

- Example

- Quantum size = 4

Process	Burst Times
P1	24
P2	3
P3	3

- Gantt chart



- Response time (between ready to first schedule)


- P1: 0, P2: 4, P3: 7. average response time =  $(0+4+7)/3 = 3.67$

- Waiting time

- P1: 6, P2: 4, P3: 7. average waiting time =  $(6+4+7)/3 = 5.67$

# How To Choose Quantum Size?

- Quantum length
  - Too short → high overhead (why?)
  - Too long → bad response time
    - Very long quantum → FIFO




# Round-Robin (RR)

- Example

- Quantum size = 2

Process	Burst Times
P1	24
P2	3
P3	3

- Gantt chart



- Response time (between ready to first schedule)

- P1: 0, P2: 2, P3: 4. average response time =  $(0+2+4)/3 = 2$

- Waiting time

- P1: 6, P2: 6, P3: 7. average waiting time =  $(6+6+7)/3 = 6.33$

# Discussion

- Comparison between FCFS, SRTF(SJF), and RR
  - What to choose for smallest average waiting time?
    - SRTF (SFJ) is the optimal
  - What to choose for better interactivity?
    - RR with small time quantum (or SRTF)
  - What to choose to minimize scheduling overhead?
    - FCFS

# Example



- Task A and B
  - CPU bound, run an hour
- Task C
  - I/O bound, repeat(1ms CPU, 9ms disk I/O)
- FCFS?
  - If A or B is scheduled first, C can begin an hour later
- RR and SRTF?




# Example Timeline




RR with 100ms time quantum

C A B ... A B C A B A B ... C A B ...



RR with 1ms time quantum

C            A            C            A            C A ...



SRTF

# Summary

- First-Come, First-Served (FCFS)
  - Run to completion in order of arrival
  - Pros: simple, low overhead, good for batch jobs
  - Cons: short jobs can stuck behind the long ones
- Round-Robin (RR)
  - FCFS with preemption. Cycle after a fixed time quantum
  - Pros: better interactivity (optimize response time)
  - Cons: performance is dependent on the quantum size
- Shortest Job First (SJF)/ Shorted Remaining Time First (SRTF)
  - Shorted job (or shortest remaining job) first
  - Pros: optimal average waiting time (turn-around time)
  - Cons: you need to know the future, long jobs can be starved by short jobs

# Agenda

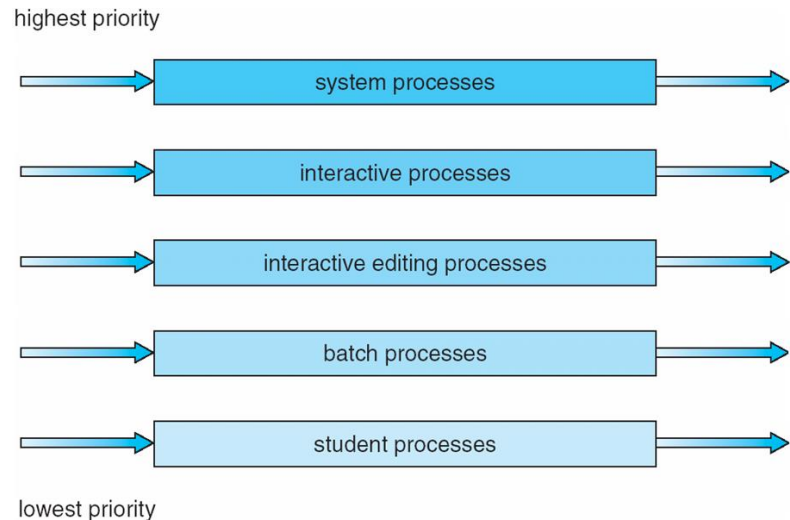
- **Multi-level queue scheduling**
- **Fair scheduling**
- Real-time scheduling
- Multicore scheduling

# Multiple Scheduling Goals

- Optimize for interactive applications
  - Round-robin
- Optimize for batch jobs
  - FCFS
- Can we do both?

# Multi-level Queue

- Ready queue is partitioned into separate queues
  - Foreground: interactive jobs
  - Background: batch jobs
- Each queue has its own scheduling algorithm
  - Foreground : RR
  - Background: FCFS
- Between the queue?




# Multi-level Queue Scheduling

- Scheduling between the queues
  - Fixed priority
    - Foreground first; schedule background only when no tasks in foreground
    - Possible starvation
  - Time slicing
    - Assign fraction of CPU time for each queue
    - 80% time for foreground; 20% time for background


# Multi-level Feedback Queue

- Each queue has a priority
- Tasks migrate across queues
  - Each job starts at the highest priority queue
  - If it uses up an entire quantum, drop one-level
  - If it finishes early, move up one-level (or stay at top)
- Benefits
  - Interactive jobs stay at high priority queues
  - Batch jobs will be at the low priority queue
  - **Automatically!**



# Example of Multilevel Feedback Queues

time = 0



■ Priority 0 (time slice = 1):

■ Priority 1 (time slice = 2):


■ Priority 2 (time slice = 4):






# Example of Multilevel Feedback Queues

■ Priority 0 (time slice = 1):



■ Priority 1 (time slice = 2):




■ Priority 2 (time slice = 4):




# Example of Multilevel Feedback Queues

time = 2

■ Priority 0 (time slice = 1):



■ Priority 1 (time slice = 2):



■ Priority 2 (time slice = 4):




# Example of Multilevel Feedback Queues

time = 3

■ Priority 0 (time slice = 1):

■ Priority 1 (time slice = 2):



■ Priority 2 (time slice = 4):




# Example of Multilevel Feedback Queues

time = 3

■ Priority 0 (time slice = 1):

■ Priority 1 (time slice = 2):



■ Priority 2 (time slice = 4):


*Suppose A is blocked on I/O*




# Example of Multilevel Feedback Queues

time = 3

■ Priority 0 (time slice = 1):



■ Priority 1 (time slice = 2):



■ Priority 2 (time slice = 4):


*Suppose A is blocked on I/O*




# Example of Multilevel Feedback Queues

time = 5

■ Priority 0 (time slice = 1):



■ Priority 1 (time slice = 2):



■ Priority 2 (time slice = 4):

*Suppose A is returned from I/O*




# Example of Multilevel Feedback Queues

time = 6

■ Priority 0 (time slice = 1):

■ Priority 1 (time slice = 2):



■ Priority 2 (time slice = 4):




# Example of Multilevel Feedback Queues

time = 8

■ Priority 0 (time slice = 1):

■ Priority 1 (time slice = 2):

■ Priority 2 (time slice = 4):






# Example of Multilevel Feedback Queues

time = 9

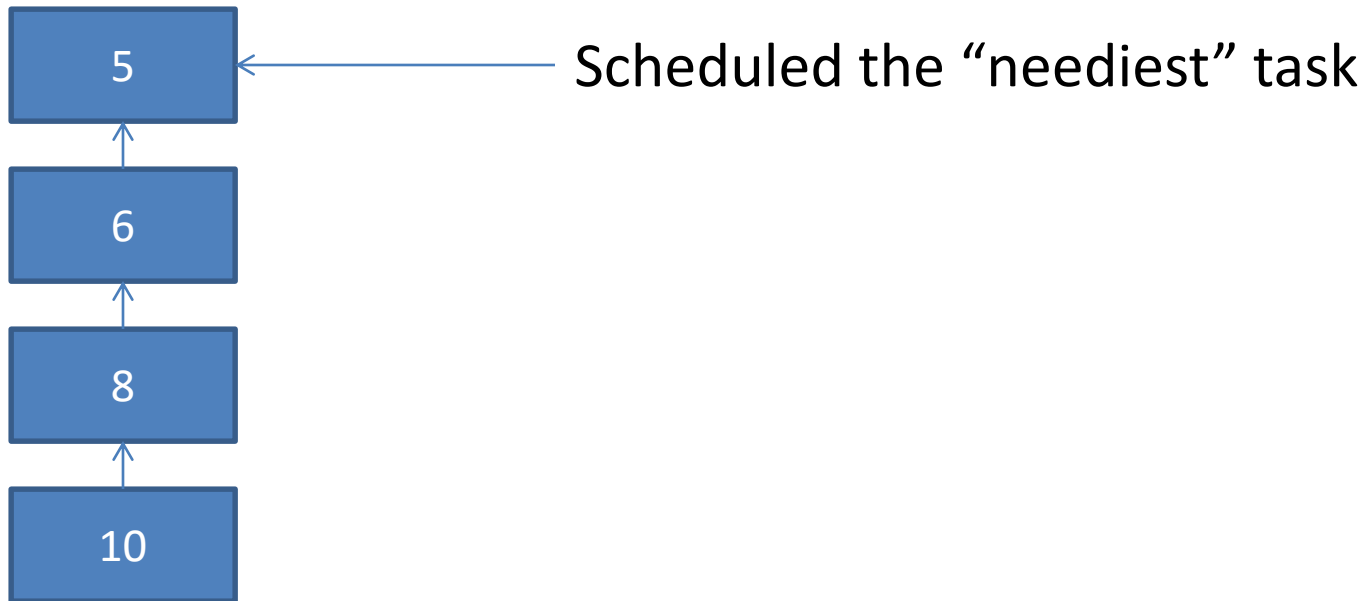
- Priority 0 (time slice = 1):
- Priority 1 (time slice = 2):
- Priority 2 (time slice = 4):



# Completely Fair Scheduler (CFS)


- Linux default scheduler, focusing on **fairness**
- Each task owns a fraction of CPU time share
  - E.g.,) A=10%, B=30%, C=60%
- Scheduling algorithm
  - Each task maintains its virtual runtime
    - Virtual runtime = executed time (x weight)
  - Pick the task with the **smallest virtual runtime**
    - Tasks are sorted according to their virtual times
  - Time slice varies depending on the #of tasks
    - Slice = target\_latency / #of tasks

# CFS Example



- Tasks are sorted according to their virtual times

# CFS Example



On a next scheduler event  
**re-sort** the list

But list is inefficient.

- Tasks are sorted according to their virtual times

# Red-black Tree





Figure source: M. Tim Jones, [“Inside the Linux 2.6 Completely Fair Scheduler”](#), IBM developerWorks

- Self-balancing binary search tree
- Insert:  $O(\log N)$ , Remove:  $O(1)$

# Weighed Fair Sharing: Example



Weights: gcc =  $2/3$ , bigsim =  $1/3$

X-axis: mcu (tick), Y-axis: virtual time

Fair in the long run

# Recap


- Multi-level queue scheduling
  - Multiple scheduling policies
    - RR for interactive tasks
    - FCFS for batch tasks
  - Multi-level feedback queue scheduling
    - Tasks automatically migrate among the queues
- CFS
  - Fair sharing of CPU time
  - Based on (virtual) runtime of each task

# Some Edge Cases

- How to set the virtual time of a new task?
  - Can't set as zero. Why?
  - System virtual time (SVT)
    - The minimum virtual time among all active tasks
    - `cfs_rq->min_vruntime`
  - The new task can “catch-up” tasks by setting its virtual time with SVT



# Weighed Fair Sharing: Example 2



Weights: gcc =  $2/3$ , bigsim =  $1/3$


X-axis: mcu (tick), Y-axis: virtual time

gcc slept 15 mcu

# Real-Time Systems

- The correctness of the system depends not only on the logical result of the computation but also on the time at which the results are produced.
- **A correct value at the wrong time is a fault.**
- Processes attempt to control or react to events that take place in the outside world
- These events occur in “real time” and tasks must be able to keep up with them
- Processes are associated with timing constraints (deadlines)

# Real-Time Spectrum

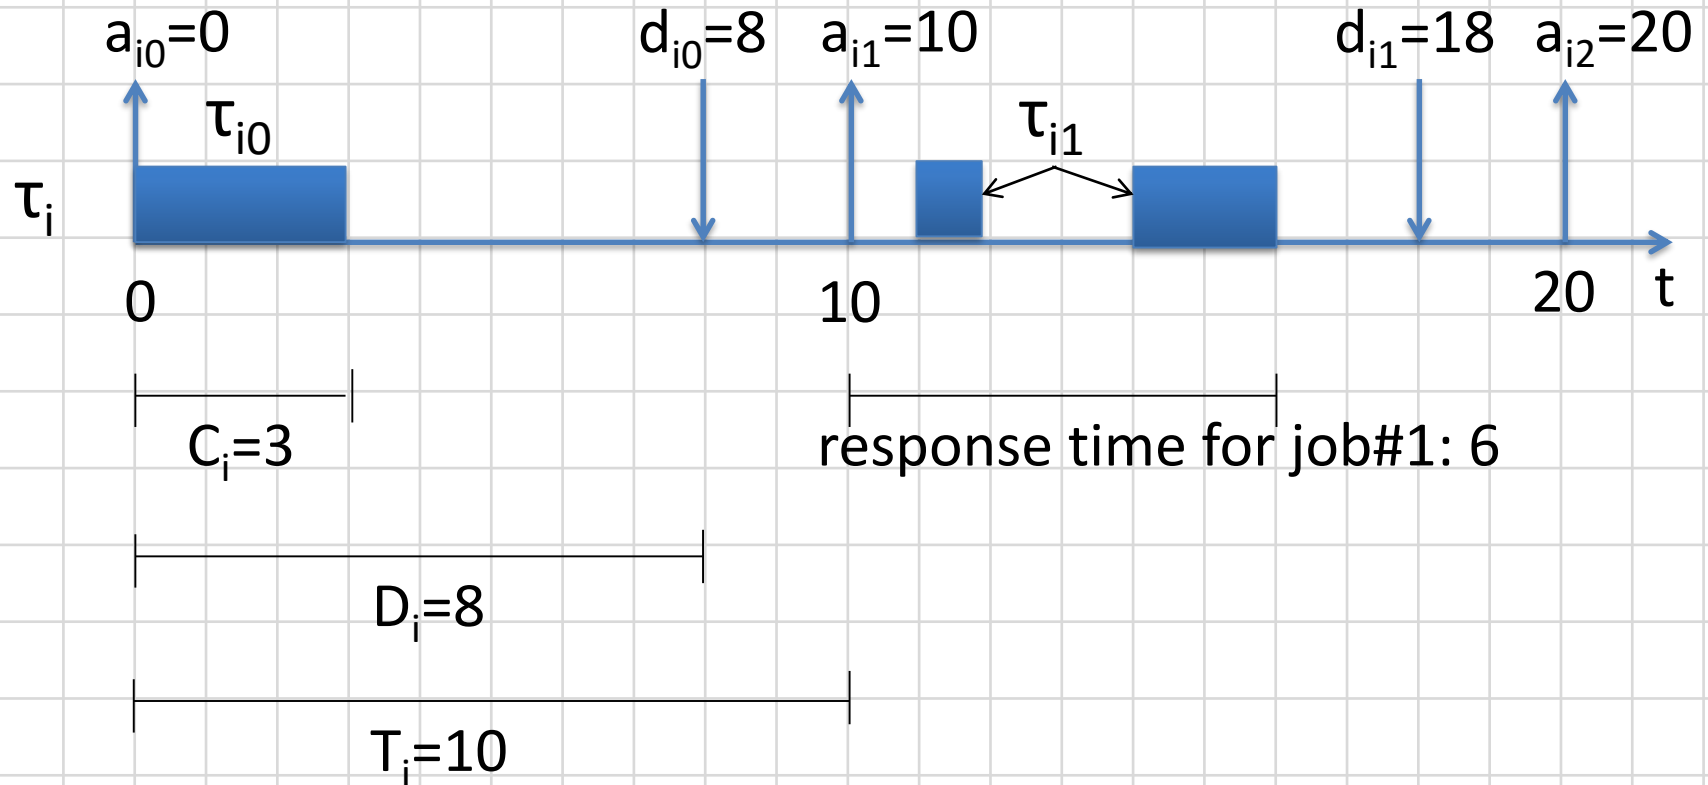


# Real-Time Scheduling

- Goal: meet the deadlines of important tasks
  - **Soft** deadline: game, video decoding, ...
  - **Hard** deadline: engine control, anti-lock break (ABS)
    - 100 ECUs (processors) in BMW i3 [\*]
- Priority scheduling
  - A high priority task preempts lower priority tasks
  - Static priority scheduling
  - Dynamic priority scheduling


[\*] Robert Leibinger, “Software Architectures for Advanced Driver Assistance Systems (ADAS)”, OSPERT’15 keynote

# Periodic Task Model




# Rate Monotonic (RM)

- Priority is assigned based on **periods**
  - Shorter period -> higher priority
  - Longer period -> lower priority
- Optimal static-priority scheduling




# Earliest Deadline First (EDF)

- Priority is assigned based on deadline
  - Shorter deadline  $\rightarrow$  higher priority
  - Longer deadline  $\rightarrow$  lower priority
- Optimal dynamic priority scheduling



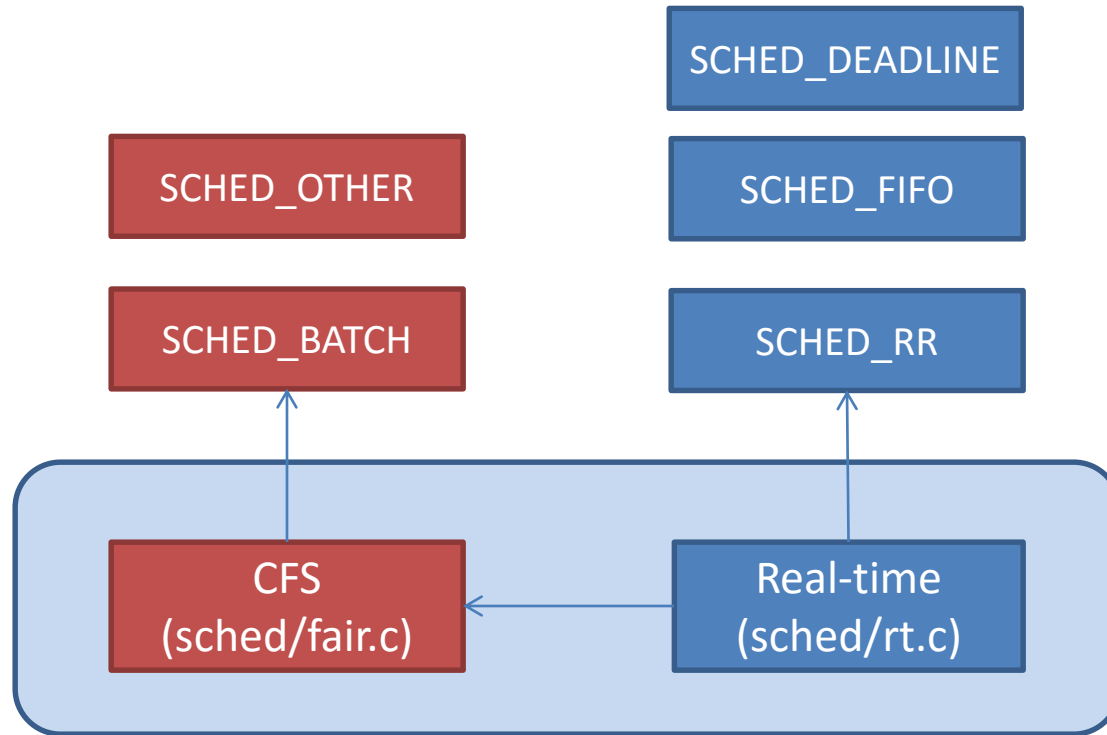
# Linux Scheduling Framework



- First, schedule real-time tasks
  - Real-time schedulers: (1) Priority based, (2) deadline based
- Then schedule normal tasks
  - Completely Fair Scheduler (CFS)
- Two-level queue scheduling
  - Between queues?



# Linux Scheduling Framework

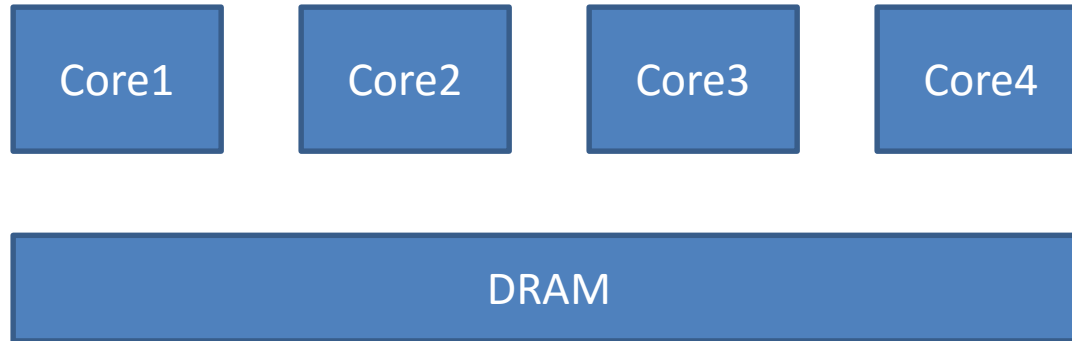


- Completely Fair Scheduler (CFS)
  - SCHED\_OTHER, SCHED\_BATCH
- Real-time Schedulers
  - SCHED\_DEADLINE, SCHED\_FIFO, SCHED\_RR

# Real-Time Schedulers in Linux


- SCHED\_FIFO
  - Static priority scheduler
- SCHED\_RR
  - Same as SCHED\_FIFO except using RR for tasks with the same priority
- SCHED\_DEADLINE
  - EDF scheduler
  - Recently merged in the Linux mainline (v3.14)

# Multiprocessor Scheduling

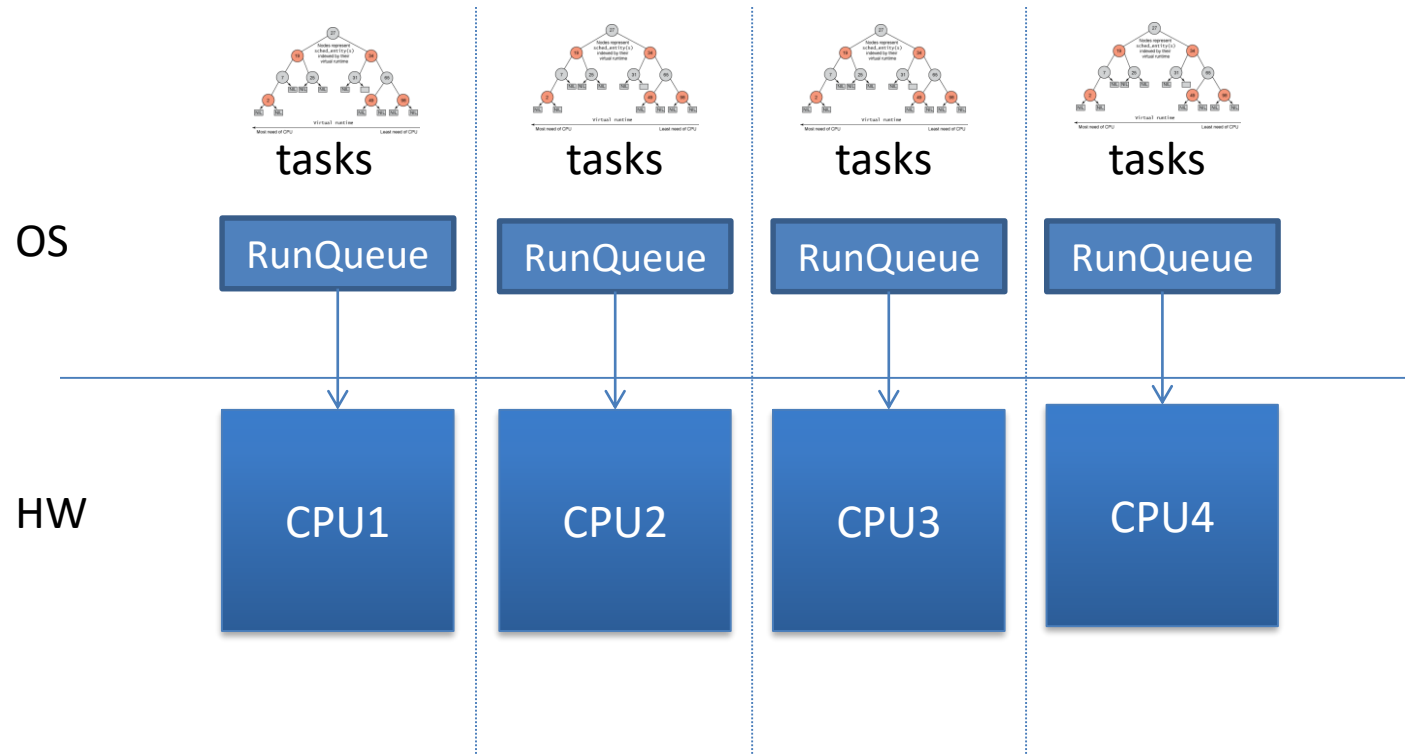


- How many scheduling queues are needed?
  - Global shared queue: all tasks are placed in a single shared queue (global scheduling)
  - Per-core queue: each core has its own scheduling queue (partitioned scheduling)

# Global Scheduling



# Partitioned Scheduling

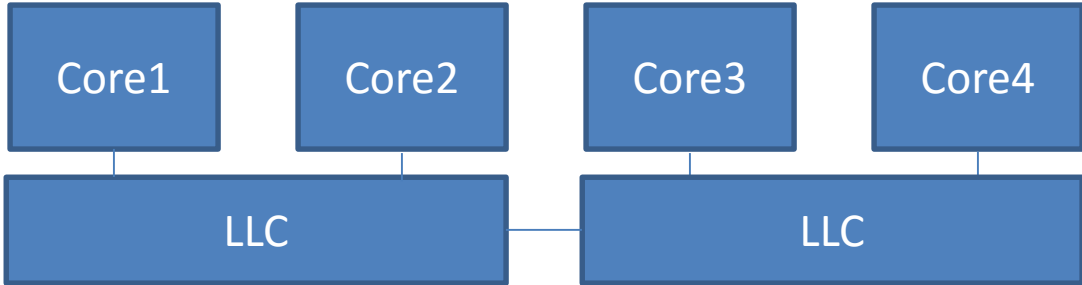


- Linux's basic design. Why?

# Load Balancing

- Undesirable situation
  - Core 1's queue: 40 tasks
  - Core 2's queue: 0 task
- Load balancing
  - Tries to balance load across all cores.
  - Not so simple, why?
    - Migration overhead: cache warmup

# Load Balancing

- More considerations
  - What if certain cores are more powerful than others?
    - E.g., ARM bigLITTLE (4 big cores, 4 small cores)
  - What if certain cores share caches while others don't?

```
graph TD; C1[Core1] --- L1[LLC]; C2[Core2] --- L1; C3[Core3] --- L2[LLC]; C4[Core4] --- L2; L1 --- L2
```
  - Which tasks to migrate?
    - Some tasks may compete for limited shared resources

# Summary

- Multi-level queue scheduling
  - Each queue has its own scheduler
  - Scheduling between the queues
- Fair scheduling (CFS)
  - Fairly allocate CPU time across all tasks
  - Pick the task with the smallest virtual time
  - Guarantee fairness and bounded response time
- Real-time scheduling
  - Static priority scheduling
  - Dynamic priority scheduling



# Summary

- Multicore scheduling
  - Global queue vs. per-core queue
    - Mostly per-core queue due to scalability
  - Load balancing
    - Balance load across all cores
    - Is complicated due to
      - Migration overhead
      - Shared hardware resources (cache, dram, etc)
      - Core architecture heterogeneity (big cores vs. small cores)
      - ...

# Acknowledgements

- Some slides are adopted from the notes of
  - Dr. Kulkarni at KU
  - Dr. Pellizzoni at Univ. of Waterloo
  - The book authors