

Memory Management

Disclaimer: some slides are adopted from book authors' slides with permission

Roadmap

- CPU management
 - Process, thread, synchronization, scheduling
- **Memory management**
 - **Virtual memory**
- Disk management
- Other topics

Memory Management

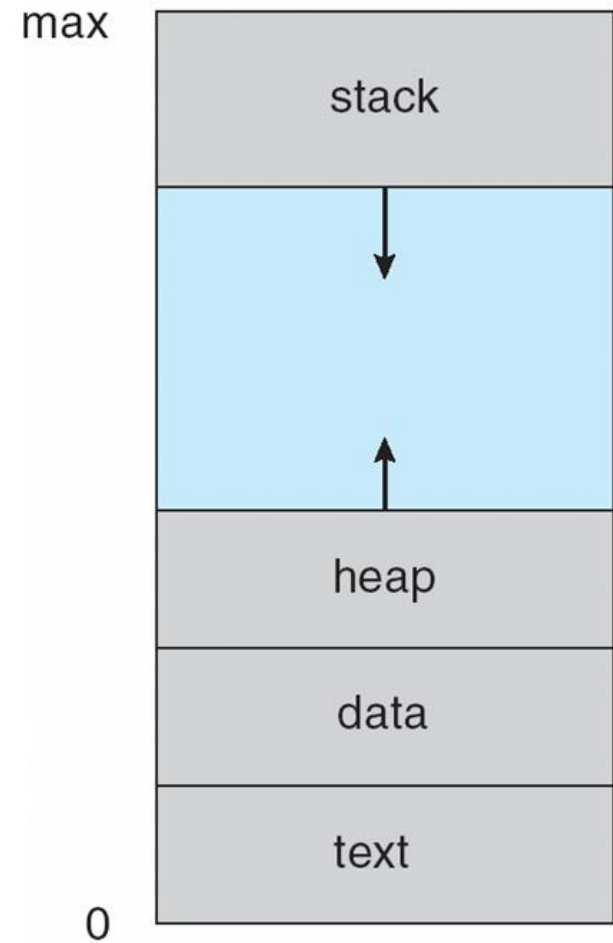
- Goals
 - Easy to use abstraction
 - Same virtual memory space for all processes
 - Isolation among processes
 - Don't corrupt each other
 - Efficient use of capacity limited physical memory
 - Don't waste memory

Concepts to Learn

- Virtual address translation
- Paging and TLB
- Page table management
- Swap

Virtual Memory (VM)

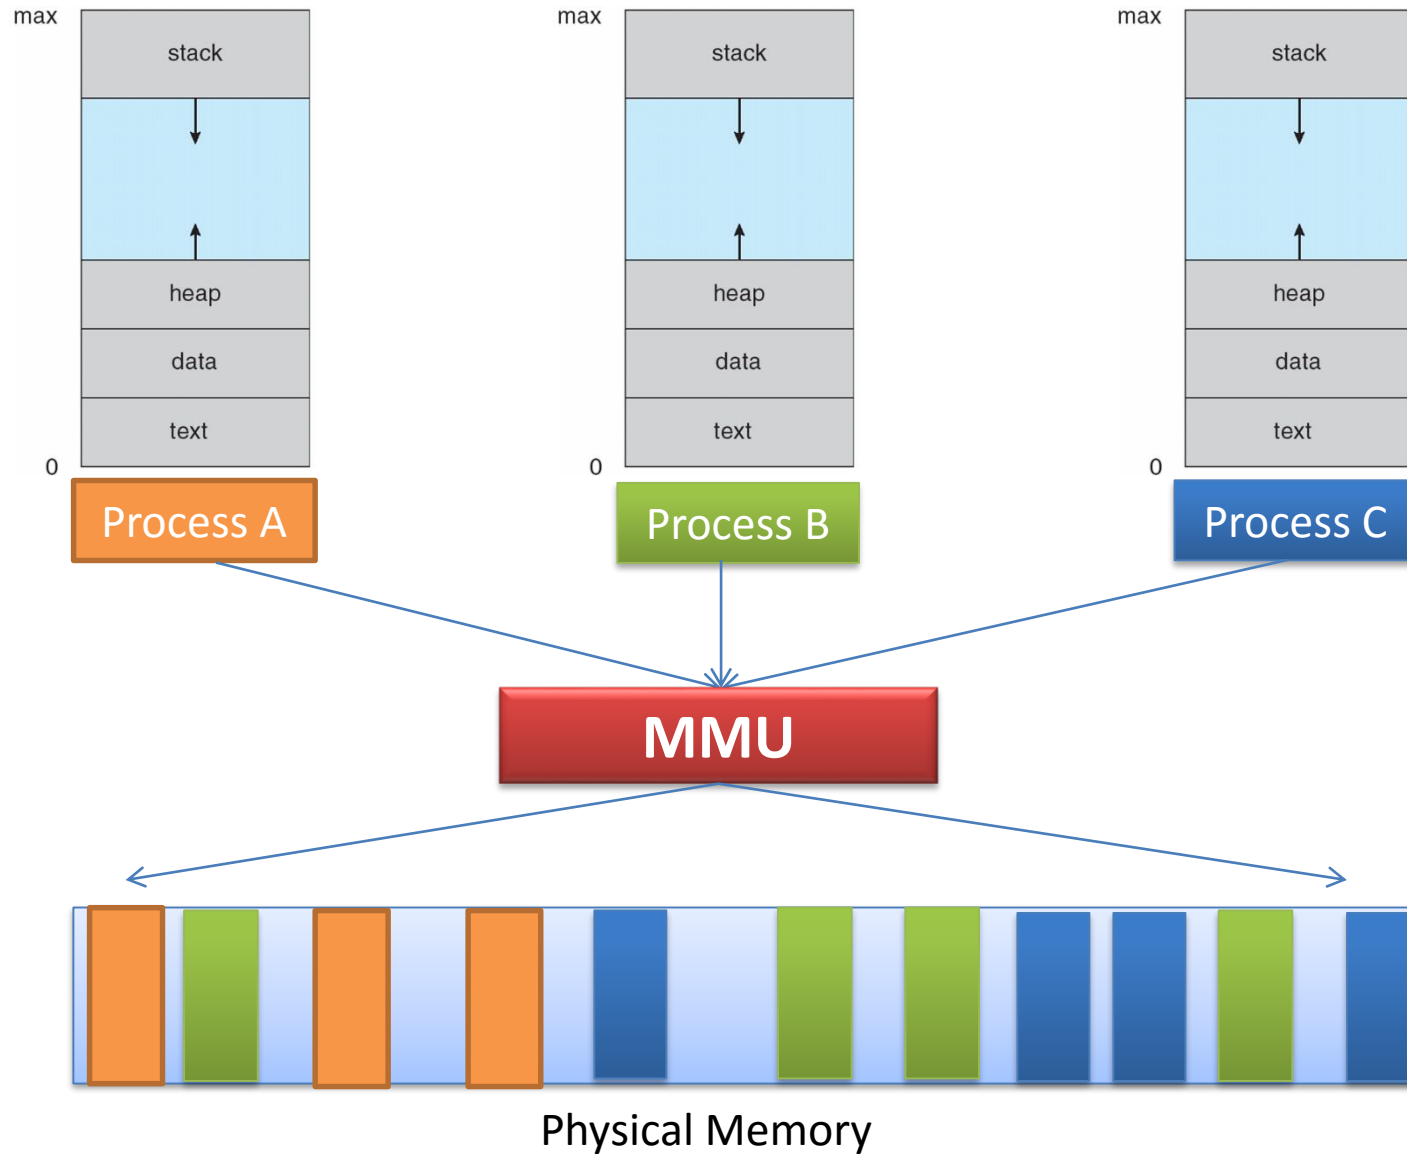
- Abstraction
 - 4GB linear address space for each process
- Reality
 - 1GB of actual physical memory shared with 20 other processes
- How?



Virtual Memory

- Hardware support
 - MMU (memory management unit)
 - TLB (translation lookaside buffer)
- OS support
 - Manage MMU (sometimes TLB)
 - Determine address mapping
- Alternatives
 - No VM: many real-time OS (RTOS) don't have VM

Virtual Address



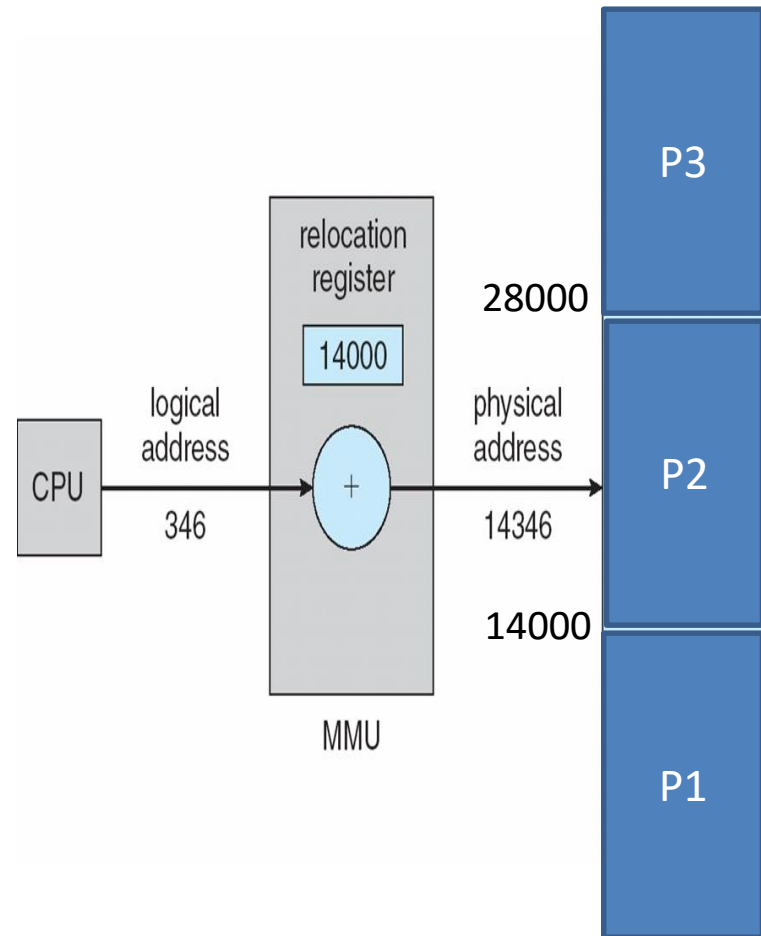
MMU

- Hardware unit that translates *virtual address* to *physical address*



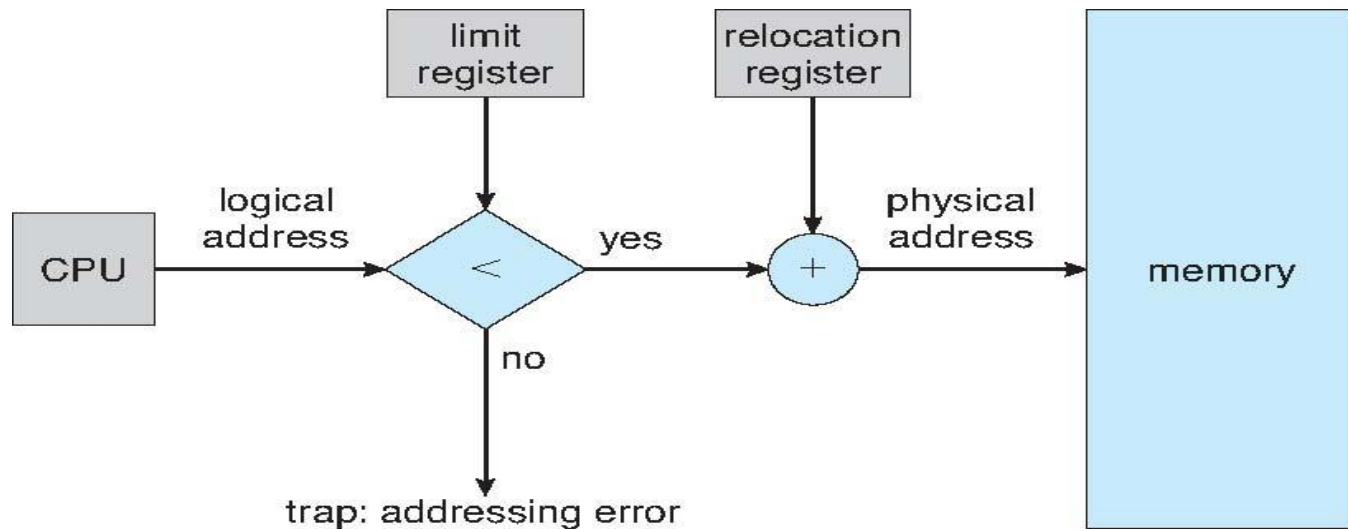
A Simple MMU

- BaseAddr: base register
- $Paddr = Vaddr + BaseAddr$
- Advantages
 - Fast
- Disadvantages
 - No protection
 - Wasteful




A Better MMU

- Base + Limit approach
 - If $Vaddr > limit$, then trap to report error
 - Else $Paddr = Vaddr + BaseAddr$




A Better MMU

- Base + Limit approach
 - If $Vaddr > limit$, then trap to report error
 - Else $Paddr = Vaddr + BaseAddr$
- Advantages
 - Support protection
 - Support variable size partitions
- Disadvantages
 - Fragmentation



Fragmentation

- External fragmentation
 - total available memory space exists to satisfy a request, but it is not contiguous

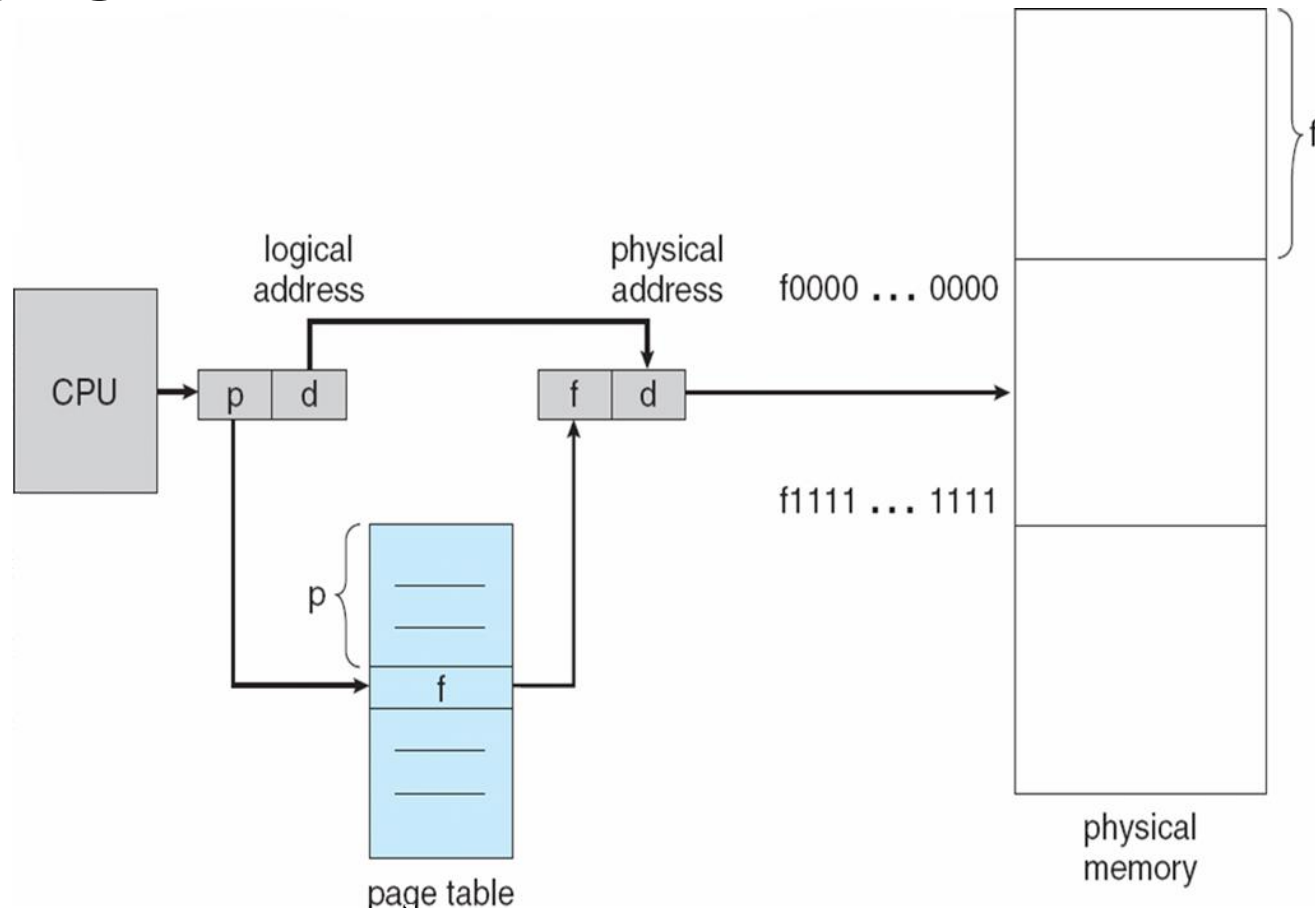


Modern MMU

- Paging approach
 - Divide physical memory into fixed-sized blocks called frames (e.g., 4KB each)
 - Divide logical memory into blocks of the same size called pages (page size = frame size)
 - Pages are mapped onto frames via a table → page table

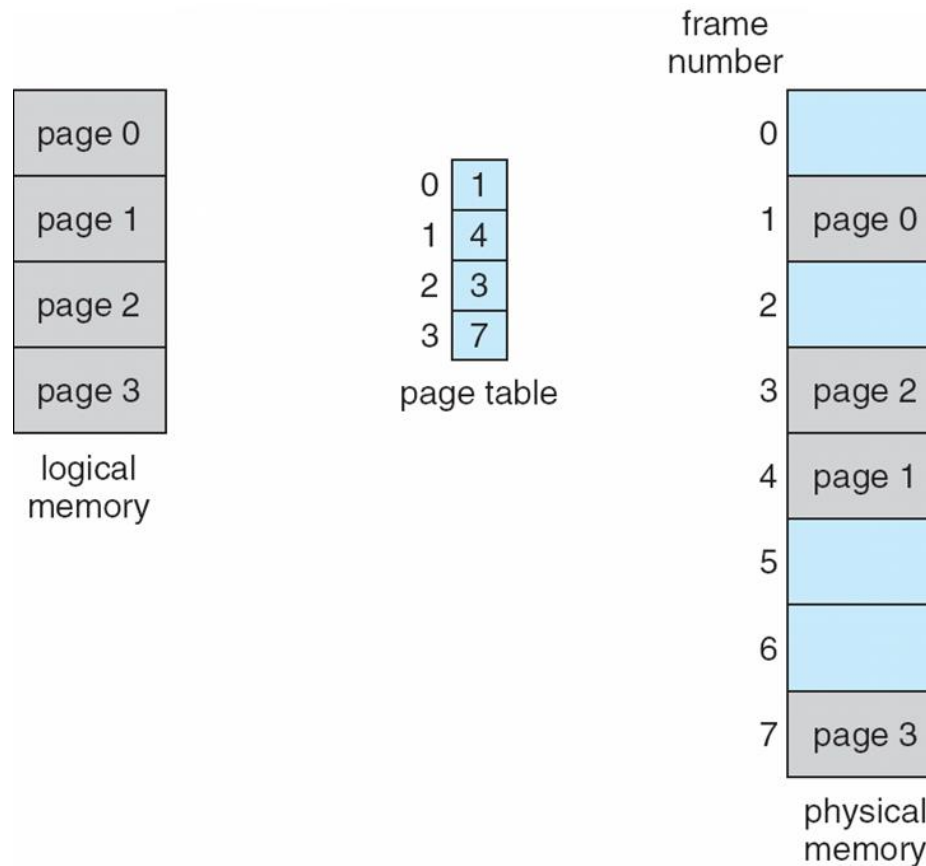
Modern MMU

- Paging hardware



Modern MMU

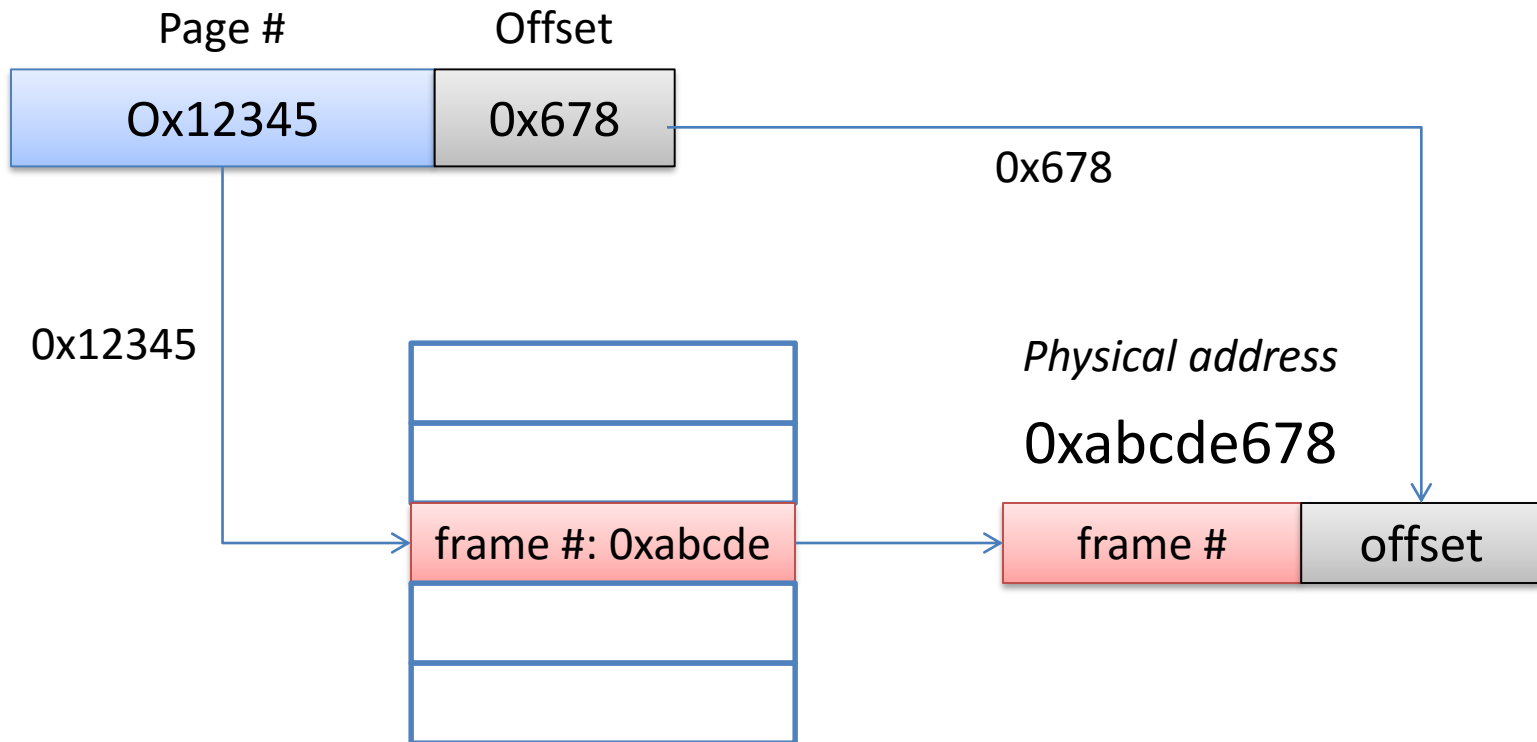
- Memory view



Virtual Address Translation

Virtual address

0x12345678



Advantages of Paging

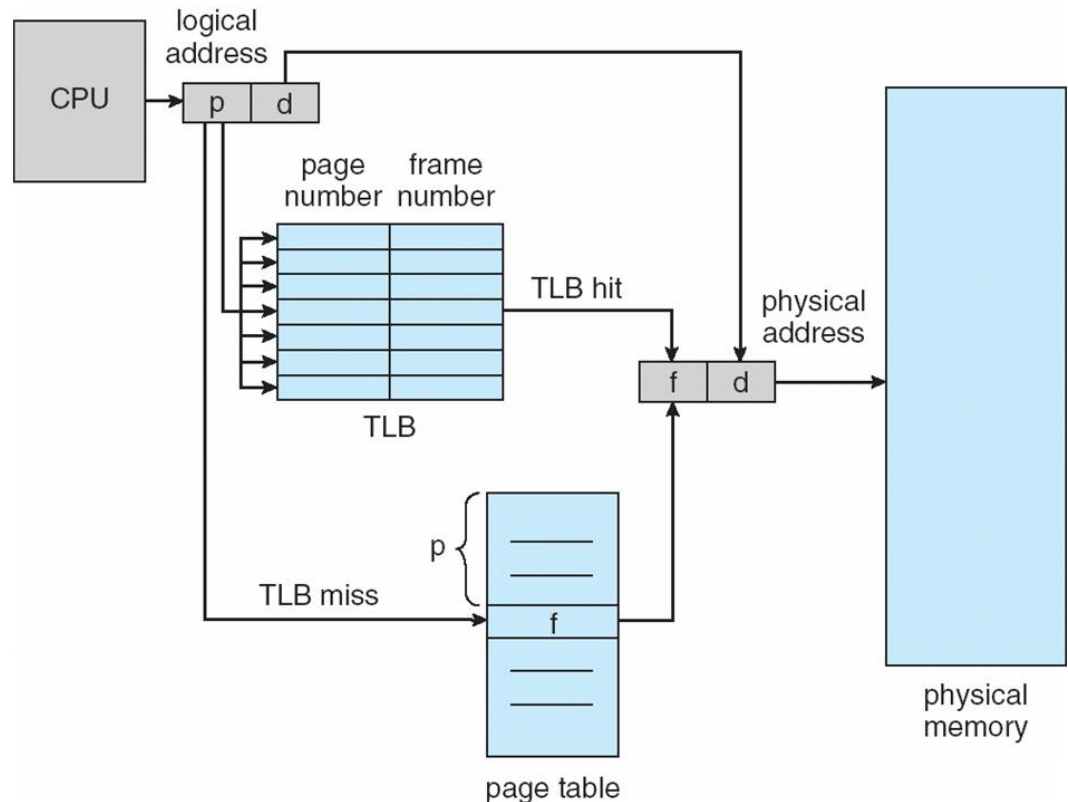
- No external fragmentation
 - Efficient use of memory
 - Internal fragmentation (waste within a page) still exists

Issues of Paging

- Translation speed
 - Each load/store instruction requires a translation
 - Table is stored in **memory**
 - Memory is **slow** to access
 - ~100 CPU cycles to access DRAM

Translation Lookaside Buffer (TLB)

- **Cache** frequent address translations
 - So that CPU don't need to access the page table all the time
 - Much faster



Issues of Paging

- Page size
 - Small: minimize space waste, requires a large table
 - Big: can waste lots of space, the table size is small
 - Typical size: 4KB
 - How many pages are needed for 4GB (32bit)?
 - $4\text{GB}/4\text{KB} = 1\text{M}$ pages
 - What is the required page table size?
 - assume 1 page table entry (PTE) is 4bytes
 - $1\text{M} * 4\text{bytes} = 4\text{MB}$
 - Btw, this is for each process. What if you have 100 processes? Or what if you have a 64bit address?

Paging

- Advantages
 - No external fragmentation
- Two main Issues
 - Translation speed can be slow
 - TLB
 - Table size is big