Memory Management

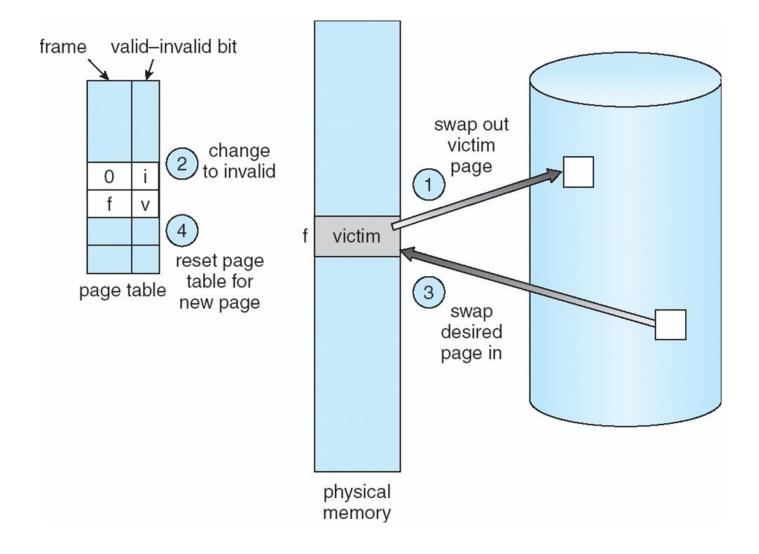


Page Replacement Procedure

- On a page fault
 - Step 1: allocate a free page frame
 - If there's a free frame, use it
 - If there's no free frame, choose a victim frame and evict it to disk (if necessary) → swap-out
 - Step 2: bring the stored page on disk (if necessary)
 - Step 3: update the PTE (mapping and valid bit)
 - Step 4: restart the instruction



Page Replacement Procedure





Page Replacement Policies

- FIFO (First In, First Out)
 - Evict the oldest page first.
 - Pros: simple, fair
 - Cons: can throw out frequently used pages
- Optimal
 - Evict the page that will not be used for the longest period
 - Pros: best performance
 - Cons: you need to know the future



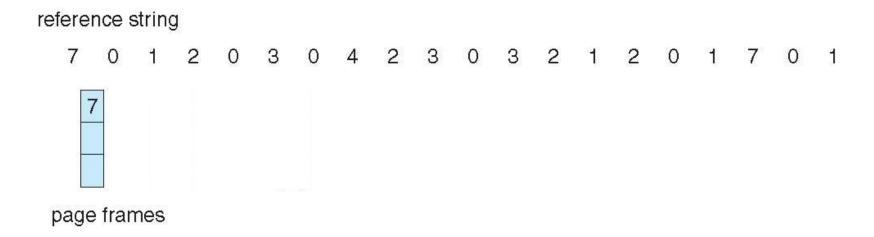
Page Replacement Policies

Random

- Randomly choose a page
- Pros: simple. TLB commonly uses this method
- Cons: unpredictable
- LRU (Least Recently Used)
 - Look at the past history, choose the one that has not been used for the longest period
 - Pros: good performance
 - Cons: complex, requires h/w support

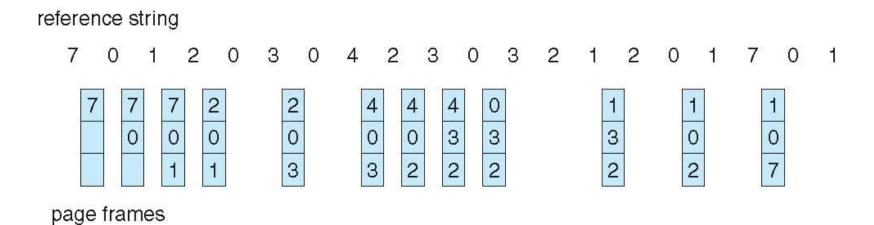


LRU Example





LRU Example





Implementing LRU

- Ideal solutions
 - Timestamp
 - Record access time of each page, and pick the page with the oldest timestamp
 - List
 - Keep a list of pages ordered by the time of reference
 - Head: recently used page, tail: least recently used page
 - Problems: very expensive (time & space & cost) to implement



Page Table Entry (PTE)

PTE format (architecture specific)



- Valid bit (V): whether the page is in memory
- Modify bit (M): whether the page is modified
- Reference bit (R): whether the page is accessed
- Protection bits(P): readable, writable, executable

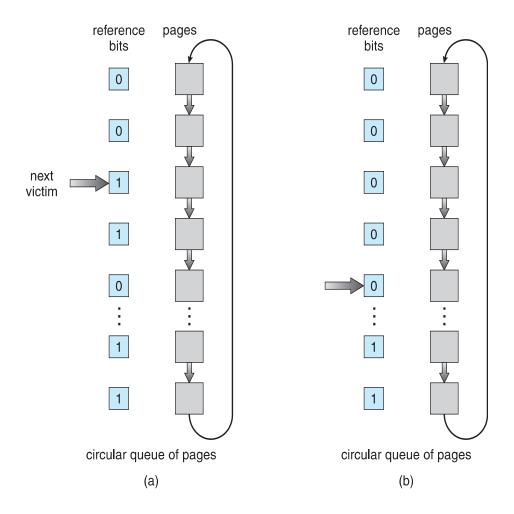


Implementing LRU: Approximation

- Second chance algorithm (or clock algorithm)
 - Replace an old page, not the oldest page
 - Use 'reference bit' set by the MMU
- Algorithm details
 - Arrange physical page frames in circle with a pointer
 - On each page fault
 - Step 1: advance the pointer by one
 - Step 2: check the reference bit of the page:
 - $1 \rightarrow$ Used recently. Clear the bit and go to Step 1
 - $0 \rightarrow$ Not used recently. Selected victim. End.



Second Chance Algorithm





Implementing LRU: Approximation

- N chance algorithm
 - OS keeps a counter per page
 - On a page fault
 - Step 1: advance the pointer by one
 - Step 2: check the reference bit of the page: check the reference bit
 - $1 \rightarrow$ reference=0; counter=0
 - 0 → counter++; if counter =N then found victim, otherwise repeat Step 1.
 - Large N → better approximation to LRU, but costly
 - Small N → more efficient but poor LRU approximation



Performance of Demand Paging

- Three major activities
 - Service the interrupt hundreds of cpu cycles
 - Read/write the page from/to disk lots of time
 - Restart the process again just a small amount of time
- Page Fault Rate $0 \le p \le 1$
 - if p = 0 no page faults
 - if p = 1, every reference is a fault
- Effective Access Time (EAT)

```
EAT = (1 - p) x memory access
+ p (page fault overhead
+ swap page out
+ swap page in )
```



Performance of Demand Paging

- Memory access time = 200 nanoseconds
- Average page-fault service time = 8 milliseconds
- EAT = $(1 p) \times 200 + p$ (8 milliseconds) = $(1 - p) \times 200 + p \times 8,000,000$ = $200 + p \times 7,999,800$
- If one access out of 1,000 causes a page fault, then
 EAT = 8.2 microseconds. → This is a slowdown by a factor of 40!!
- If want performance degradation < 10 percent
 - 220 > 200 + 7,999,800 x p20 > 7,999,800 x p
 - p < .0000025
 - < one page fault in every 400,000 memory accesses



Thrashing

- A process is busy swapping pages in and out
 - Don't make much progress
 - Happens when a process do not have "enough" pages in memory
 - Very high page fault rate
 - Low CPU utilization (why?)
 - CPU utilization based admission control may bring more programs to increase the utilization → more page faults



Thrashing

