

Memory Management

Disclaimer: some slides are adopted from book authors' slides with permission


Recap: Performance of Demand Paging

- Memory access time = 200 nanoseconds
- Average page-fault service time = 8 milliseconds
- $EAT = (1 - p) \times 200 + p (8 \text{ milliseconds})$
 $= (1 - p) \times 200 + p \times 8,000,000$
 $= 200 + p \times 7,999,800$
- If one access out of 1,000 causes a page fault, then
EAT = 8.2 microseconds. \rightarrow This is a slowdown by a factor of 40!!
- If want performance degradation < 10 percent
 - $220 > 200 + 7,999,800 \times p$
 $20 > 7,999,800 \times p$
 - $p < .0000025$
 - < one page fault in every 400,000 memory accesses

Thrashing

- A process is busy swapping pages in and out
 - Don't make much progress
 - Happens when a process do not have “enough” pages in memory
 - Very high page fault rate
 - Low CPU utilization (why?)
 - CPU utilization based admission control may bring more programs to increase the utilization → more page faults

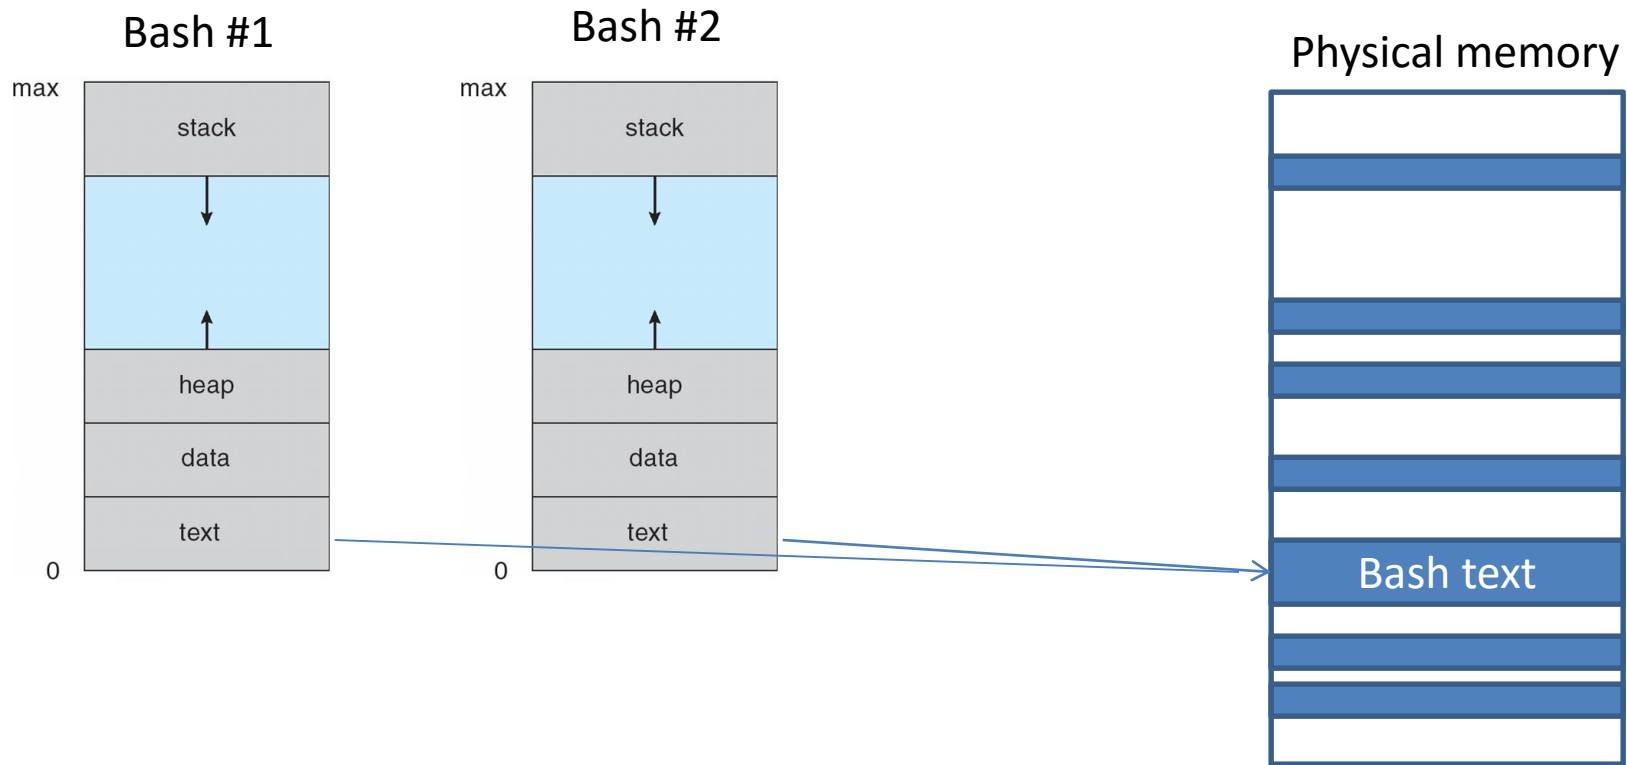
Thrashing



Concepts to Learn

- Memory-mapped I/O
- Copy-on-Write (COW)
- Memory allocator


Recap: Program Binary Sharing



- Multiple instances of the same program
 - E.g., 10 bash shells

Memory Mapped I/O

- Idea: map a file on disk onto the memory space



Memory Mapped I/O


- Benefits: you don't need to use read()/write() system calls, just directly access data in the file via memory instructions
- How it works?
 - Just like demand paging of an executable file
 - What about writes?
 - Mark the modified (M) bit in the PTE
 - Write back the modified pages back to the original file

Copy-on-Write (COW)

- Fork() creates a copy of a parent process
 - Copy the entire pages on new page frames?
 - If the parent uses 1GB memory, then a fork() call would take a while
 - Then, suppose you immediately call exec(). Was it of any use to copy the 1GB of parent process's memory?


Copy-on-Write

- Better way: copy the page table of the parent
 - Page table is much smaller (so copy is faster)
 - Both parent and child point to the exactly same physical page frames



Copy-on-Write

- What happens when the parent/child reads?
- What happens when the parent/child writes?
 - Trouble!!!



Page Table Entry (PTE)

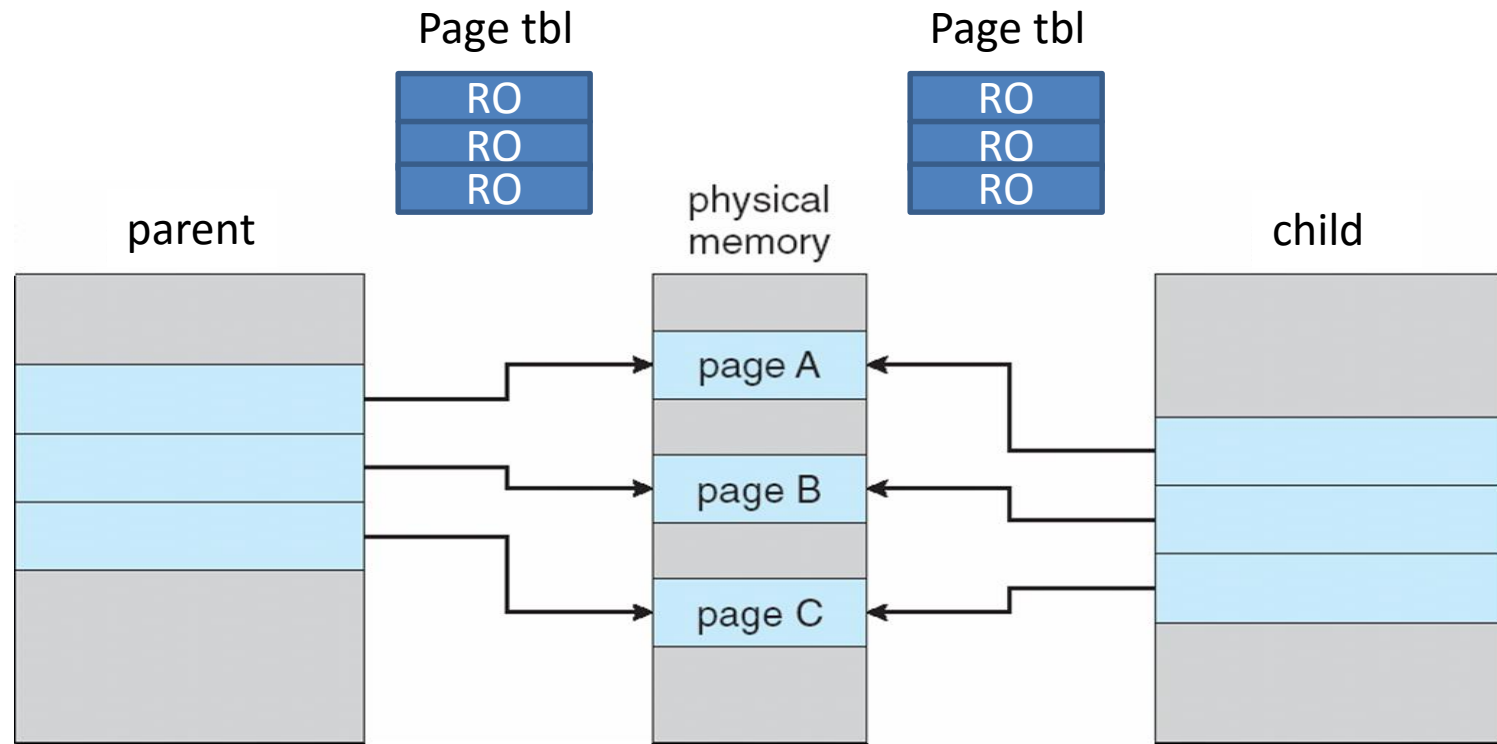
- PTE format (architecture specific)



- Valid bit (V): whether the page is in memory
- Modify bit (M): whether the page is modified
- Reference bit (R): whether the page is accessed
- **Protection bits(P): readable, writable, executable**

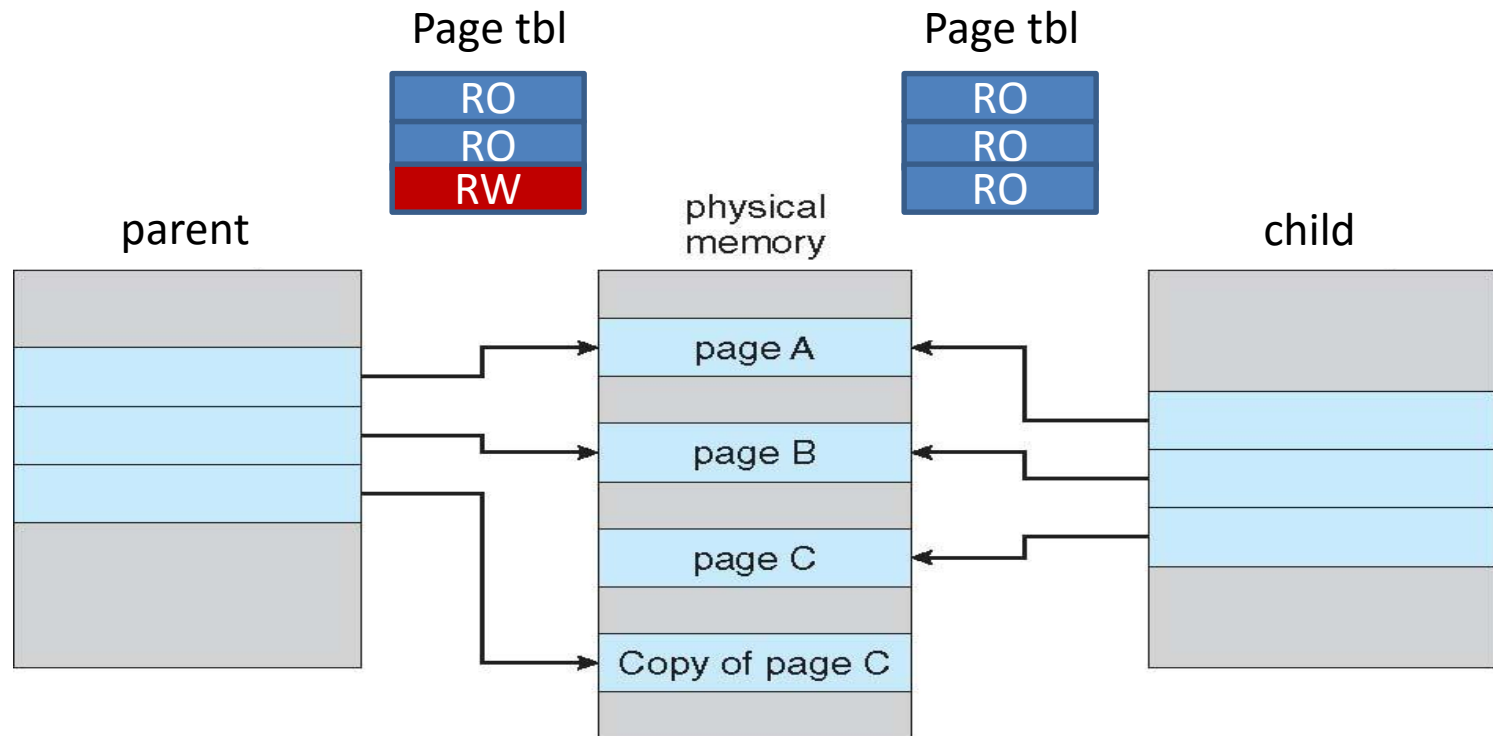
Copy-on-Write

- All pages are marked as read-only




Copy-on-Write

- Up on a write, a page fault occurs and the OS copies the page on a new frame and maps to it with R/W protection setting



User-level Memory Allocation

- When a process actually allocate a memory from the kernel?
 - On a page fault
 - Allocate a page (e.g., 4KB)
- What does **malloc()** do?
 - Manage a process's heap
 - Variable size objects in heap




Kernel-level Memory Allocation

- Page-level allocator
 - Page frame allocation/free (fixed size)
 - Users: page fault handler, kernel-memory allocator
- Kernel-memory allocator (KMA)
 - Typical kernel object size \ll page size
 - File descriptor, inode, task_struct, ...
 - KMA \leftarrow kernel-level malloc
 - In Linux: buddy allocator, SLAB


Buddy Allocator

- Allocate physically contiguous pages
 - Satisfies requests in units sized as power of 2
 - Request rounded up to next highest power of 2
 - When smaller allocation needed than is available, current chunk split into two buddies of next-lower power of 2
 - Quickly expand/shrink across the lists



Buddy Allocator

- Example
 - Assume 256KB chunk available, kernel requests 21KB



Buddy Allocator

- Example
 - Free A

32 A	32 Free	64 Free	128 Free
---------	------------	------------	-------------

32 Free	32 Free	64 Free	128 Free
------------	------------	------------	-------------

64 Free	64 Free	128 Free
------------	------------	-------------

128 Free	128 Free
-------------	-------------

256 Free


Virtual Memory Summary

- MMU and address translation
- Paging
- Demand paging
- Copy-on-write
- Page replacement

Quiz: Address Translation



Virtual address format (24bits)



Page table entry (8bit)

Vaddr: 0x0703FE

Paddr: 0x3FE

Vaddr: 0x072370

Paddr: ???

Vaddr: 0x082370

Paddr: ???


Page-table base address = 0x100

Addr	+0	+1	+2	+3	+4	+5	+6	+7	+8	+A	+B	+C	+D	+E	+F
0x000				31											
0x010															
0x020				41											
..															
0x100	00	01						01	00			01			
..															
0x200															

Quiz: Address Translation



Virtual address format (24bits)



Page table entry (8bit)

Vaddr: 0x0703FE

Vaddr: 0x072370

Vaddr: 0x082370

Paddr: 0x3FE

Paddr: **0x470**

Paddr: **invalid**

Page-table base address = 0x100

Addr	+0	+1	+2	+3	+4	+5	+6	+7	+8	+A	+B	+C	+D	+E	+F
0x000				31											
0x010															
0x020				41											
..															
0x100	00	01						01	00			01			
..															
0x200															