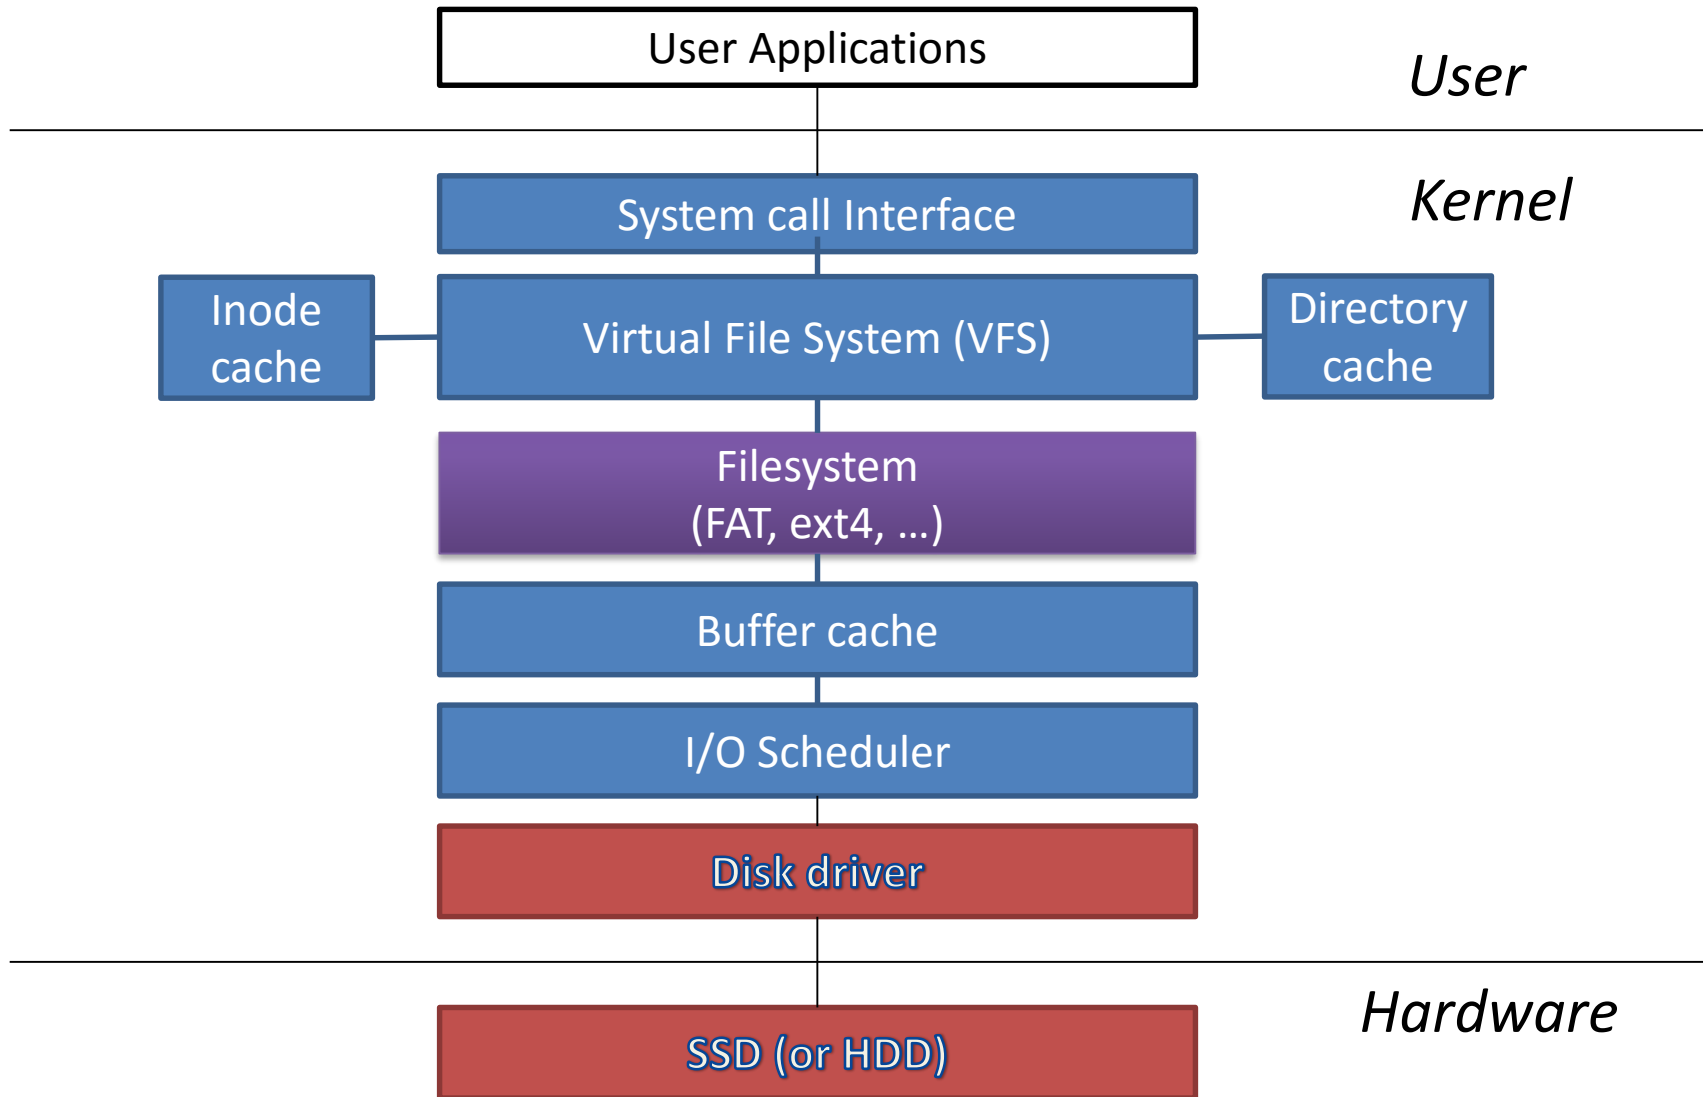


Filesystem

Disclaimer: some slides are adopted from book authors' slides with permission

Storage Subsystem in Linux OS



Filesystem

- Definition
 - An OS layer that provides **file** and **directory** abstractions on disks
- File
 - User's view: a collection of **bytes** (non-volatile)
 - OS's view: a collection of **blocks**
 - A block is a logical transfer unit of the kernel (typically block size \geq sector size)

Filesystem


- File types
 - Executables, DLLs, text, word,
 - Filesystems mostly don't care
- File attributes (metadata)
 - Name, location, size, protection, ...
- File operations
 - Create, read, write, delete, seek, truncate, ...

How to Design a Filesystem?

- What to do?
 - Map disk blocks to each file
 - Need to track free disk blocks
 - Need to organize files into directories
- Requirements
 - Should not waste space
 - Should be fast

Access Pattern

- Sequential access
 - E.g.,) read next 1000 bytes



- Random access
 - E.g,) Read 10 bytes at the offset 300
- Remember that random access is especially slow in HDD.

File Usage Patterns

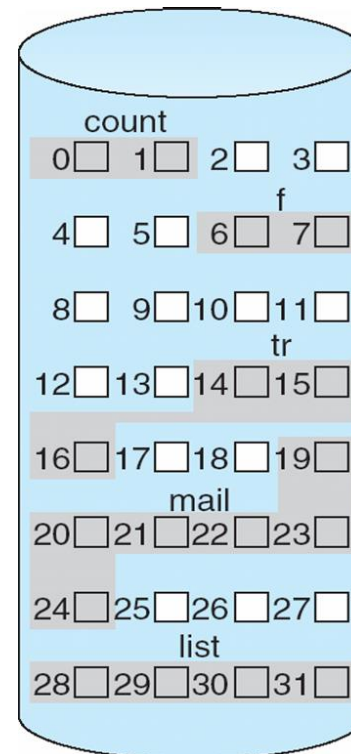
- Most files are small
 - .c, .h, .txt, .log, .ico, ...
 - Also more frequently accessed
 - If the block size is too big, It wastes space (why?)
- Large files use most of the space
 - .avi, .mp3, .jpg,
 - If the block size is too small, mapping information can be huge (performance and space overhead)

Disk Allocation

- How to map disk blocks to files?
 - Each file may have very different size
 - The size of a file may change over time (grow or shrink)
- Disk allocation methods
 - Continuous allocation
 - Linked allocation
 - Indexed allocation

Continuous Allocation

- Use continuous ranges of blocks
 - Users declare the size of a file in advance
 - File header: first block #, #of blocks
 - Similar to malloc()
- Pros
 - Fast sequential access
 - easy random access
- Cons
 - External fragmentation
 - difficult to increase

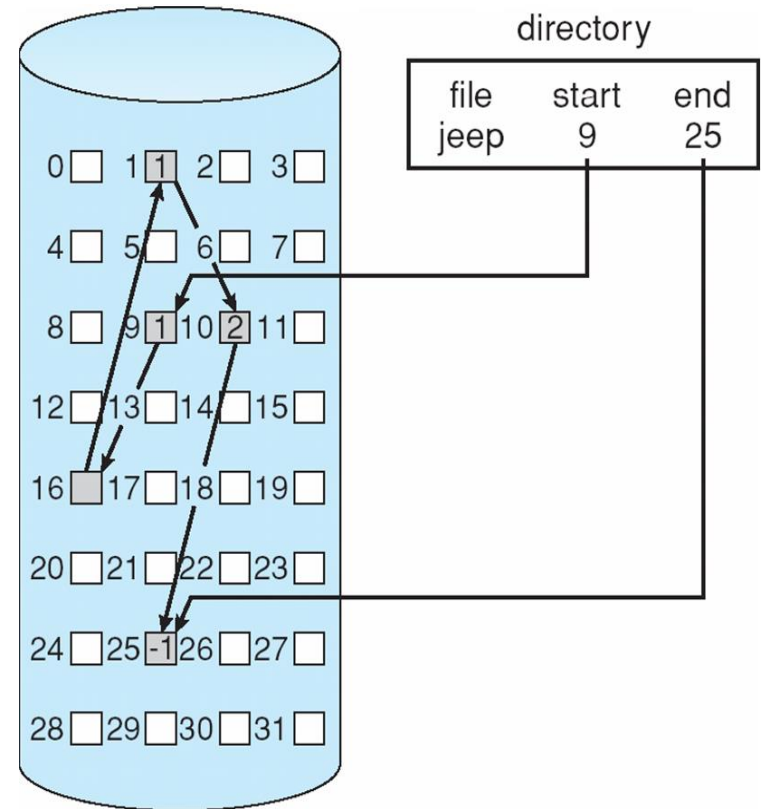


directory

file	start	length
count	0	2
tr	14	3
mail	19	6
list	28	4
f	6	2


Linked-List Allocation

- Each block holds a pointer to the next block in the file
- Pros
 - Can grow easily
- Cons
 - Bad sequential access perf.
 - Unreliable (why?)



File Allocation Table (FAT)

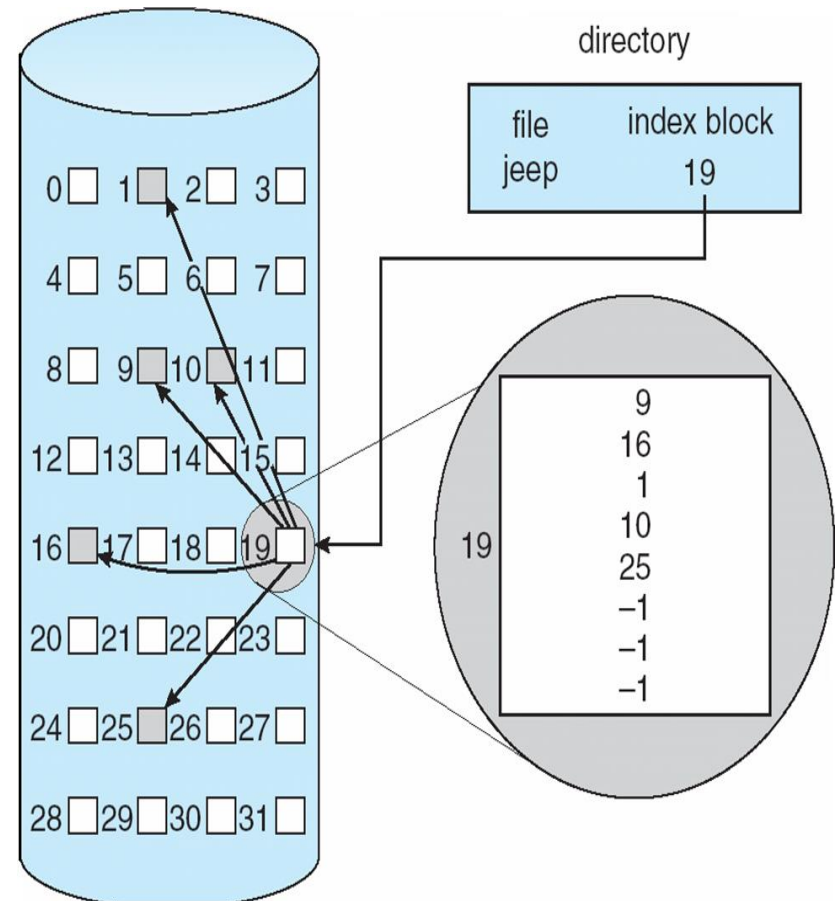
- A variation of linked allocation
 - Links are not stored in data blocks but in a separate table FAT[#of blocks]



- Directory entry points to the first block (217)
- FAT entry points to the next block (FAT[217] = 618)

Indexed Allocation


- Use per-file index block which holds block pointers for the file
 - Directory entry points to an index block (block 19)
 - The index block points to all blocks used by the file
- Pros
 - No external fragmentation
 - Fast random access
- Cons
 - Space overhead
 - File size limit (why?)



Multilevel Indexed Allocation

- Direct mapping for small files
- Indirect (2 or 3 level) mapping for large files

- 10 blocks are directly mapped
- 1 indirect pointer
 - 256 blocks
- 1 double indirect pointer
 - 64K blocks
- 1 triple indirect pointer
 - 16M blocks



Multilevel Indexed Allocation

- Direct mapping for small files
- Indirect (2 or 3 level) mapping for large files

- Pros
 - Easy to expand
 - Small files are fast (why?)
- Cons
 - Large files are costly (why?)
 - Still has size limit (e.g., 16GB)

