

# Filesystem

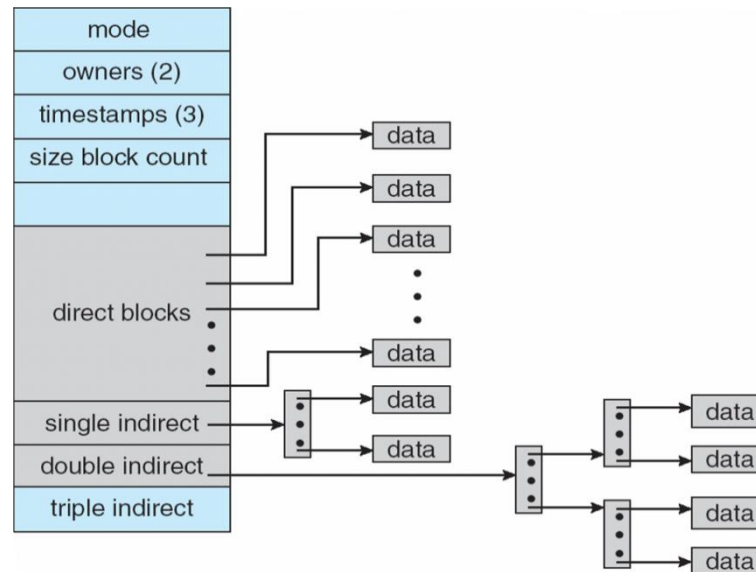
Disclaimer: some slides are adopted from book authors' slides with permission

# Concepts to Learn

- **Directory**
- **Caching**
- **Virtual File System**
- Putting it all together: FAT32 and Ext2
- Journaling
- Network filesystem (NFS)

# How To Access Files?

- Filename (e.g., “project2.c”)
  - Must be converted to the file header (inode)



- How to find the inode for a given filename?

# Directory

- A special file contains a table of
  - Filename (directory name) & inode number pairs

```
$ ls -li project2/  
24242928 directory  
25311615 dot_vimrc  
25311394 linux-2.6.32.60.tar.gz  
22148028 scheduling.html  
25311610 kvm-kernel-build  
22147399 project2.pdf  
25311133 scheduling.pdf  
25311604 kvm-kernel.config  
25311612 reinstall-kernel  
25311606 thread_runner.tar.gz
```



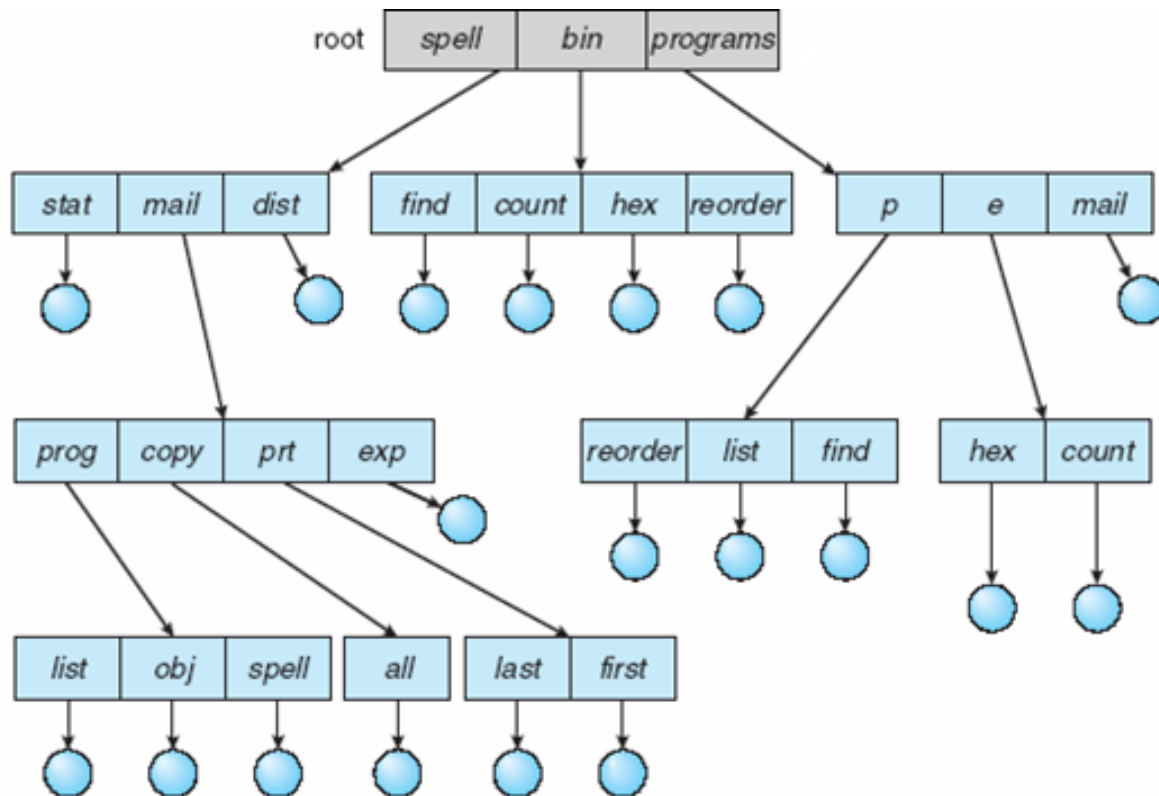
Inode  
number



Filename  
(or dirname)

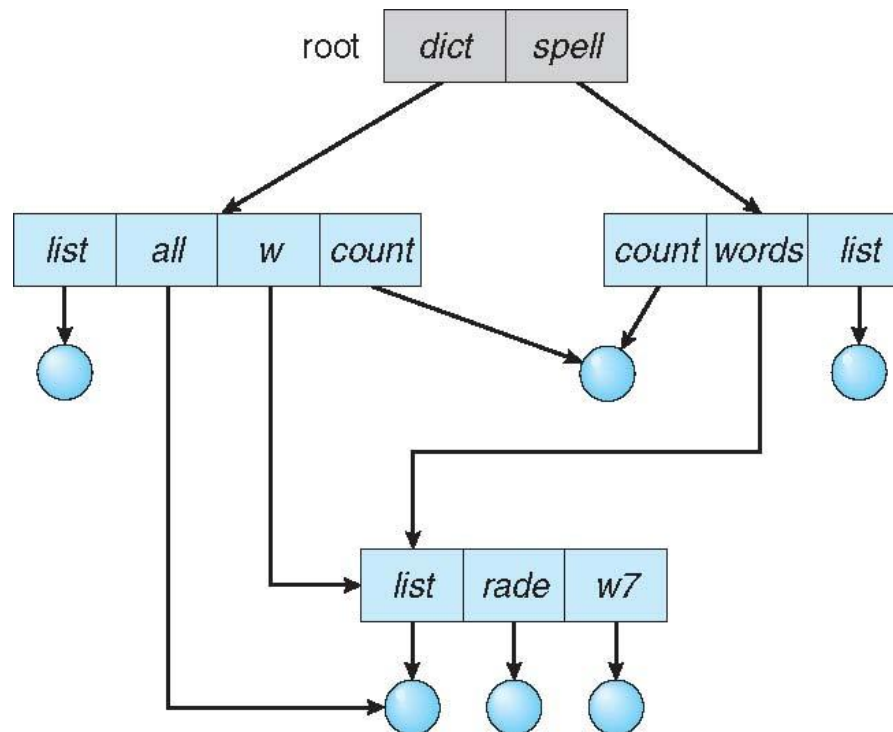
# Directory Organization

- Typically a tree structure



# Directory Organization

- Some filesystems support links → graph
  - Hard link: different names for a single file
  - Symbolic link: pointer to another file (“shortcut”)



# Name Resolution

- Path
  - A unique name of a file or directory in a filesystem
    - E.g., /usr/bin/top
- Name resolution
  - Process of converting a path into an inode
  - How many disk accesses to resolve “/usr/bin/top”?

# Name Resolution

- How many disk accesses to resolve “/usr/bin/top”?
  - Read “/” directory inode
  - Read first data block of “/” and search “usr”
  - Read “usr” directory inode
  - Read first data block of “usr” and search “bin”
  - Read “bin” directory inode
  - Read first block of “bin” and search “top”
  - Read “top” file inode
  - Total 7 disk reads!!!
    - This is the minimum. Why? Hint: imagine 10000 entries in each directory




# Directory Cache

- Directory name → inode number
  - Speedup name resolution process
    - When you first list a directory, it could be slow; next time you do, it would be much faster
  - Hashing
  - Keep only frequently used directory names in memory cache (how? LRU)

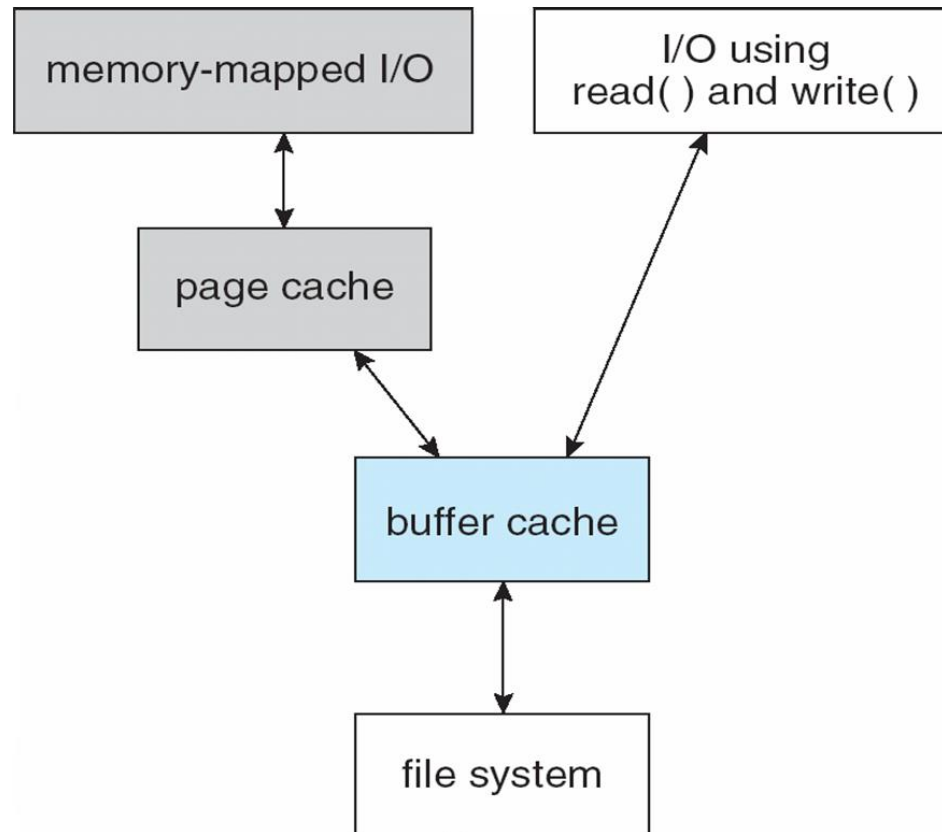
# Filesystem Related Caches

- Buffer cache
  - Caching frequently accessed disk blocks
- Page cache
  - Remember memory mapped files?
  - Map pages to files using virtual memory




# Non Unified Caches (Pre Linux 2.4)

- Problem: double caching

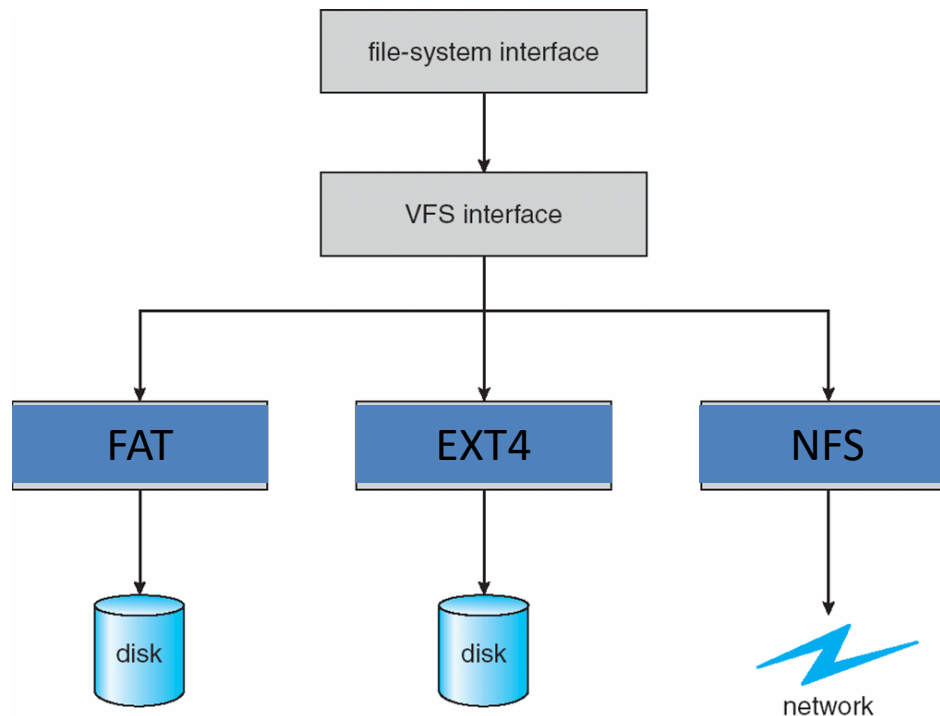


# Unified Buffer Cache



# Virtual Filesystem (VFS)

- Provides the same filesystem interface for different types of file systems



# Virtual Filesystem (VFS)

- VFS defined APIs
  - `int open(. . .)` —Open a file
  - `int close(. . .)` —Close an already-open file
  - `ssize_t read(. . .)` —Read from a file
  - `ssize_t write(. . .)` —Write to a file
  - `int mmap(. . .)` —Memory-map a file
  - ...
- All filesystems support the VFS apis

# Storage System Layers (in Linux)

