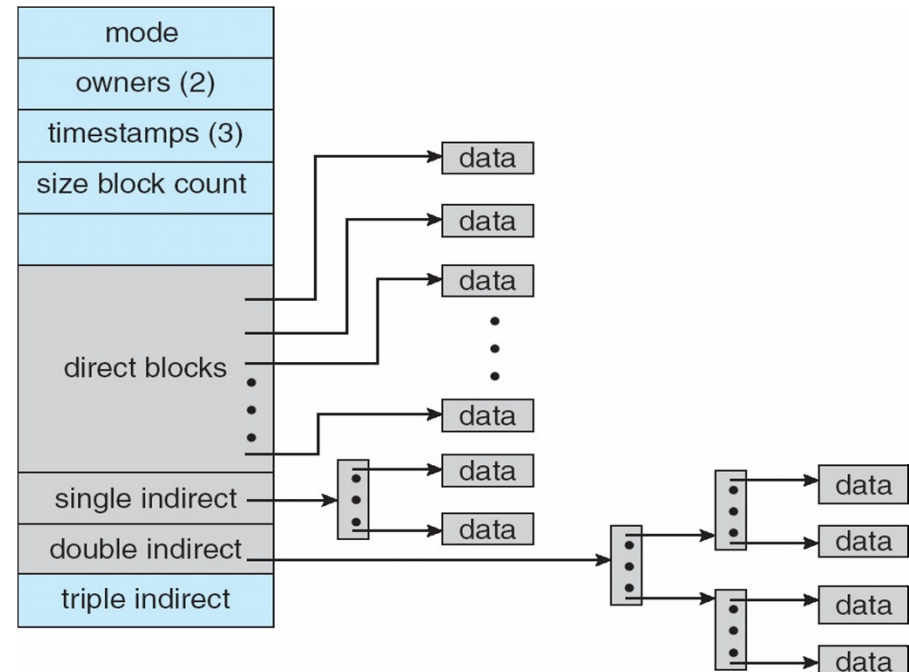


Filesystem

Disclaimer: some slides are adopted from book authors' slides with permission

Recap: Ext2

- Inode
 - 12 blocks are directly mapped, 1 indirect pointer, 1 double indirect pointer, 1 triple indirect pointer
- Maximum file size?
 - $\min(((b/4)^3 + (b/4)^2 + b/4 + 12) * b, (2^{32} - 1) * 512)$
 - 1K block size
 - $1K * (12 + 256 + 256^2 + 256^3) = 16GB$
 - 2K block size
 - $2K * (12 + 512 + 512^2 + 512^3) = 256G$
 - 4K block size
 - $(2^{32} - 1) * 512 = 2TB$



Journaling

- What happens if you lost power while updating to the filesystem?
 - Example
 - Create many files in a directory
 - System crashed while updating the directory entry
 - All new files are now “lost”
 - Recovery (fsck)
 - May not be possible
 - Even if it is possible to a certain degree, it may take very long time

Journaling

- Idea
 - First, write a log (journal) that describes all changes to the filesystem, then update the actual filesystem sometime later
- Procedure
 - Begin transaction
 - Write changes to the log (**journal**)
 - End transaction (**commit**)
 - At some point (**checkpoint**), synchronize the log with the filesystem

Recovery in Journaling Filesystems

- Check logs since the last checkpoint
- If a transaction log was committed, apply the changes to the filesystem
- If a transaction log was not committed, simply ignore the transaction

Types of Journaling

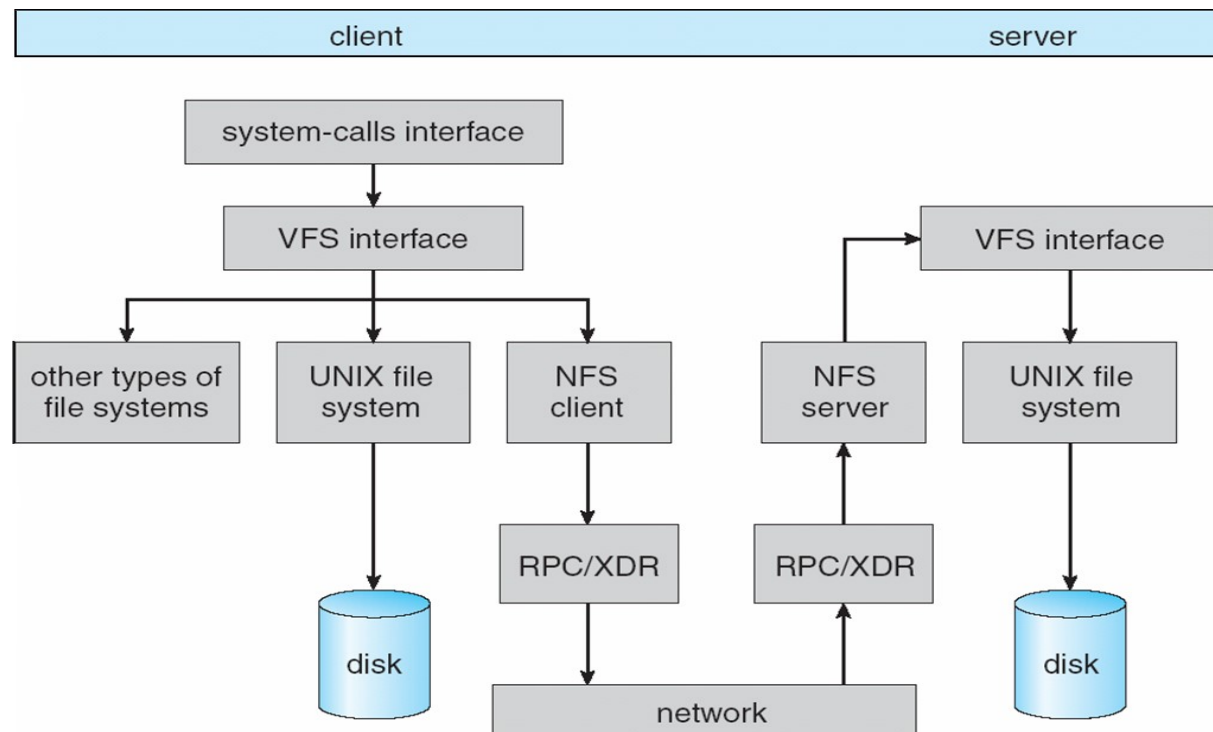
- Full journaling
 - All data & metadata are written twice
- Metadata journaling
 - Only write metadata of a file to the journal

Ext3 Filesystem

- Ext3 = Ext2 + Journaling
- Journal is stored in a special file
- Supported journaling modes
 - Write-back (metadata journaling)
 - Ordered (metadata journaling)
 - Data blocks are written to disk first
 - Metadata is written to journal
 - Data (full journaling)
 - Data and metadata are written to journal


Network File System (NFS)

- Developed in mid 80s by Sun Microsystems
- RPC based server/client architecture
- Attach a remote filesystem as part of a local filesystem



NFS Mounting Example

- Mount S1:/usr/share /usr/local



NFS vs. Dropbox

- NFS
 - All data is stored in a remote server
 - Client doesn't have any data on its local storage
 - Network failure → no access to data
- Dropbox
 - Client store data in its own local storage
 - Differences between the server and the client are exchanges to synchronize
 - Network failure → still can work on local data. Changes are synchronized when the network is recovered
- Which approach do you like more and why?

Summary

- I/O mechanisms
- Disk
- Disk allocation methods
- Directory
- Caching
- Virtual File System
- FAT and Ext2 filesystem
- Journaling
- Network filesystem (NFS)

Distributed Systems

Roadmap


- CPU management
- Memory management
- Disk management
- **Distributed System**
- Protection & Security
- Virtual machine

Today

- Distributed systems overview
- Basic network concepts
- TCP/IP protocol
- Sending/Receiving a packet in Linux

Distributed Systems

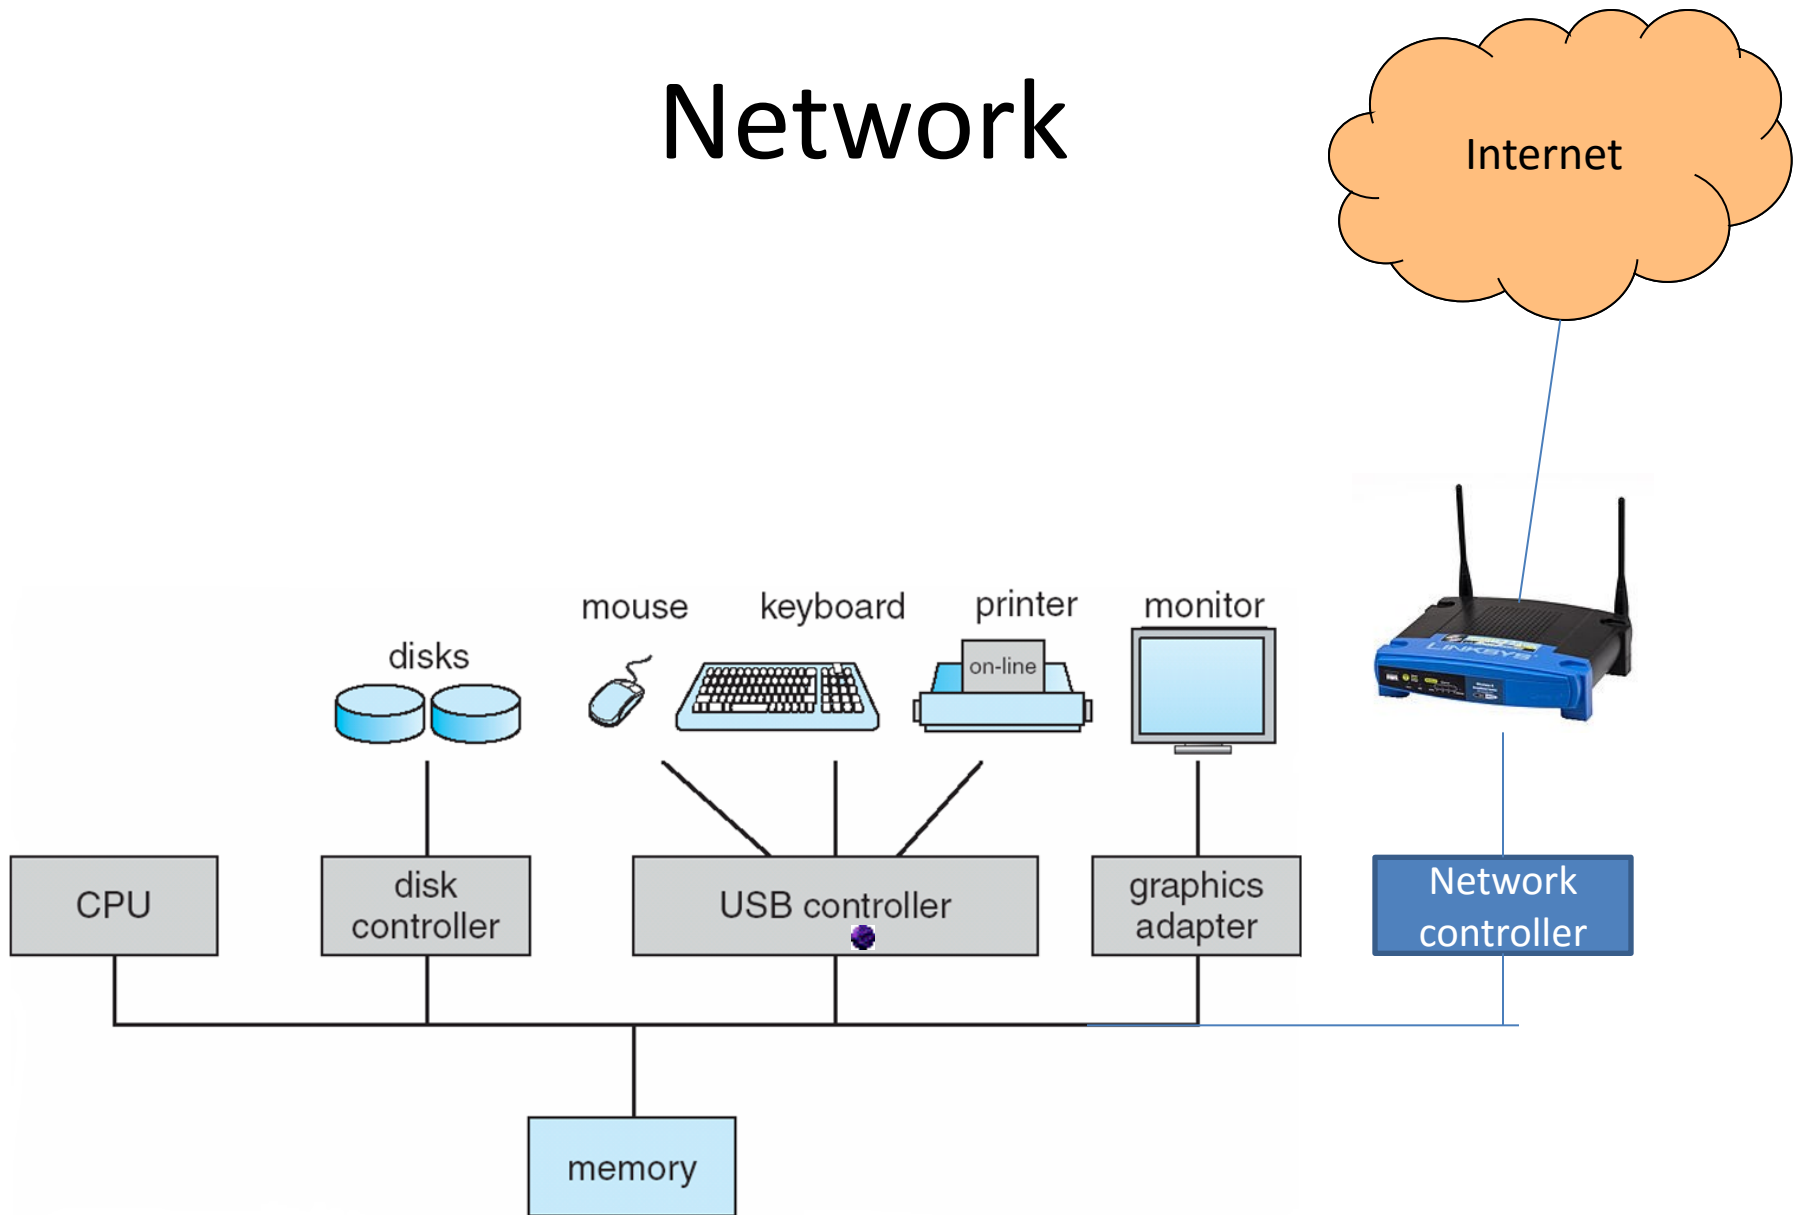
- A collection of connected computers



Why Distributed Computing?

- Resource sharing
 - Sharing and printing files at remote sites
 - Processing information in a distributed database
 - Using remote specialized hardware devices
- Performance
 - More computers → more performance
- Reliability
 - Detect and recover from site failure, function transfer, reintegrate failed site

Network

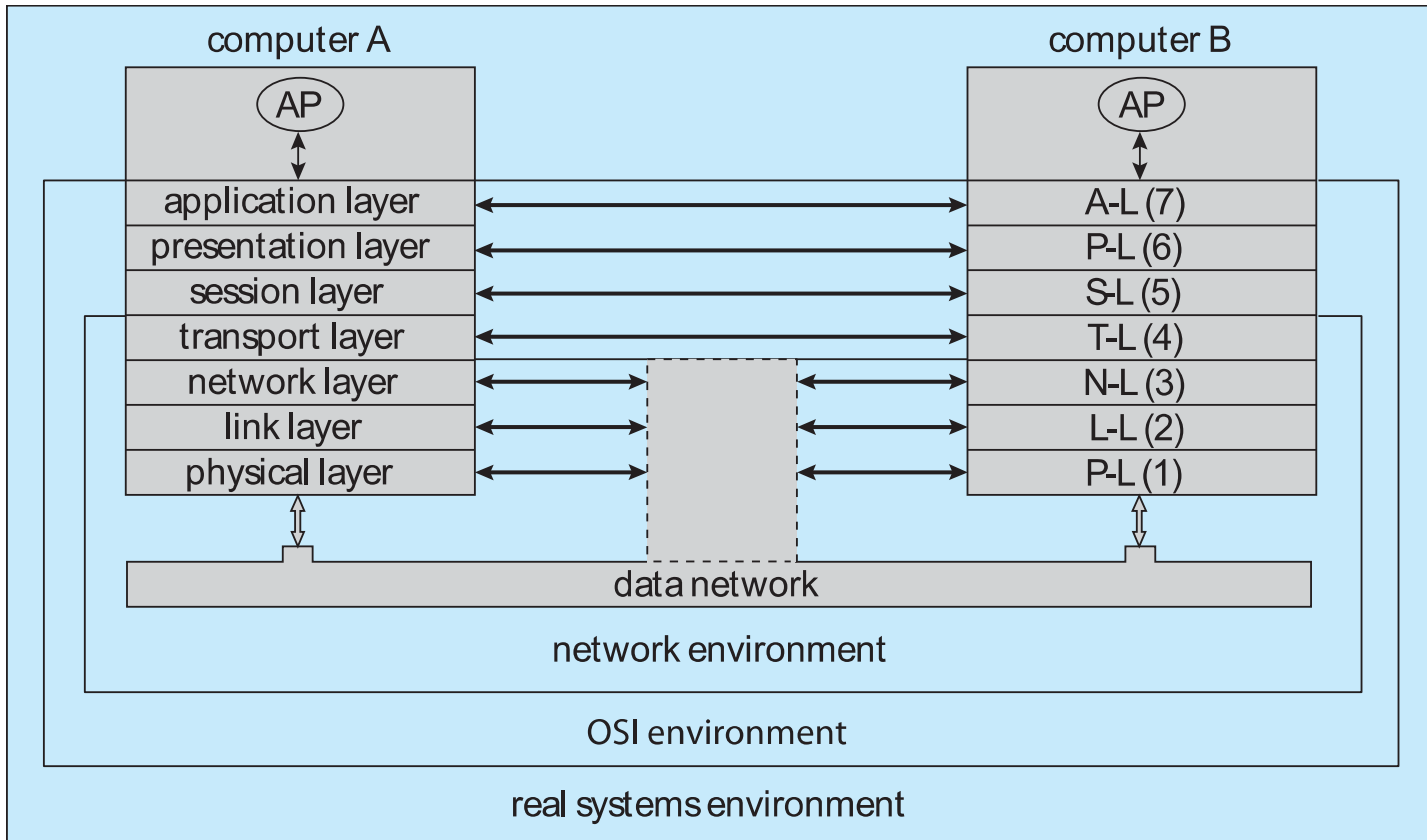


Terminologies

- Network
 - Physical medium of data transfer among multiple computers (e.g., Ethernet, CDMA,...)
- Packet
 - A unit of transfer in the network
- Protocol
 - A contract on how to transfer and receive data among the computers in the network

Communication Protocol

- Layered architecture



OSI 7 Layer communication model

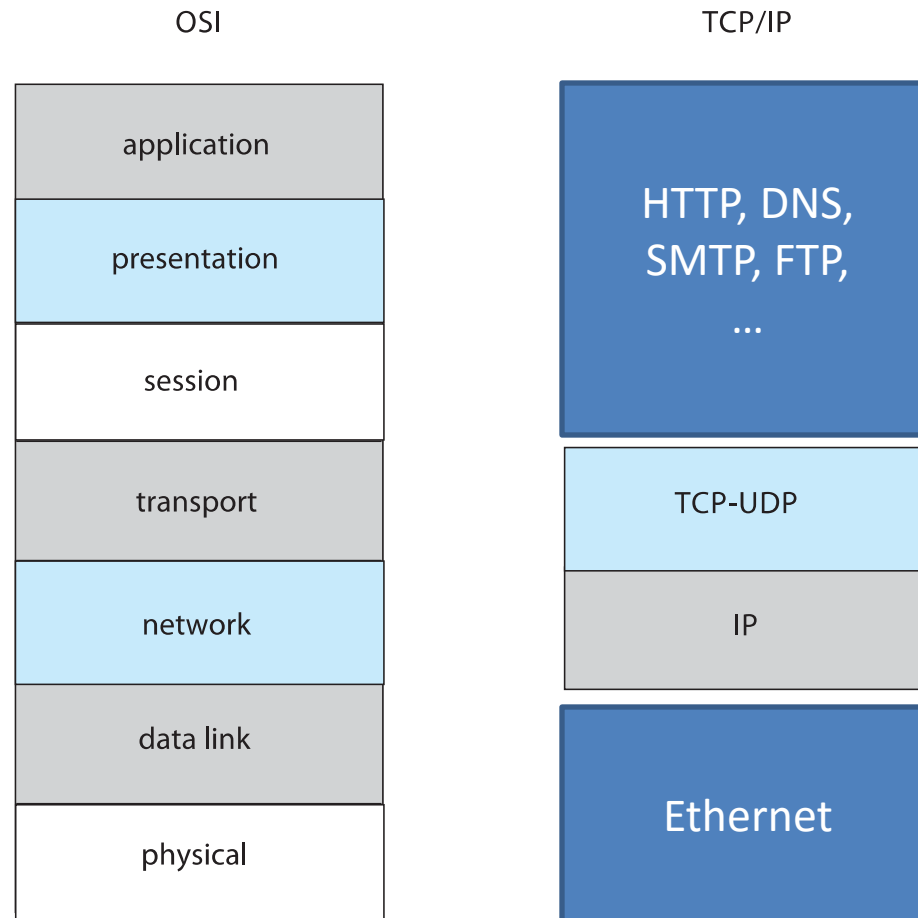
OSI Layers

1. Physical - electrical details of the physical transmission of a bit stream
2. Data-link - reliable data delivery on the physical medium
3. Network - addressing, routing, and delivery of packets
4. Transport – reliable delivery over the network
5. Session – session management among applications
6. Presentation – data representation, encryption
7. Application – application specific

■ Pros and Cons

- Pros: separation of concerns
- Cons: overhead, duplication

TCP/IP Protocol Layers



A Packet




Image source: http://en.wikipedia.org/wiki/Internet_protocol_suite


An Ethernet Frame

bytes

7	preamble—start of packet	each byte pattern 10101010
1	start of frame delimiter	pattern 10101011
2 or 6	destination address	Ethernet address or broadcast
2 or 6	source address	Ethernet address
2	length of data section	length in bytes
0–1500	data	message data
0–46	pad (optional)	message must be > 63 bytes long
4	frame checksum	for error detection

Internet Protocol (IP)

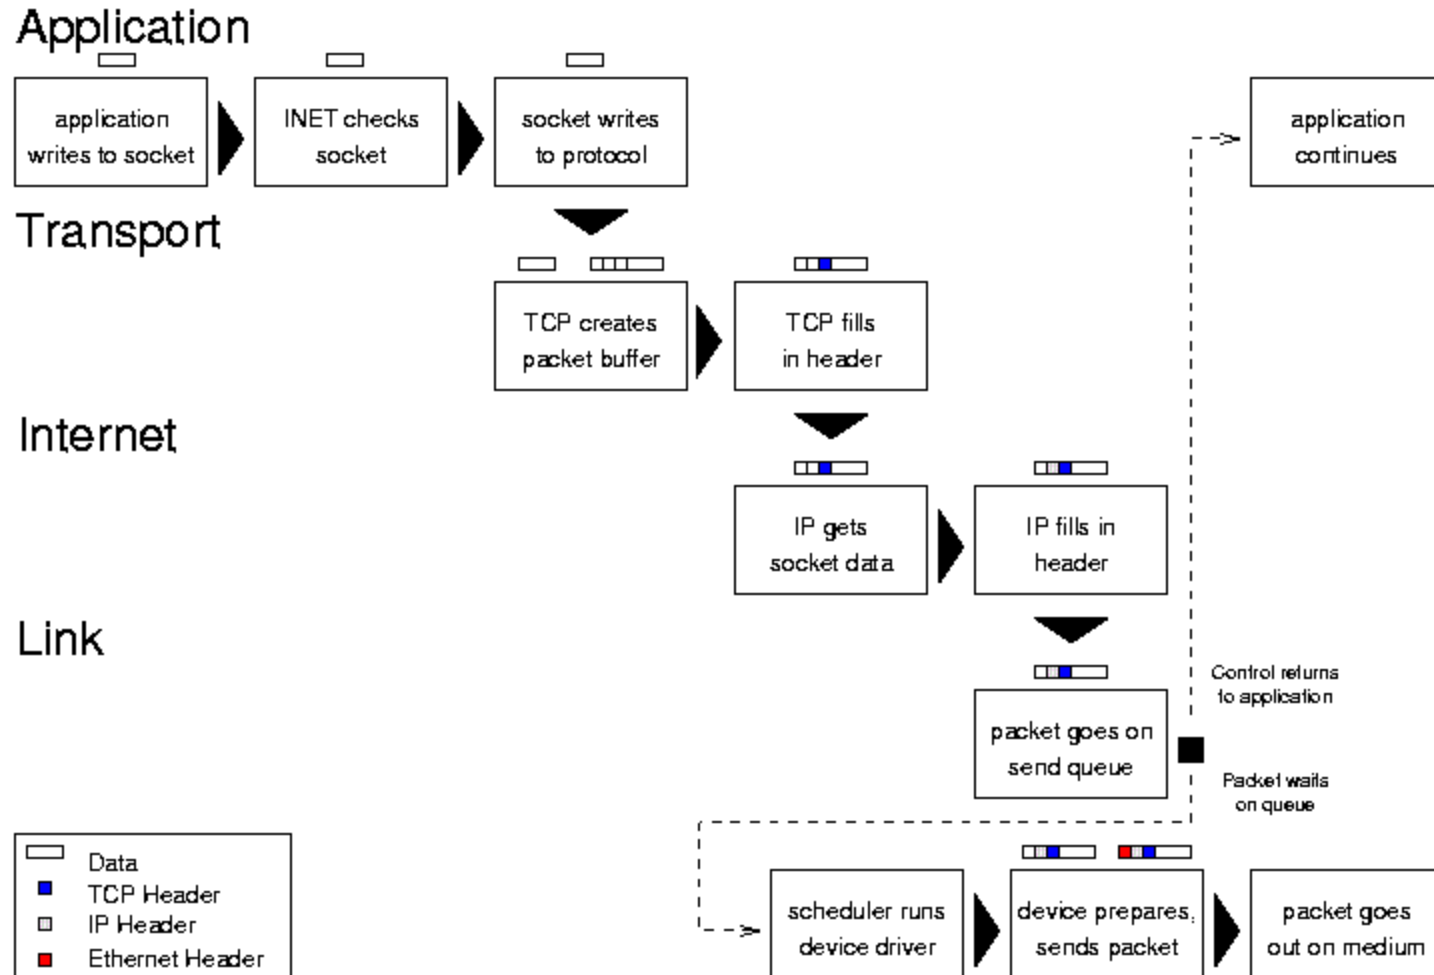
- Addressing
 - 32 bit (4 bytes) address: e.g., 129.237.123.1
- Routing
 - Forwarding packets through routers to reach their destination



Domain Name System (DNS)

- Domain name
 - Human readable internet address:
e.g., www.ku.edu
- How to map domain names to IP addresses?
 - www.ku.edu → 129.237.11.182
 - www.google.com → may vary depending on your location, server load, etc.
- Domain Name System
 - A distributed database of domain name, IP addr.


Sending a Packet




Source: G. Herrin, [Linux IP Networking: A Guide to the Implementation and Modification of the Linux Protocol Stack](#), 2000

Receiving a Packet


Application




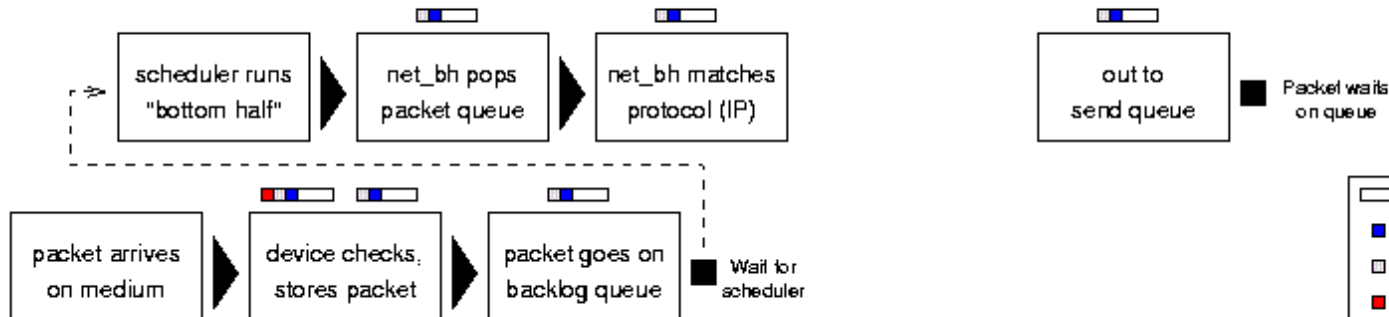
Transport



Internet



Link



Source: G. Herrin, [Linux IP Networking: A Guide to the Implementation and Modification of the Linux Protocol Stack](#), 2000