

# Virtual Machine Monitor

# Recap

- Emulator
  - Emulate hardware (e.g., Nintendo) in software
- Java virtual machine (JVM)
  - Platform neutral machine (typically s/w) that runs java bytecode
- Virtual machine monitor (VMM)
  - A kind of OS for virtual machines

# Java Bytecode

## Java Source

```
int f(){  
    int a,b,c;  
    ..  
    c = a + b + 1;  
    ..  
}
```



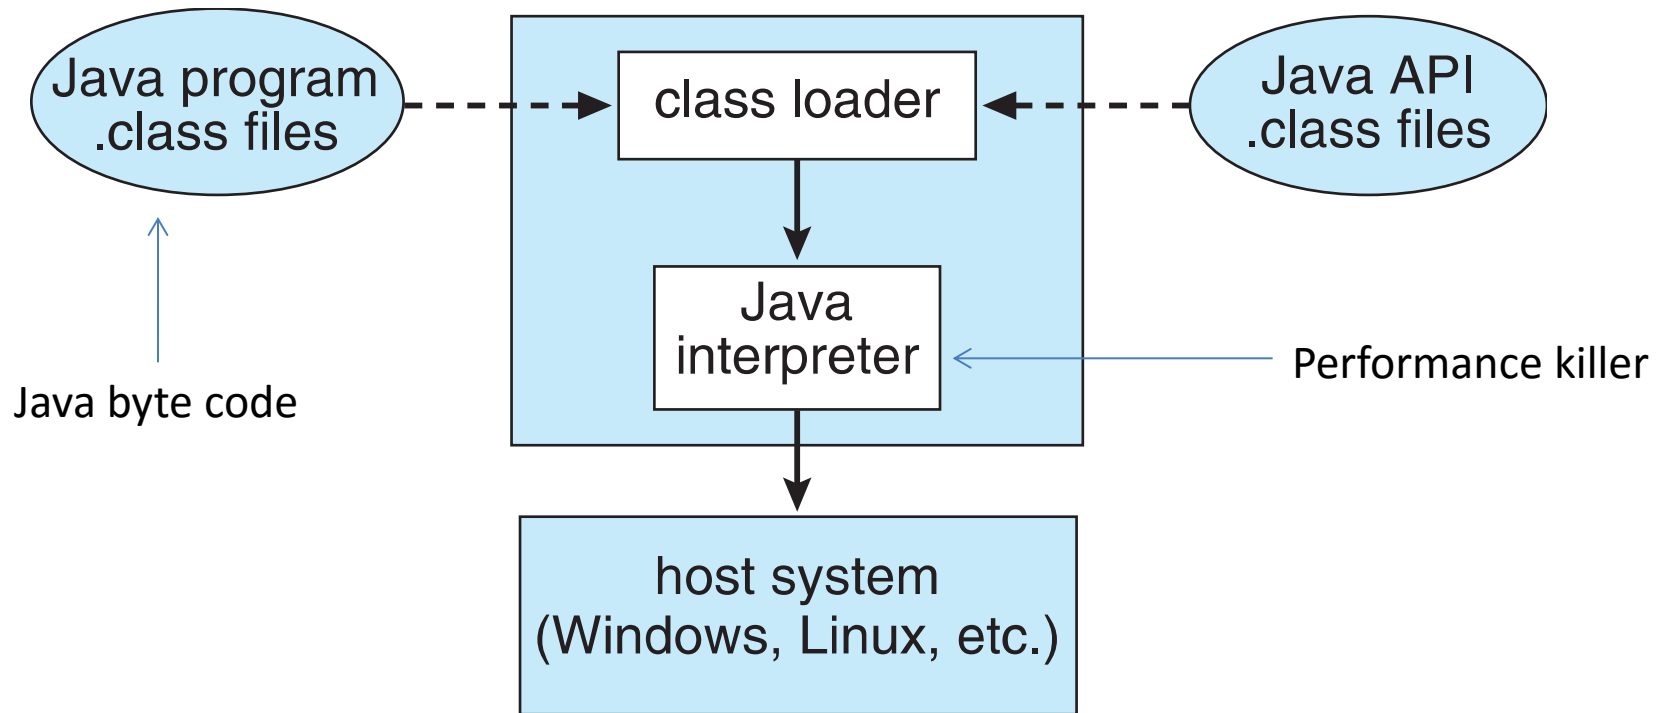
javac compiler

## Java Bytecode

```
int f();  
iload a  
iload b  
iconst 1  
iadd  
iadd  
istore c
```



- Java bytecode = instructions for JVM
- One byte opcode + variable length args
  - 198 instructions are in use

# Recap: Java Virtual Machine



# Recap: Types of VMM

- Native (or Type 1) VMM
  - VMM runs directly on top of bare hardware
  - Vmware ESX, Microsoft Hyper-V
  - VMM is a kind of a OS on its own right
- Hosted (or Type 2) VMM
  - VMM runs within an OS
  - VirtualBox, VMWare Workstation
  - VMM relies on functionalities of the host OS



# Today

- Virtualization
- Container
- Docker

# Virtualizing Interrupts & I/O

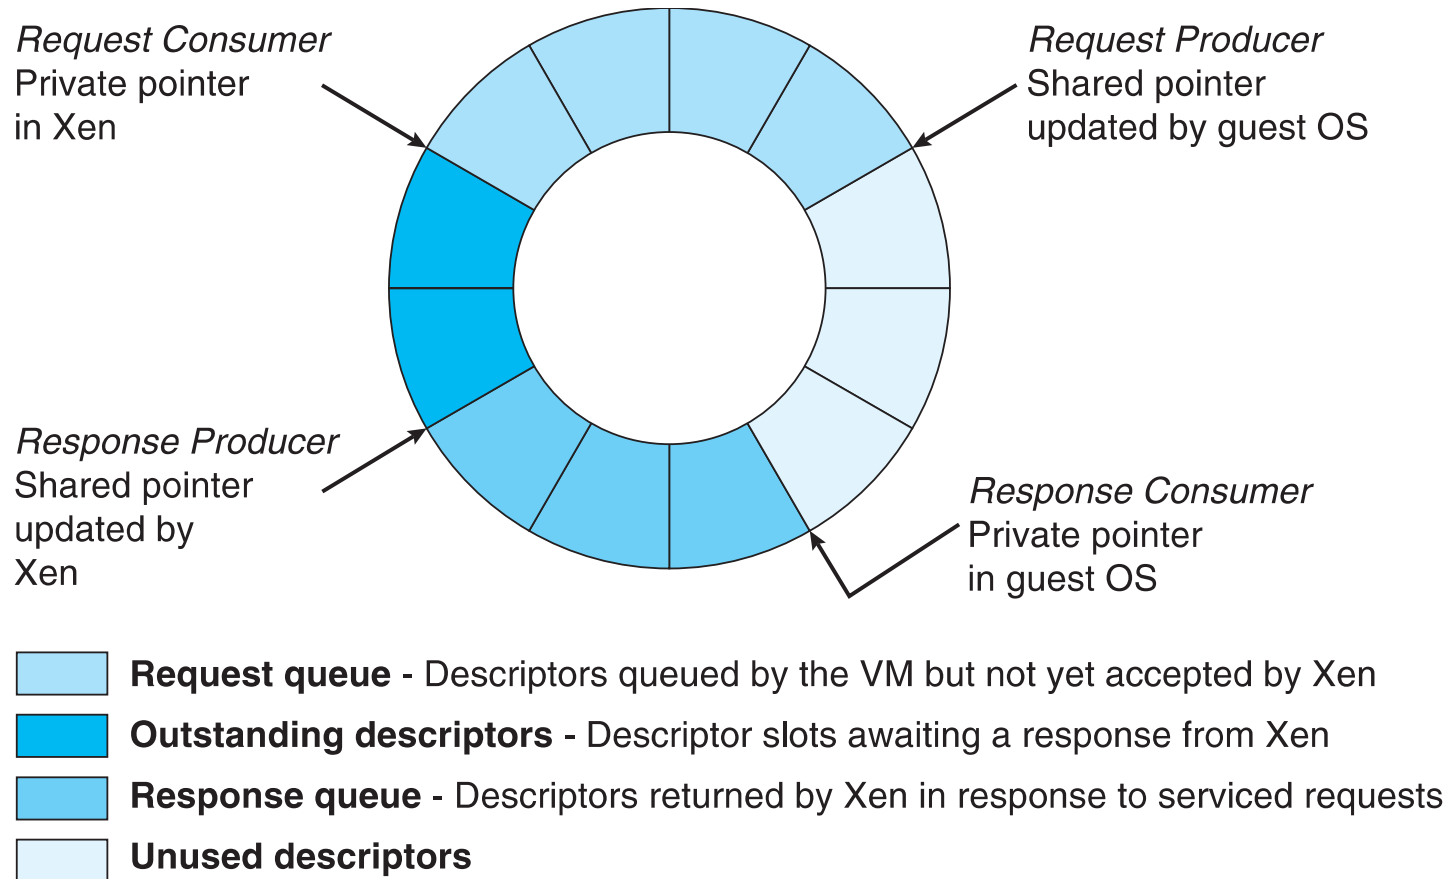
- VMM receives h/w interrupts
  - Determines which VM to receive
  - Emulate interrupt controller for the VM
- VMM emulate a specific h/w devices
  - Guest OS → VMM → devices
    - E.g., AMD Lance PCNet ethernet device
- Lots of I/O → performance killers

# Para-virtualization

- Idea: provides simple/fast APIs to guests
  - Instead of emulating actual hardware (e.g., PCNet32 ethernet card)
  - Pros
    - can be a lot faster (more efficient I/O)
  - Cons
    - need to modify the guest OS




# I/O in Xen via Shared Buffer



- Request queue** - Descriptors queued by the VM but not yet accepted by Xen
- Outstanding descriptors** - Descriptor slots awaiting a response from Xen
- Response queue** - Descriptors returned by Xen in response to serviced requests
- Unused descriptors**


# IOMMU

- Problem: How to do DMA in a VM?
  - DMA controller needs host physical address, not guest physical address
- IOMMU
  - MMU for IO devices
  - maps guest physical → host physical for the I/O devices




[https://en.wikipedia.org/wiki/Input%E2%80%93output\\_memory\\_management\\_unit#/media/File:MMU\\_and\\_IOMMU.svg](https://en.wikipedia.org/wiki/Input%E2%80%93output_memory_management_unit#/media/File:MMU_and_IOMMU.svg)

# LXC: OS (Linux) Container



- Same kernel, separate user-space
- Virtualize OS, not machine
- Low overhead, flexible

# Docker: Application Container



- A container contain one application (process)
- Built on top of OS containers
- Even more flexible

# Summary

- Virtual Machine (hardware virtualization)
  - Trap & emulate
  - Binary translation
  - Para-virtualization
  - Hardware support for virtualization
- Containers
  - OS container: same kernel, different user-space
  - App container: same kernel, per-process space