

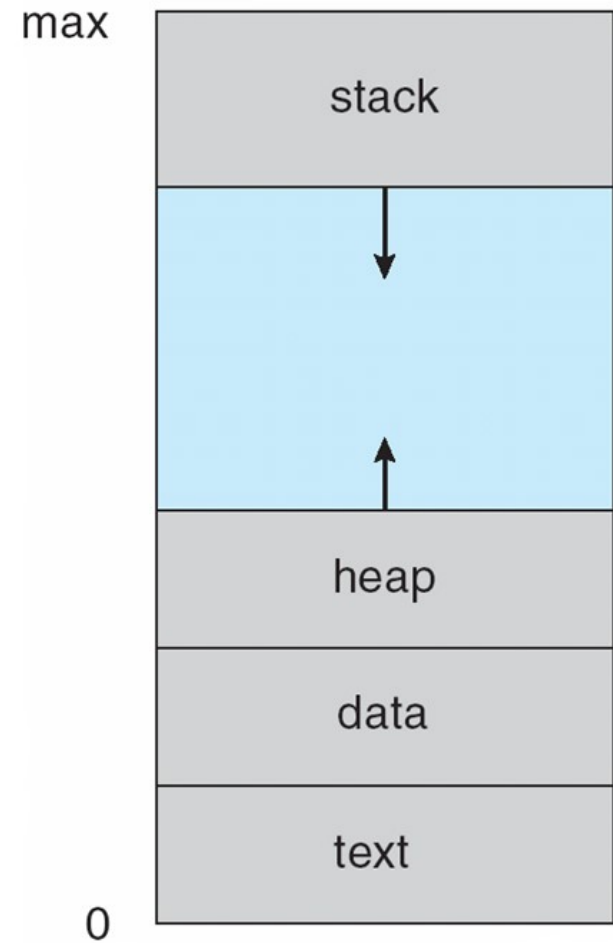
Final Review

2nd Half

- Memory management
- Disk management
- Network and Security
- Virtual machine

Virtual Memory (VM)

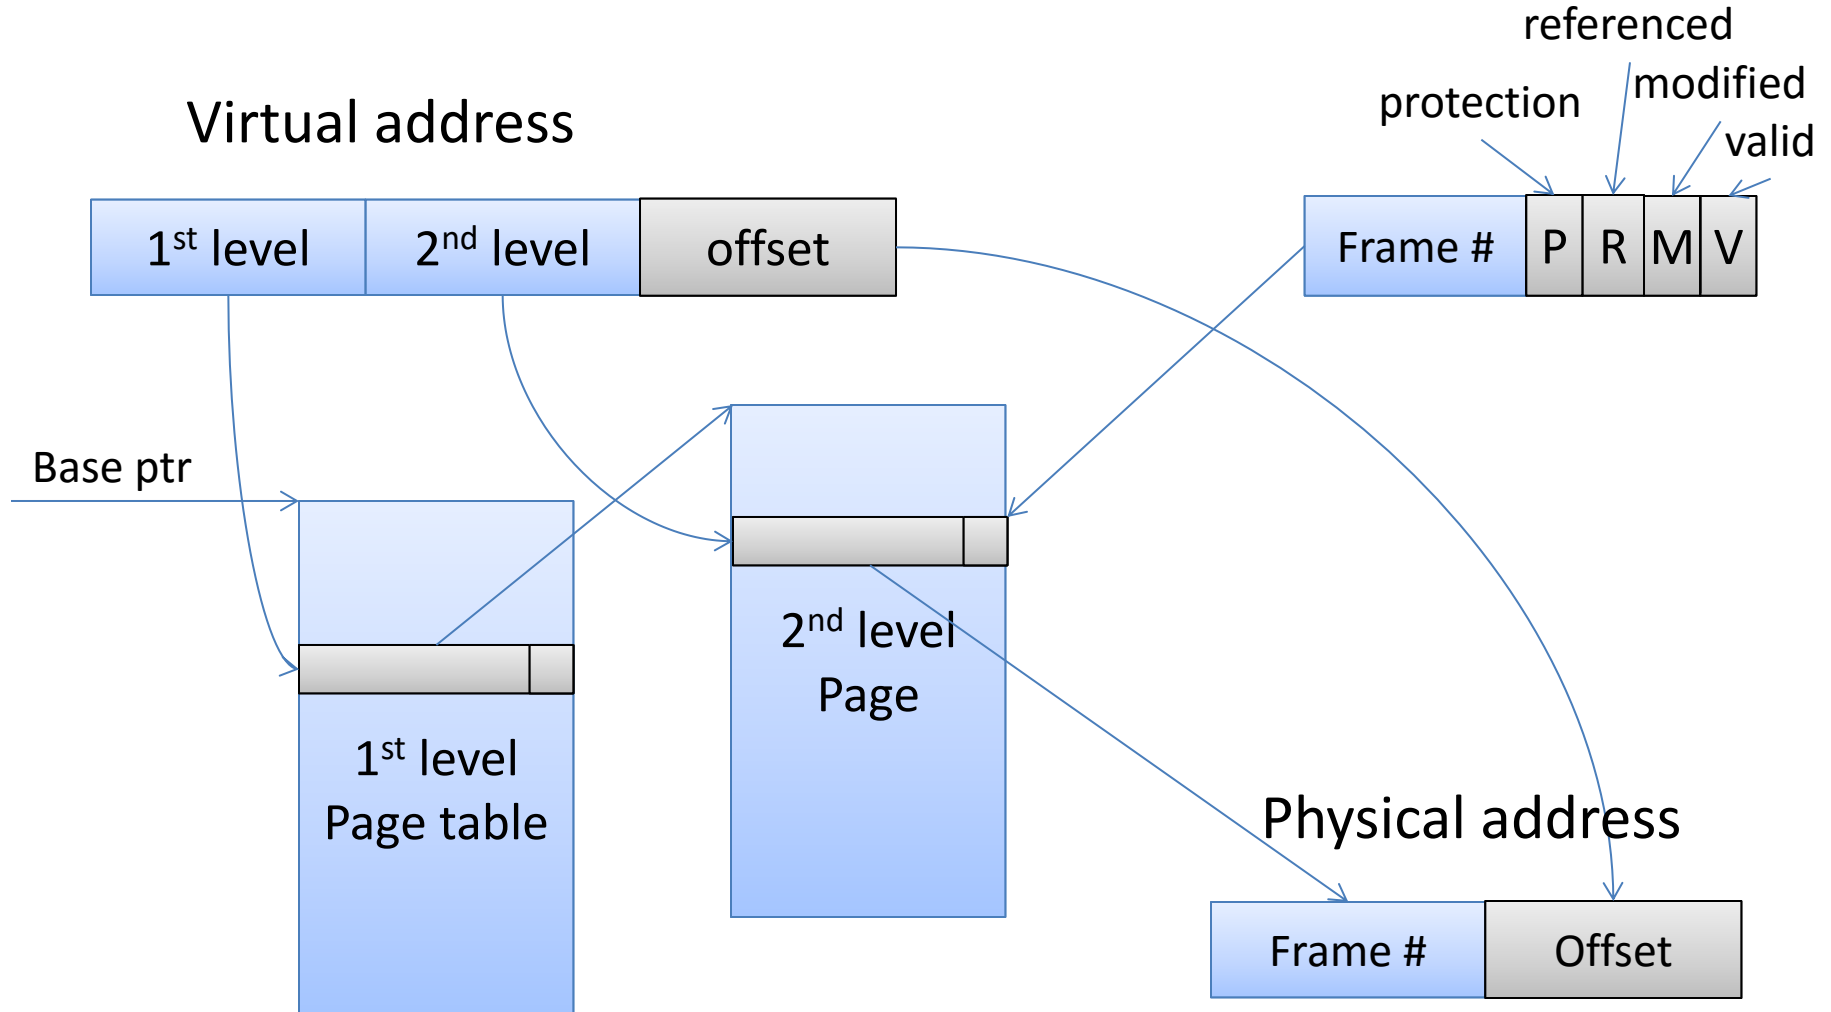
- Abstraction
 - 4GB (32bit) linear address space for each process
- Reality
 - 1GB of actual physical memory shared with 20 other processes



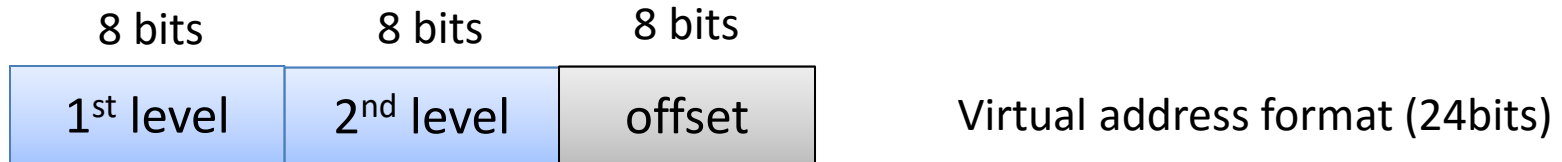
Paging

- MMU (memory management unit)
 - Segmentation (base+limit),
 - Paging (page table)
 - TLB (translation lookaside buffer)
- Page table
 - Virtual addr → physical address translation table
- Other concepts to know
 - Fragmentation

Two Level Address Translation



Quiz: Address Translation

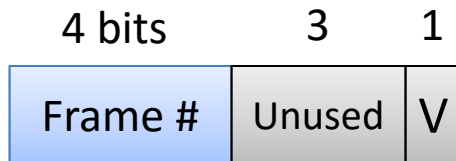


- What is the size of a single page?
 - 256 bytes
- What is the size of the virtual address space?
 - $2^{24} = 16\text{MB}$
- What is the size of the physical address space?
 - $2^{12} = 4\text{KB}$

Quiz: Address Translation



Virtual address format (24bits)



Page table entry (8bit)

Vaddr: 0x0703FE

Paddr: 0x3FE

Vaddr: 0x072370

Paddr: ???

Vaddr: 0x082370

Paddr: ???

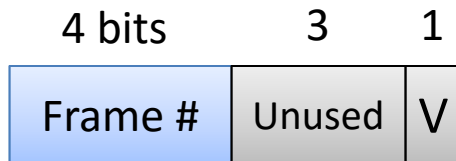
Page-table base address = 0x100

Addr	+0	+1	+2	+3	+4	+5	+6	+7	+8	+A	+B	+C	+D	+E	+F
0x000				31											
0x010															
0x020				41											
..															
0x100	00	01						01	00			01			
..															
0x200															

Quiz: Address Translation



Virtual address format (24bits)



Page table entry (8bit)

Vaddr: 0x0703FE

Paddr: 0x3FE

Vaddr: 0x072370

Paddr: **0x470**

Vaddr: 0x082370

Paddr: **invalid**

Page-table base address = 0x100

Addr	+0	+1	+2	+3	+4	+5	+6	+7	+8	+A	+B	+C	+D	+E	+F
0x000				31											
0x010															
0x020				41											
..															
0x100	00	01						01	00			01			
..															
0x200															

Demand Paging

- Key idea
 - no need to load all pages on memory.
- Page fault
 - Occurs when there's no valid virt -> phys addr translation
 - Allocate a free frame [and load a page from the disk]
- Copy-on-write
 - Fork()/exec() semantic
- Other concepts to know
 - Page Table Entry (PTE) bits: valid/invalid, ...

Page Replacement & Swapping

- Least Recently Used (LRU)
 - Approximation: clock algorithm, N-clock algorithm
 - When swapping occurs in the page fault handler?
- Other concepts to know
 - Thrashing
 - FIFO, MIN (optimal) replacement algorithms

Exercise: LRU Replacement

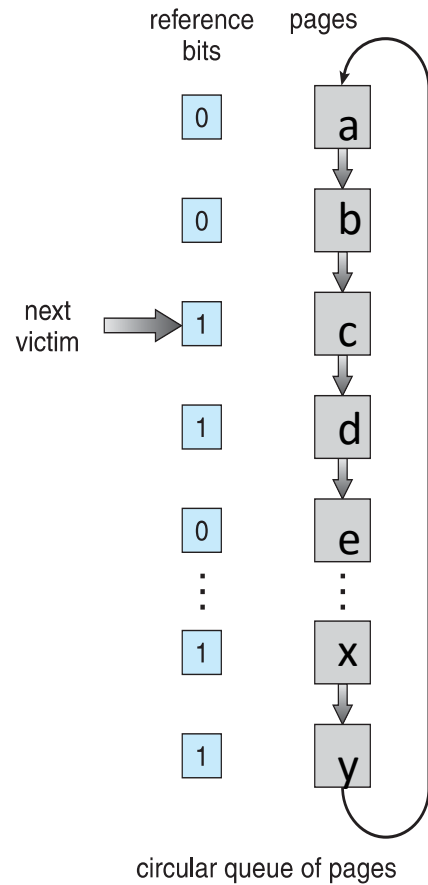
reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



page frames

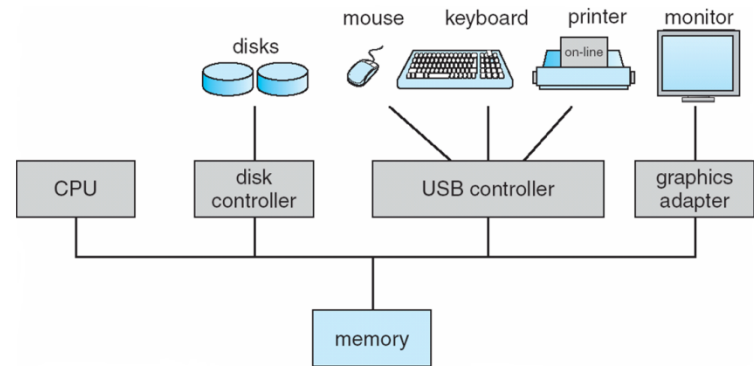
Second Chance Algorithm



(a)

I/O Devices

- Block devices
 - E.g., disk, cd-rom, USB stick
 - High speed, block (sector) level accesses
- Character devices
 - E.g., keyboard, mouse, joystick
 - Low speed, character level accesses
- Network devices
 - E.g., ethernet, wifi, bluetooth
 - Socket interface

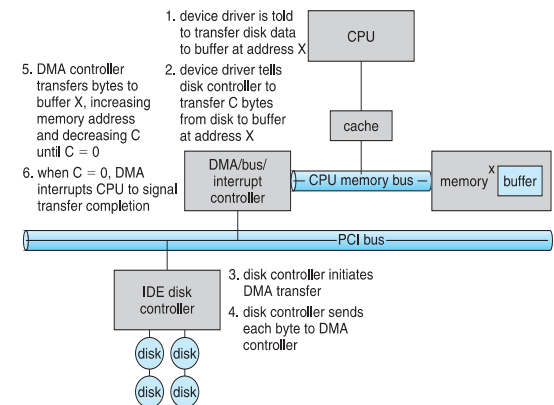


Disk

- Magnetic disk
 - Key characteristic: long seek time
- Solid state disk (SSD)
 - Key characteristic: zero seek time
- Performance metrics
 - Access latency, price/storage space

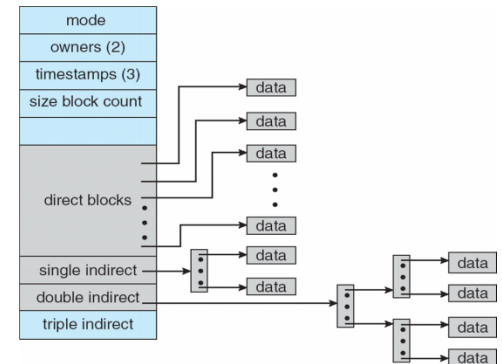
I/O Mechanisms

- Memory-mapped I/O vs. I/O instructions
 - MMIO: Address space → registers in the hardware
 - I/o instructions: *inb* or *outb* in x86
- Programmed I/O (PIO) vs. DMA
 - PIO: Reads/writes by CPU
 - DMA: DMA hw operation steps



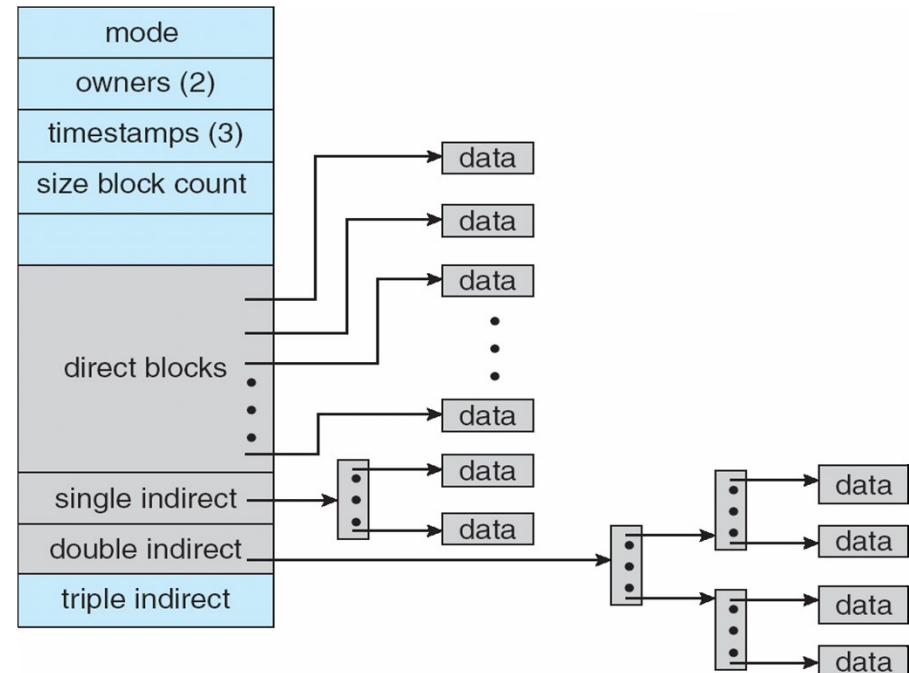
Filesystem

- Disk block allocation methods
 - Continuous allocation
 - **Linked allocation**
 - **Indexed allocation**
- Other concepts to know
 - Metadata, inode, directory
 - **Direct/single/double/triple indirect blocks**
 - How they differ. Why different levels?



Recap: Ext2

- Inode
 - 12 blocks are directly mapped, 1 indirect pointer. 1 double indirect pointer, 1 triple indirect pointer. 32bit pointer
- Maximum file size?
 - $\min(((b/4)^3 + (b/4)^2 + b/4 + 12) * b, (2^{32} - 1) * 512)$
 - 1K block size
 - $1K * (12 + 256 + 256^2 + 256^3) = 16GB$
 - 2K block size
 - $2K * (12 + 512 + 512^2 + 512^3) = 256G$
 - 4K block size
 - $(2^{32} - 1) * 512 = 2TB$
- Maximum disk size?
 - $2^{32} * \text{block size}$



Name Resolution

- How many disk accesses to resolve “/usr/bin/top”?
 - Read “/” directory inode
 - Read first data block of “/” and search “usr”
 - Read “usr” directory inode
 - Read first data block of “usr” and search “bin”
 - Read “bin” directory inode
 - Read first block of “bin” and search “top”
 - Read “top” file inode
 - Total 7 disk reads!!!
 - This is the minimum. Why? Hint: imagine 10000 entries in each directory

Storage caches

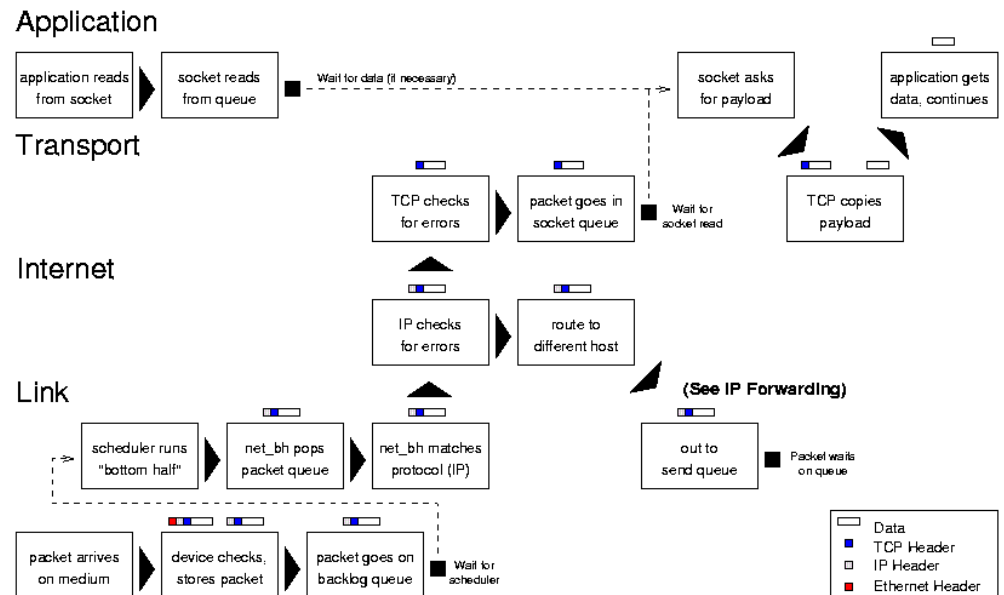
- Directory cache
 - For faster name resolution
 - Steps to resolve a file name to inode
- Buffer cache
 - Frequently used disk blocks
- Other concepts to know
 - Unified buffer and page cache
 - Journaling

Types of Journaling

- Full journaling
 - All data & metadata are written twice
- Metadata journaling
 - Only write metadata of a file to the journal

Network

- OSI 7 layers vs. TCP/IP 5 layers
- Steps for sending/receiving a packet



Security & Virtual Machine

- Buffer overflow bugs
 - Look at the example code
- Virtual machine monitor (VMM)
- Native VMM vs. hosted VMM
- Full virtualization vs. para virtualization

1st Half

OS Structure

- User mode/ kernel mode
 - Memory protection, privileged instructions
- System call
 - Definition, examples, how it works?
- Other concepts to know
 - Monolithic kernel vs. Micro kernel

Process/Thread

- Address space layout
 - Code, data, heap, stack
- Process states
 - new, ready, running, waiting, terminated
- Other concepts to know
 - Process Control Block
 - Context switch
 - Zombie, Orphan
 - Communication overheads of processes vs. threads

Synchronization

- Spinlock
 - Semantics of test & set, compare & swap
 - Implementation using compare & swap
- Blocking synchronization primitives
 - Busy waiting vs. blocking
 - Mutex, semaphore, monitor
- Example problems
 - Bounded buffer, read/writer

Deadlock

- Deadlock conditions
 - mutual exclusion, no preemption, hold and wait, circular wait
- Deadlock prevention
- **Banker's algorithm**
- Other concepts to know
 - Starvation vs. deadlock

Scheduling

- Three main schedulers
 - FCFS, SJF/SRTF, RR
 - Gant chart examples
- Other concepts to know
 - Fair scheduling (CFS)
 - Fixed priority scheduling
 - Multi-level queue scheduling
 - Load balancing and multicore scheduling

Thank you!