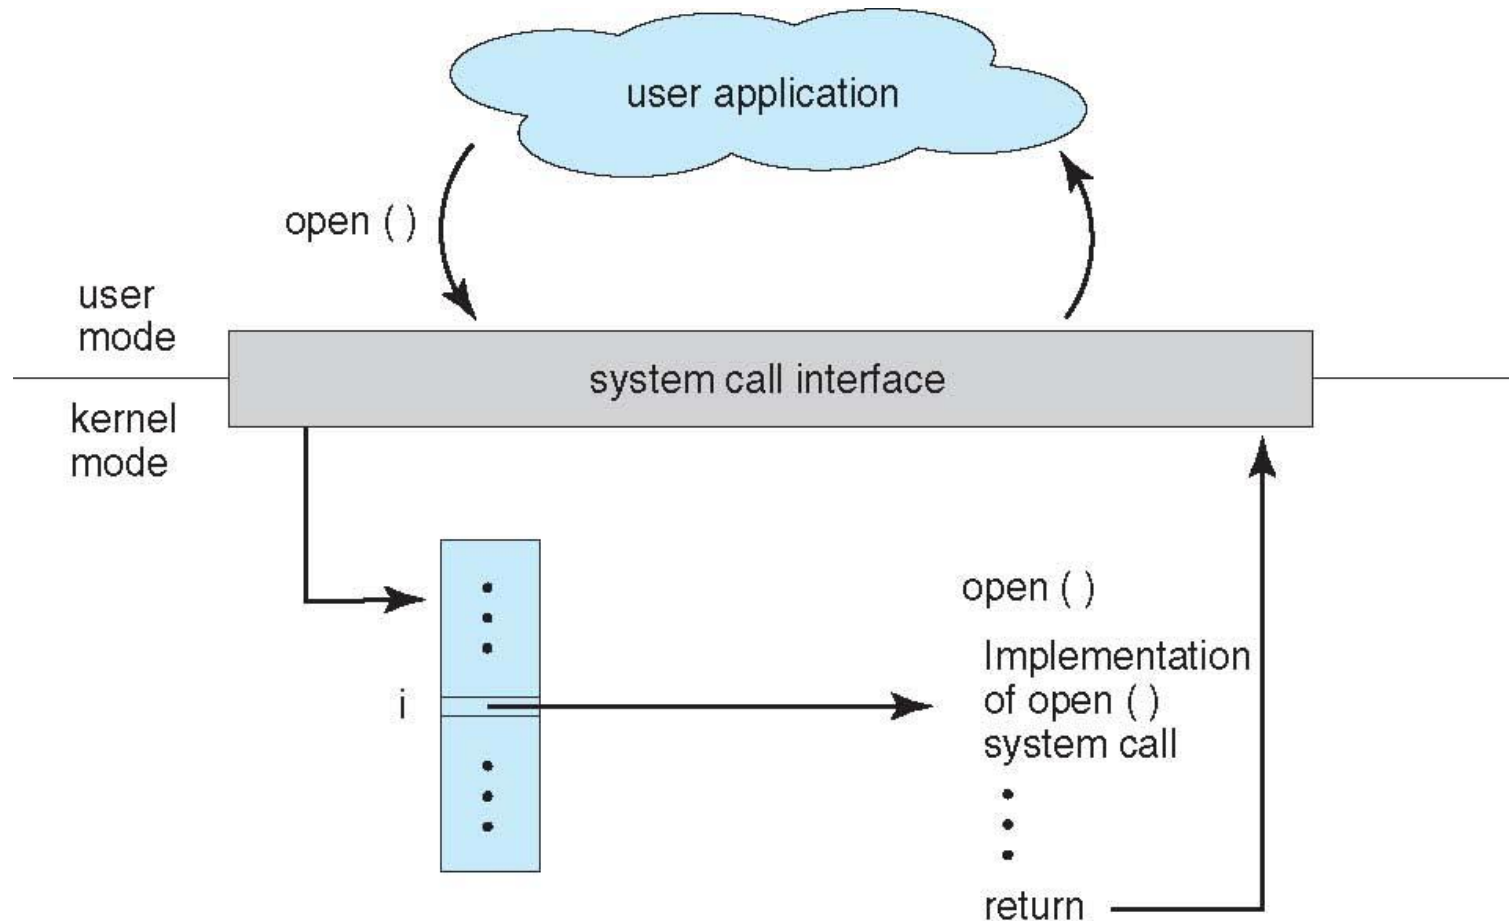


# Midterm Review

# OS Structure

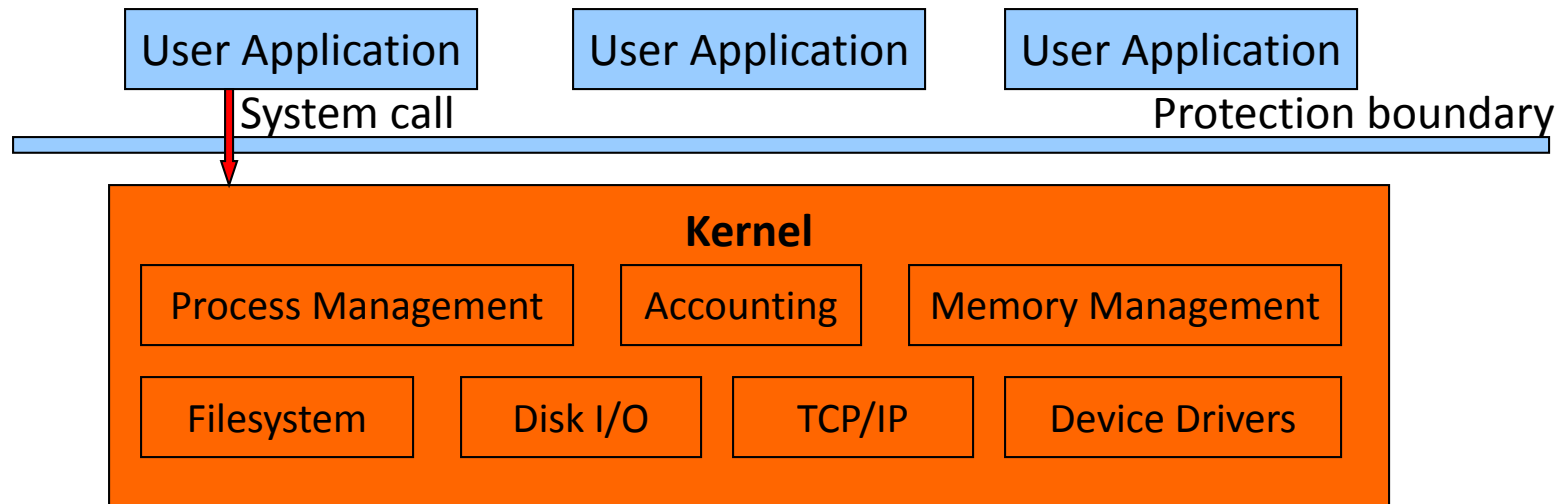
- User mode/ kernel mode
  - Memory protection, privileged instructions
- System call
  - Definition, examples, how it works?
- Other concepts to know
  - Monolithic kernel vs. Micro kernel

# API - System Call - OS



# UNIX: Monolithic Kernel

- Implements CPU scheduling, memory management, filesystems, and other OS modules all in a single big chunk




- Pros and Cons
  - + Overhead is low
  - + Data sharing among the modules is easy
  - Too big. (device drivers!!!)
  - A bug in one part of the kernel can crash the entire system

# Process

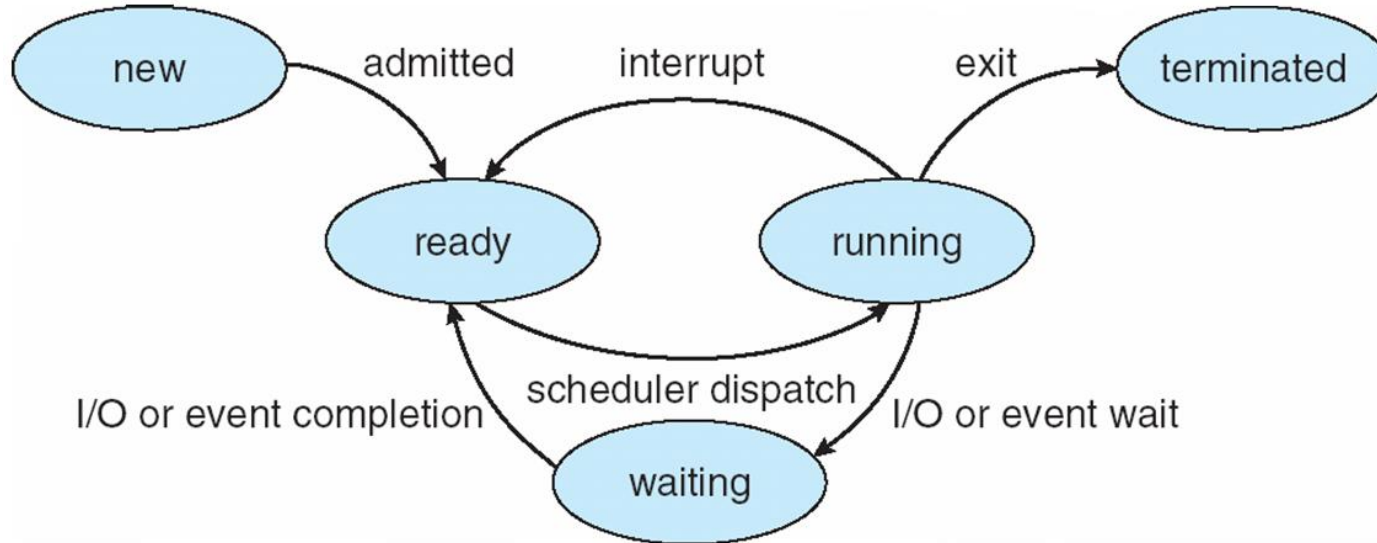
- Address space layout
  - Code, data, heap, stack
- Process states
  - new, ready, running, waiting, terminated
- Other concepts to know
  - Process Control Block
  - Context switch
  - Zombie, Orphan
  - Communication overheads of processes vs. threads

# Process Address Space

- Text
  - Program code
- Data
  - Global variables
- Heap
  - Dynamically allocated memory
    - i.e., Malloc()
- Stack
  - Temporary data
  - Grow at each function call



# Process State



- **running**: Instructions are being executed
- **waiting**: The process is waiting for some event to occur
- **ready**: The process is waiting to be assigned to a processor

# Quiz

```
int count = 0;
int main()
{
    int pid = fork();
    if (pid == 0){
        count++;
        printf("Child: %d\n", count);
    } else{
        wait(NULL);
        count++;
        printf("Parent: %d\n", count);
    }
    count++;
    printf("Main: %d\n", count);
    return 0;
}
```


- Hints
  - Each process has its own private address space
  - Wait() blocks until the child finish
- Describe the output.
  - Child: 1
  - Main: 2
  - Parent: 1
  - Main: 2



# Inter-Process Communication


- Shared memory
- Message passing

# Models of IPC



(a)

message passing



(b)

shared memory

# Quiz

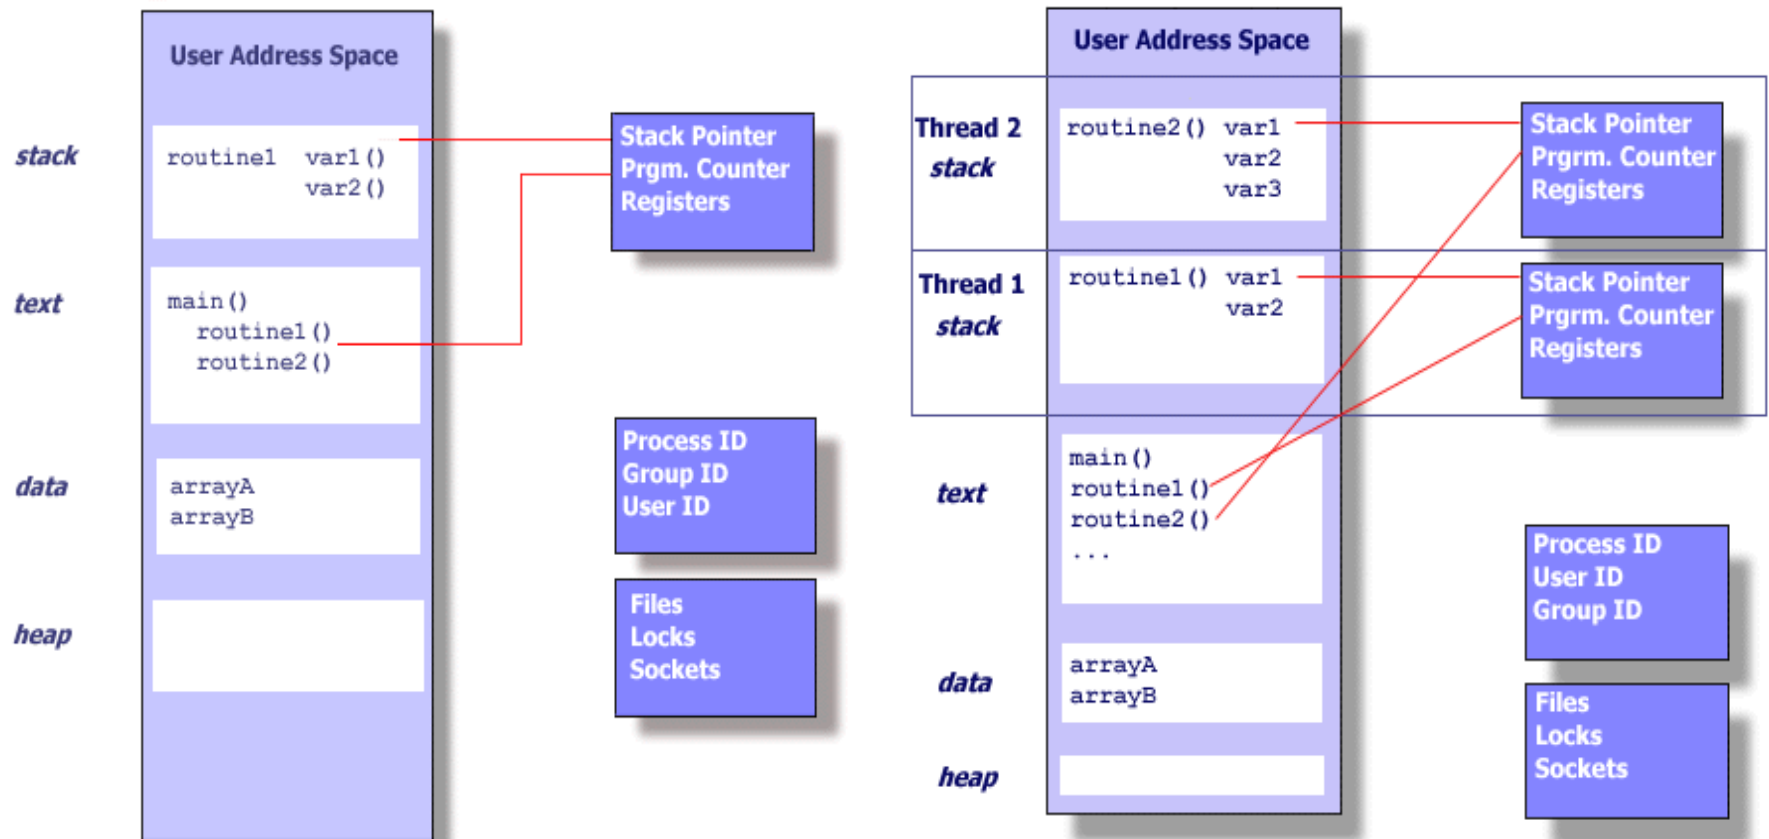
- A process produces 100MB data in memory. You want to share the data with two other processes so that each of which can access half the data (50MB each). What IPC mechanism will you use and why?
- IPC mechanism: POSIX Shared memory
- Reasons: (1) large data → need high performance, (2) no need for synchronization,

# Threads

- Definition
- Key differences compared to process
- Communication between the threads
- User threads vs. Kernel threads
- Key benefits over processes?



# Single and Multithreaded Process



source: <https://computing.llnl.gov/tutorials/pthreads/>

# Recap: Multi-threads vs. Multi-processes

- Multi-processes
  - (+) protection
  - (-) performance (?)
- Multi-threads
  - (+) performance
  - (-) protection



Process-per-tab



Single-process multi-threads

# Synchronization

- Race condition
- Synchronization instructions
  - *test&set, compare&swap*
- Spinlock
  - Spin on wait
  - Good for short critical section but can be wasteful
- Mutex
  - Block (sleep) on wait
  - Good for long critical section but bad for short one

# Race Condition

Initial condition: *counter* = 5

Thread 1

```
R1 = load (counter);  
R1 = R1 + 1;  
counter = store (R1);
```

Thread 2

```
R2 = load (counter);  
R2 = R2 - 1;  
counter = store (R2);
```

- What are the possible outcome?



# Quiz

- Write the pseudo-code definition of *TestAndSet* instruction

```
Int TestAndSet(int *lock)
{
    int ret;
    _____
    _____
    return ret;
}
```

- Complete the following lock implementation using *TestAndSet* instruction

```
void init_lock(int *mutex)
{
    *mutex = 0;
}
void lock(int *mutex)
{
    while (_____);
}
void unlock(int *mutex)
{
    _____
}
```

# Quiz

- Pseudo code of test&set

```
int TestAndSet(int *lock) {  
    int ret = *lock;  
    *lock = 1;  
    return ret;  
}
```

- Spinlock implementation using *TestAndSet* instruction

```
void init_lock(int *lock) {  
    *lock = 0;  
}
```

```
void lock(int *lock) {  
    while (TestAndSet(lock));  
}
```

```
void unlock(int *lock) {  
    *lock = 0;  
}
```

```
void mutex_init (mutex_t *lock)
{
    lock->value = 0;
    list_init(&lock->wait_list);
    spin_lock_init(&lock->wait_lock);
}
```

```
void mutex_lock (mutex_t *lock)
{
    ...
    while(TestAndSet(&lock->value)) {
        current->state = WAITING;
        list_add(&lock->wait_list, current);
    ...
        schedule();
    ...
    }
    ...
}
```

```
void mutex_unlock (mutex_t *lock)
{
    ...
    lock->value = 0;
    if (!list_empty(&lock->wait_list))
        wake_up_process(&lock->wait_list)
    ...
}
```

```
void mutex_init (mutex_t *lock)
```

```
{
```

```
    lock->value = 0;
```

```
    list_init(&lock->wait_list);
```

```
    spin_lock_init(&lock->wait_lock);
```

```
}
```

← Thread waiting list

← To protect waiting list

```
void mutex_lock (mutex_t *lock)
```

```
{
```

```
    spin_lock(&lock->wait_lock);
```

```
    while(TestAndSet(&lock->value)) {
```

```
        current->state = WAITING;
```

```
        list_add(&lock->wait_list, current);
```

```
        spin_unlock(&lock->wait_lock);
```

```
        schedule();
```

```
        spin_lock(&lock->wait_lock);
```

```
    }
```

```
    spin_unlock(&lock->wait_lock);
```

```
}
```

← Thread state change

← Add the current thread to the waiting list

← Sleep or schedule another thread

```
void mutex_unlock (mutex_t *lock)
```

```
{
```

```
    spin_lock(&lock->wait_lock);
```

```
    lock->value = 0;
```

```
    if (!list_empty(&lock->wait_list))
```

```
        wake_up_process(&lock->wait_list)
```

```
    spin_unlock(&lock->wait_lock);
```

```
}
```

← Someone is waiting for the lock

← Wake-up a waiting thread

# Recap: Bounded Buffer Problem Revisit

## Monitor version

```
Mutex lock;
Condition full, empty;

produce (item)
{
    lock.acquire();
    while (queue.isFull())
        empty.wait(&lock);
    queue.enqueue(item);
    full.signal();
    lock.release();
}

consume()
{
    lock.acquire();
    while (queue.isEmpty())
        full.wait(&lock);
    item = queue.dequeue(item);
    empty.signal();
    lock.release();
    return item;
}
```

## Semaphore version


```
Semaphore mutex = 1, full = 0,
empty = N;

produce (item)
{
    empty.P();
    mutex.P();
    queue.enqueue(item);
    mutex.V();
    full.V();
}

consume()
{
    full.P();
    mutex.P();
    item = queue.dequeue();
    mutex.V();
    empty.V();
    return item;
}
```

# Deadlock

- Deadlock conditions
- Resource allocation graph
- Banker's algorithm
- Dining philosopher example
- Other concepts to know
  - Starvation vs. deadlock



# Conditions for Deadlocks

- Mutual exclusion
  - only one process at a time can use a resource
- No preemption
  - resources cannot be preempted, release must be voluntary
- Hold and wait
  - a process must be holding at least one resource, and waiting to acquire additional resources held by other processes
- Circular wait
  - There must be a circular dependency. For example, A waits B, B waits C, and C waits A.
- **All four conditions must simultaneously hold**

# Resource-Allocation Graph


■ Process





■ Resource Type with 4 instances



■  $P_i$  requests instance of  $R_j$




■  $P_i$  is holding an instance of  $R_j$





# Quiz

Draw a resource allocation graph for the following.



# Quiz

- Using Banker's algorithm, determine whether this state is safe or unsafe.

Total resources: 10

**Avail** resources: 1

Process	Max	Alloc
$P_0$	10	4
$P_1$	3	1
$P_2$	6	4

# Scheduling

- Three main schedulers
  - FCFS, SJF/SRTF, RR
  - Gant chart examples
- Other concepts to know
  - Fair scheduling (CFS)
  - Fixed priority scheduling
  - Multi-level queue scheduling
  - Load balancing and multicore scheduling


# Round-Robin (RR)

- Example

- Quantum size = 2

Process	Burst Times
P1	24
P2	3
P3	3

- Gantt chart



- **Response time** (between ready to first schedule)

- P1: 0, P2: 2, P3: 4. average response time =  $(0+2+4)/3 = 2$

- **Waiting time**

- P1: 6, P2: 6, P3: 7. average waiting time =  $(6+6+7)/3 = 6.33$